# X-Mem

2.2.7

Generated by Doxygen 1.8.6

Tue Jun 30 2015 20:25:40

# Contents


**1 README** 1

**2 Hierarchical Index** 7

2.1 Class Hierarchy . . . 7

**3 Class Index** 9

3.1 Class List . . . 9

**4 File Index** 11

4.1 File List . . . 11

**5 Class Documentation** 13

5.1 xmem::Parser::Action Struct Reference . . . 13

5.1.1 Member Function Documentation . . . 13

5.1.1.1 finished . . . 13

5.1.1.2 perform . . . 14

5.2 xmem::Arg Struct Reference . . . 14

5.2.1 Detailed Description . . . 15

5.3 xmem::Benchmark Class Reference . . . 15

5.3.1 Detailed Description . . . 18

5.3.2 Constructor & Destructor Documentation . . . 18

5.3.2.1 Benchmark . . . 18

5.3.3 Member Function Documentation . . . 18

5.3.3.1 _run_core . . . 18

5.3.3.2 _start_power_threads . . . 18

5.3.3.3 _stop_power_threads . . . 19

5.3.3.4 getAverageDRAMPower . . . 19

5.3.3.5 getAverageMetric . . . 19

5.3.3.6 getChunkSize . . . 19

5.3.3.7 getCPUNode . . . 19

5.3.3.8 getIterations . . . 19

5.3.3.9 getLen . . . 20

5.3.3.10 getMemNode . . . 20

5.3.3.11 getMetricOnIter . . . 20
5.3.3.12 getMetricUnits . . . 20
5.3.3.13 getName . . . 20
5.3.3.14 getNumThreads . . . 20
5.3.3.15 getPatternMode . . . 21
5.3.3.16 getPeakDRAMPower . . . 21
5.3.3.17 getRWMode . . . 21
5.3.3.18 getStrideSize . . . 21
5.3.3.19 hasRun . . . 21
5.3.3.20 isValid . . . 21
5.3.3.21 run . . . 22
5.3.4 Member Data Documentation . . . 22
5.3.4.1 _average_dram_power_socket . . . 22
5.3.4.2 _averageMetric . . . 22
5.3.4.3 _chunk_size . . . 22
5.3.4.4 _cpu_node . . . 22
5.3.4.5 _dram_power_readers . . . 22
5.3.4.6 _dram_power_threads . . . 22
5.3.4.7 _hasRun . . . 23
5.3.4.8 _iterations . . . 23
5.3.4.9 _len . . . 23
5.3.4.10 _mem_array . . . 23
5.3.4.11 _mem_node . . . 23
5.3.4.12 _metricOnIter . . . 23
5.3.4.13 _metricUnits . . . 23
5.3.4.14 _name . . . 23
5.3.4.15 _num_worker_threads . . . 23
5.3.4.16 _obj_valid . . . 23
5.3.4.17 _pattern_mode . . . 23
5.3.4.18 _peak_dram_power_socket . . . 23
5.3.4.19 _rw_mode . . . 24
5.3.4.20 _stride_size . . . 24
5.3.4.21 _warning . . . 24
5.4 xmem::BenchmarkManager Class Reference . . . 24
5.4.1 Detailed Description . . . 24
5.4.2 Constructor & Destructor Documentation . . . 24
5.4.2.1 BenchmarkManager . . . 24
5.4.3 Member Function Documentation . . . 25
5.4.3.1 runAll . . . 25
5.4.3.2 runLatencyBenchmarks . . . 25

5.4.3.3 runThroughputBenchmarks . . . 26
5.5 xmem::Configurator Class Reference . . . 26
5.5.1 Detailed Description . . . 28
5.5.2 Member Function Documentation . . . 28
5.5.2.1 configureFromInput . . . 28
5.5.2.2 extensionsEnabled . . . 28
5.5.2.3 getIterationsPerTest . . . 29
5.5.2.4 getNumWorkerThreads . . . 29
5.5.2.5 getOutputFilename . . . 29
5.5.2.6 getStartingTestIndex . . . 29
5.5.2.7 getWorkingSetSizePerThread . . . 29
5.5.2.8 isNUMAEnabled . . . 29
5.5.2.9 latencyTestSelected . . . 29
5.5.2.10 runExtDelayInjectedLoadedLatencyBenchmark . . . 30
5.5.2.11 setUseOutputFile . . . 30
5.5.2.12 throughputTestSelected . . . 30
5.5.2.13 useChunk32b . . . 30
5.5.2.14 useLargePages . . . 30
5.5.2.15 useOutputFile . . . 30
5.5.2.16 useRandomAccessPattern . . . 30
5.5.2.17 useReads . . . 31
5.5.2.18 useSequentialAccessPattern . . . 31
5.5.2.19 useStrideN1 . . . 31
5.5.2.20 useStrideN16 . . . 31
5.5.2.21 useStrideN2 . . . 31
5.5.2.22 useStrideN4 . . . 31
5.5.2.23 useStrideN8 . . . 31
5.5.2.24 useStrideP1 . . . 32
5.5.2.25 useStrideP16 . . . 32
5.5.2.26 useStrideP2 . . . 32
5.5.2.27 useStrideP4 . . . 32
5.5.2.28 useStrideP8 . . . 32
5.5.2.29 useWrites . . . 32
5.5.2.30 verboseMode . . . 33
5.6 xmem::Stats::CountOptionsAction Class Reference . . . 33
5.6.1 Constructor & Destructor Documentation . . . 34
5.6.1.1 CountOptionsAction . . . 34
5.6.2 Member Function Documentation . . . 34
5.6.2.1 perform . . . 34
5.7 xmem::Descriptor Struct Reference . . . 34

5.7.1 Detailed Description . . . 34
5.7.2 Member Data Documentation . . . 35
5.7.2.1 check_arg . . . 35
5.7.2.2 help . . . 35
5.7.2.3 index . . . 35
5.7.2.4 longopt . . . 35
5.7.2.5 shortopt . . . 36
5.7.2.6 type . . . 36
5.8 xmem::ExampleArg Class Reference . . . 37
5.9 xmem::PrintUsageImplementation::FunctionWriter< Function > Struct Template Reference . . . 38
5.10 xmem::PrintUsageImplementation::IStringWriter Struct Reference . . . 39
5.11 xmem::LatencyBenchmark Class Reference . . . 39
5.11.1 Detailed Description . . . 41
5.11.2 Member Function Documentation . . . 41
5.11.2.1 _run_core . . . 41
5.11.2.2 getAvgLoadMetric . . . 41
5.11.2.3 getLoadMetricOnIter . . . 41
5.11.3 Member Data Documentation . . . 42
5.11.3.1 _averageLoadMetric . . . 42
5.11.3.2 _loadMetricOnIter . . . 42
5.12 xmem::LatencyWorker Class Reference . . . 42
5.12.1 Detailed Description . . . 43
5.12.2 Constructor & Destructor Documentation . . . 43
5.12.2.1 LatencyWorker . . . 43
5.13 xmem::PrintUsageImplementation::LinePartIterator Class Reference . . . 44
5.13.1 Member Function Documentation . . . 44
5.13.1.1 next . . . 44
5.13.1.2 nextRow . . . 44
5.13.1.3 nextTable . . . 45
5.14 xmem::PrintUsageImplementation::LineWrapper Class Reference . . . 45
5.14.1 Constructor & Destructor Documentation . . . 45
5.14.1.1 LineWrapper . . . 45
5.14.2 Member Function Documentation . . . 46
5.14.2.1 process . . . 46
5.15 xmem::LoadWorker Class Reference . . . 46
5.15.1 Detailed Description . . . 48
5.15.2 Constructor & Destructor Documentation . . . 48
5.15.2.1 LoadWorker . . . 48
5.15.2.2 LoadWorker . . . 48
5.16 xmem::MemoryWorker Class Reference . . . 48

5.16.1 Detailed Description . . . 50
5.16.2 Constructor & Destructor Documentation . . . 50
5.16.2.1 MemoryWorker . . . 50
5.16.3 Member Function Documentation . . . 50
5.16.3.1 getAdjustedTicks . . . 50
5.16.3.2 getBytesPerPass . . . 51
5.16.3.3 getElapsedDummyTicks . . . 51
5.16.3.4 getElapsedTicks . . . 52
5.16.3.5 getLen . . . 52
5.16.3.6 getPasses . . . 53
5.16.3.7 hadWarning . . . 53
5.16.4 Member Data Documentation . . . 54
5.16.4.1 _adjusted_ticks . . . 54
5.16.4.2 _bytes_per_pass . . . 54
5.16.4.3 _completed . . . 54
5.16.4.4 _cpu_affinity . . . 54
5.16.4.5 _elapsed_dummy_ticks . . . 54
5.16.4.6 _elapsed_ticks . . . 54
5.16.4.7 _len . . . 54
5.16.4.8 _mem_array . . . 54
5.16.4.9 _passes . . . 55
5.16.4.10 _warning . . . 55
5.17 xmem::MyArg Class Reference . . . 55
5.18 xmem::Option Class Reference . . . 56
5.18.1 Detailed Description . . . 58
5.18.2 Constructor & Destructor Documentation . . . 58
5.18.2.1 Option . . . 58
5.18.2.2 Option . . . 58
5.18.3 Member Function Documentation . . . 59
5.18.3.1 append . . . 59
5.18.3.2 count . . . 59
5.18.3.3 first . . . 59
5.18.3.4 isFirst . . . 60
5.18.3.5 isLast . . . 60
5.18.3.6 last . . . 60
5.18.3.7 next . . . 61
5.18.3.8 nextwrap . . . 61
5.18.3.9 operator const Option * . . . 61
5.18.3.10 operator Option * . . . 62
5.18.3.11 operator= . . . 62

5.18.3.12 prev . . . 62
5.18.3.13 prevwrap . . . 62
5.18.3.14 type . . . 63
5.18.4 Member Data Documentation . . . 63
5.18.4.1 arg . . . 63
5.18.4.2 desc . . . 63
5.18.4.3 name . . . 63
5.18.4.4 namelen . . . 63
5.19 xmem::PrintUsageImplementation::OStreamWriter< OStream > Struct Template Reference . . . 64
5.20 xmem::Parser Class Reference . . . 65
5.20.1 Detailed Description . . . 66
5.20.2 Constructor & Destructor Documentation . . . 66
5.20.2.1 Parser . . . 66
5.20.3 Member Function Documentation . . . 67
5.20.3.1 error . . . 68
5.20.3.2 nonOptions . . . 68
5.20.3.3 nonOptionsCount . . . 68
5.20.3.4 optionsCount . . . 68
5.20.3.5 parse . . . 68
5.21 xmem::PowerReader Class Reference . . . 69
5.21.1 Detailed Description . . . 71
5.21.2 Constructor & Destructor Documentation . . . 71
5.21.2.1 PowerReader . . . 71
5.21.3 Member Function Documentation . . . 72
5.21.3.1 calculateMetrics . . . 72
5.21.3.2 clear . . . 72
5.21.3.3 clear_and_reset . . . 72
5.21.3.4 getAveragePower . . . 73
5.21.3.5 getLastSample . . . 73
5.21.3.6 getNumSamples . . . 74
5.21.3.7 getPeakPower . . . 74
5.21.3.8 getPowerTrace . . . 75
5.21.3.9 getPowerUnits . . . 75
5.21.3.10 getSamplingPeriod . . . 76
5.21.3.11 name . . . 76
5.21.3.12 stop . . . 77
5.21.4 Member Data Documentation . . . 77
5.21.4.1 _average_power . . . 77
5.21.4.2 _cpu_affinity . . . 77
5.21.4.3 _name . . . 77

5.21.4.4 _num_samples . . . 78
5.21.4.5 _peak_power . . . 78
5.21.4.6 _power_trace . . . 78
5.21.4.7 _power_units . . . 78
5.21.4.8 _sampling_period . . . 78
5.21.4.9 _stop_signal . . . 78
5.22 xmem::PrintUsageImplementation Struct Reference . . . 78
5.22.1 Member Function Documentation . . . 79
5.22.1.1 isWideChar . . . 79
5.23 xmem::Runnable Class Reference . . . 79
5.23.1 Detailed Description . . . 80
5.23.2 Member Function Documentation . . . 80
5.23.2.1 _acquireLock . . . 80
5.23.2.2 _releaseLock . . . 80
5.24 xmem::Stats Struct Reference . . . 80
5.24.1 Detailed Description . . . 81
5.24.2 Constructor & Destructor Documentation . . . 81
5.24.2.1 Stats . . . 81
5.24.3 Member Function Documentation . . . 82
5.24.3.1 add . . . 82
5.24.4 Member Data Documentation . . . 82
5.24.4.1 buffer_max . . . 82
5.24.4.2 options_max . . . 82
5.25 xmem::Parser::StoreOptionAction Class Reference . . . 83
5.25.1 Constructor & Destructor Documentation . . . 83
5.25.1.1 StoreOptionAction . . . 83
5.25.2 Member Function Documentation . . . 84
5.25.2.1 finished . . . 84
5.25.2.2 perform . . . 84
5.26 xmem::PrintUsageImplementation::StreamWriter< Function, Stream > Struct Template Reference 85
5.27 xmem::PrintUsageImplementation::SyscallWriter< Syscall > Struct Template Reference . . . 86
5.28 xmem::PrintUsageImplementation::TemporaryWriter< Temporary > Struct Template Reference . . . 87
5.29 xmem::Thread Class Reference . . . 88
5.29.1 Detailed Description . . . 88
5.29.2 Constructor & Destructor Documentation . . . 88
5.29.2.1 Thread . . . 88
5.29.2.2 ~Thread . . . 88
5.29.3 Member Function Documentation . . . 89
5.29.3.1 cancel . . . 89
5.29.3.2 completed . . . 89

5.29.3.3 create_and_start . . . 89
5.29.3.4 created . . . 89
5.29.3.5 getExitCode . . . 89
5.29.3.6 getTarget . . . 89
5.29.3.7 isThreadRunning . . . 89
5.29.3.8 isThreadSuspended . . . 90
5.29.3.9 join . . . 90
5.29.3.10 started . . . 90
5.29.3.11 validTarget . . . 90
5.30 xmem::ThroughputBenchmark Class Reference . . . 90
5.30.1 Detailed Description . . . 91
5.30.2 Member Function Documentation . . . 91
5.30.2.1 _run_core . . . 91
5.31 xmem::Timer Class Reference . . . 92
5.31.1 Detailed Description . . . 92
5.31.2 Member Function Documentation . . . 93
5.31.2.1 get_ns_per_tick . . . 93
5.31.2.2 get_ticks_per_ms . . . 93
5.31.3 Member Data Documentation . . . 93
5.31.3.1 _ns_per_tick . . . 93
5.31.3.2 _ticks_per_ms . . . 93

**6 File Documentation** **95**

6.1 src/Benchmark.cpp File Reference . . . 95
6.1.1 Detailed Description . . . 95
6.2 src/benchmark_kernels.cpp File Reference . . . 95
6.2.1 Detailed Description . . . 96
6.3 src/BenchmarkManager.cpp File Reference . . . 96
6.3.1 Detailed Description . . . 97
6.4 src/common.cpp File Reference . . . 97
6.4.1 Detailed Description . . . 98
6.5 src/Configurator.cpp File Reference . . . 98
6.5.1 Detailed Description . . . 98
6.6 src/ext/DelayInjectedLoadedLatencyBenchmark/delay_injected_benchmark_kernels.cpp File Reference . . . 99
6.6.1 Detailed Description . . . 99
6.7 src/ext/DelayInjectedLoadedLatencyBenchmark/DelayInjectedLoadedLatencyBenchmark.cpp File Reference . . . 99
6.7.1 Detailed Description . . . 100
6.8 src/include/Benchmark.h File Reference . . . 100
6.8.1 Detailed Description . . . 101

6.9 src/include/benchmark_kernels.h File Reference . . . 101
6.9.1 Detailed Description . . . 104
6.10 src/include/BenchmarkManager.h File Reference . . . 104
6.10.1 Detailed Description . . . 105
6.11 src/include/common.h File Reference . . . 105
6.11.1 Detailed Description . . . 108
6.11.2 Macro Definition Documentation . . . 108
6.11.2.1 BENCHMARK_DURATION_MS . . . 108
6.11.2.2 DEFAULT_LARGE_PAGE_SIZE . . . 108
6.11.2.3 DEFAULT_NUM_L1_CACHES . . . 108
6.11.2.4 DEFAULT_NUM_L2_CACHES . . . 108
6.11.2.5 DEFAULT_NUM_L3_CACHES . . . 109
6.11.2.6 DEFAULT_NUM_L4_CACHES . . . 109
6.11.2.7 DEFAULT_NUM_LOGICAL_CPUS . . . 109
6.11.2.8 DEFAULT_NUM_NODES . . . 109
6.11.2.9 DEFAULT_NUM_PHYSICAL_CPUS . . . 109
6.11.2.10 DEFAULT_NUM_PHYSICAL_PACKAGES . . . 109
6.11.2.11 DEFAULT_NUM_WORKER_THREADS . . . 109
6.11.2.12 DEFAULT_PAGE_SIZE . . . 109
6.11.2.13 DEFAULT_WORKING_SET_SIZE_PER_THREAD . . . 109
6.11.2.14 EXT_DELAY_INJECTED_LOADED_LATENCY_BENCHMARK . . . 109
6.11.2.15 LATENCY_BENCHMARK_UNROLL_LENGTH . . . 109
6.11.2.16 MIN_ELAPSED_TICKS . . . 109
6.11.2.17 POWER_SAMPLING_PERIOD_MS . . . 110
6.11.2.18 THROUGHPUT_BENCHMARK_BYTES_PER_PASS . . . 110
6.11.2.19 USE_OS_TIMER . . . 110
6.12 src/include/Configurator.h File Reference . . . 110
6.12.1 Detailed Description . . . 111
6.13 src/include/ExampleArg.h File Reference . . . 111
6.13.1 Detailed Description . . . 112
6.14 src/include/ext/DelayInjectedLoadedLatencyBenchmark/delay_injected_benchmark_kernels.h File Reference . . . 112
6.14.1 Detailed Description . . . 115
6.15 src/include/ext/DelayInjectedLoadedLatencyBenchmark/DelayInjectedLoadedLatencyBenchmark.h File Reference . . . 115
6.15.1 Detailed Description . . . 115
6.16 src/include/LatencyBenchmark.h File Reference . . . 115
6.16.1 Detailed Description . . . 116
6.17 src/include/LatencyWorker.h File Reference . . . 117
6.17.1 Detailed Description . . . 117

6.18 src/include/LoadWorker.h File Reference . . . 118
6.18.1 Detailed Description . . . 118
6.19 src/include/MemoryWorker.h File Reference . . . 119
6.19.1 Detailed Description . . . 119
6.20 src/include/MyArg.h File Reference . . . 119
6.20.1 Detailed Description . . . 120
6.21 src/include/optionparser.h File Reference . . . 121
6.21.1 Detailed Description . . . 122
6.22 src/include/PowerReader.h File Reference . . . 125
6.22.1 Detailed Description . . . 125
6.23 src/include/Runnable.h File Reference . . . 126
6.23.1 Detailed Description . . . 126
6.24 src/include/Thread.h File Reference . . . 126
6.24.1 Detailed Description . . . 127
6.25 src/include/ThroughputBenchmark.h File Reference . . . 127
6.25.1 Detailed Description . . . 128
6.26 src/include/Timer.h File Reference . . . 128
6.26.1 Detailed Description . . . 129
6.27 src/include/win/win_common_third_party.h File Reference . . . 129
6.27.1 Detailed Description . . . 130
6.28 src/include/win/win_CPdhQuery.h File Reference . . . 130
6.28.1 Detailed Description . . . 130
6.29 src/include/win/WindowsDRAMPowerReader.h File Reference . . . 130
6.29.1 Detailed Description . . . 130
6.30 src/LatencyBenchmark.cpp File Reference . . . 130
6.30.1 Detailed Description . . . 131
6.31 src/LatencyWorker.cpp File Reference . . . 131
6.31.1 Detailed Description . . . 131
6.32 src/LoadWorker.cpp File Reference . . . 131
6.32.1 Detailed Description . . . 132
6.33 src/main.cpp File Reference . . . 132
6.33.1 Detailed Description . . . 133
6.34 src/MemoryWorker.cpp File Reference . . . 133
6.34.1 Detailed Description . . . 133
6.35 src/PowerReader.cpp File Reference . . . 133
6.35.1 Detailed Description . . . 134
6.36 src/Runnable.cpp File Reference . . . 134
6.36.1 Detailed Description . . . 135
6.37 src/Thread.cpp File Reference . . . 135
6.37.1 Detailed Description . . . 135

6.38 src/ThroughputBenchmark.cpp File Reference . . . . . 135
6.38.1 Detailed Description . . . . . 136
6.39 src/Timer.cpp File Reference . . . . . 136
6.39.1 Detailed Description . . . . . 137
6.40 src/win/win_common_third_party.cpp File Reference . . . . . 137
6.40.1 Detailed Description . . . . . 137
6.41 src/win/WindowsDRAMPowerReader.cpp File Reference . . . . . 137
6.41.1 Detailed Description . . . . . 137

**Index** **138**

# Chapter 1

# README

### X-Mem: A Cross-Platform and Extensible Memory Characterization Tool for the Cloud v2.2.7

X-Mem is a flexible open-source research tool for characterizing memory hierarchy throughput, latency, power, and more. The tool was developed jointly by Microsoft and the UCLA NanoCAD Lab. This project was started by Mark Gottscho (Email: mgottscho@ucla.edu) as a Summer 2014 PhD intern at Microsoft Research. X-Mem is released freely and open-source under the MIT License. The project is under active development. Stay tuned for more updates.

PROJECT REVISION DATE: June 30, 2015.

### VERSION CONTROL AND OBTAINING SOURCE CODE

This project is under version control using git. Its master repository is hosted at https://github.com/Microsoft/X-Mem.git. There is also another mirrored fork at https://github.com/nanocad-lab/X-Mem.git.

### OBTAINING PRE-BUILT BINARIES

If you don't want to build X-Mem from source, you can get regular releases of pre-built binaries for Windows and GNU/Linux from the project homepage, located at https://nanocad-lab.github.io/X-Mem.

### LICENSE

The MIT License (MIT)

Copyright (c) 2014 Microsoft

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

## DESIGN PHILOSOPHY

Previous memory benchmarking tools such as STREAM/STREAM2, lmbench3, and mlc all had various shortcomings that made them inadequate for emerging needs in the cloud. In particular, cloud systems present four key challenges that we addressed with X-Mem. Stay tuned for updates – more information will be coming soon.

## FEATURES

This tool is provided as open source with the hope of being useful to the broader research and development community. Here are some of X-Mem's features.

Flexibility: Easy reconfiguration for different combinations of tests

- Working sets in increments of 4KB, allowing cache up to main memory-level benchmarking
- NUMA support
- Multi-threading support
- Large page support

Extensibility: Modularity via C++ object-oriented principles

- Supports rapid addition of new benchmark kernel routines
- Example: stream triad algorithm, impact of false sharing, etc. are possible with minor changes

Cross-platform: Currently implemented for two OSes and architecture families

- Windows: Intel x86 (32-bit), x86-64, and x86-64 with AVX extensions
- GNU/Linux: Intel x86 (32-bit), x86-64, and x86-64 with AVX extensions, ARM (32-bit), ARM (32-bit) with NEON, ARMv8 (64-bit)
- Designed to allow straightforward porting to other operating systems and ISAs
- ARM on Windows currently not possible due to some incompatibilites and/or lack of support. This may be resolved in the future.

Memory throughput:

- Accurate measurement of sustained memory throughput to all levels of cache and memory
- Regular access patterns: forward & reverse sequential as well as strides of 2, 4, 8, and 16 words
- Random access patterns
- Read and write
- 32, 64, 128, 256-bit width memory instructions where applicable on each architecture

Memory latency:

- Accurate measurement of round-trip memory latency to all levels of cache and memory
- Loaded and unloaded latency via use of multithreaded load generation

Memory power:

- Support custom power instrumentation through a simple interface that end-users can implement
- Can collect DRAM power via custom driver exposed in Windows performance counter API

Documentation:

- Extensive Doxygen source code comments, PDF manual, HTML

INCLUDED EXTENSIONS (under src/include/ext and src/ext directories):

- Loaded latency benchmark variant with load delays inserted as nop instructions between memory instructions. This is done for 32, 64, 128, and 256-bit load chunk sizes where applicable using the forward sequential read pattern.

For feature requests, please refer to the contact information at the end of this README.

## RUNTIME PREREQUISITES

There are a few runtime prerequisites in order for the software to run correctly.

HARDWARE:

- Intel x86, x86-64, or x86-64 with AVX CPU. AMD CPUs should also work although this has not been tested!
- ARM Cortex-A series processors with VFP and NEON extensions. Specifically tested on ARM Cortex A9 (32-bit) which is ARMv7. 64-bit builds for ARMv8-A should also work but have not been tested. GNU/Linux builds only. You can attempt Windows, but I have issues linking the 32-bit version and compiling the NEON version. Also, as far as I can see there is no Windows support for ARMv8 at all.

WINDOWS:

- Microsoft Windows 8.0 64-bit or later, Server 2012 or later.
- Microsoft Visual C++ 2013 Redistributables (32-bit) – for x86 (32-bit) builds
- Microsoft Visual C++ 2013 Redistributables (64-bit) – for x86-64 and x86-64 with AVX builds
- You MAY need Administrator privileges, in order to:
  - Use large pages, if the –large_pages option is selected (see USAGE, below)
  - The first time you use –large_pages on a given Windows machine, you may need to ensure that your Windows user account has the necessary rights to allow lockable memory pages. To do this on Windows 8, run gpedit.msc –> Local Computer Policy –> Computer Configuration –> Windows Settings –> Security Settings –> Local Policies –> User Rights Assignment –> Add your username to "Lock pages in memory". Then log out and then log back in.
  - Use the PowerReader interface, depending on end-user implementation
  - Elevate thread priority and pin threads to logical CPUs for improved performance and benchmarking consistency

GNU/LINUX:

- GNU utilities with support for C++11. Tested with gcc 4.8.2 on Ubuntu 14.04 LTS for x86 (32-bit), x86-64, and x86-64 with AVX on Intel Sandy Bridge, Ivy Bridge, and Haswell families.
- libhugetlbfs. You can obtain it at http://libhugetlbfs.sourceforge.net. On Ubuntu systems, you can install using "sudo apt-get install libhugetlbfs0".
- Potentially, administrator privileges, if you plan to use the –large_pages option.
  - During runtime, if the –large_pages option is selected, you may need to first manually ensure that large pages are available from the OS. This can be done by running "hugeadm --pool-list". It is recommended to set minimum pool to 1GB (in order to measure DRAM effectively). If needed, this can be done by running "hugeadm --pool-pages-min 2MB:512". Alternatively, run the linux_setup_runtime_hugetlbfs.sh script that is provided with X-Mem.

### INSTALLATION

The only file that is needed to run on Windows is the respective executable xmem-win-<ARCH>.exe on Windows, and xmem-linux-<ARCH> on GNU/Linux. It has no other dependencies aside from the pre-installed system pre-requisites listed above.

## USAGE

To get help using the tool, simply run it with the "-h" or "--help" option. It includes a list of all available options and provides usage examples.

## BUILDING FROM SOURCE

After you have set the desired compile-time options, build the source. On Windows, running build-win.bat should suffice. On GNU/Linux, run build-linux.sh. Each takes a single argument specifying the target architecture.

If you customize your build, make sure you use the "Release" mode for your OS/compiler. Do not include debug capabilities as it can dramatically affect performance of the benchmarks, leading to pessimistic results.

## BUILD PREREQUISITES

There are a few software build prerequisites, depending on your platform.

WINDOWS:

- Any version of Visual Studio 2013 64-bit (also known as version 12.0).
- Python 2.7. You can obtain it at http://www.python.org.
- SCons build system. You can obtain it at http://www.scons.org. Build tested with SCons 2.3.4.

GNU/LINUX:

- bash shell. Other shells could work but are untested.
- gcc with support for the C++11 standard. Tested with gcc version 4.8.2 on Ubuntu 14.04 LTS for x86 (32-bit), x86-64, x86-64 with AVX builds.
- gcc cross-compiler for ARM targets (assumed build on x86-64 Ubuntu host).
- Python 2.7. You can obtain it at http://www.python.org. On Ubuntu systems, you can install using "sudo apt-get install python2.7". You may need some other Python 2.7 packages as well.
- SCons build system. You can obtain it at http://www.scons.org. On Ubuntu systems, you can install using "sudo apt-get install scons". Build tested with SCons 2.3.4.
- Kernel support for large (huge) pages. This support can be verified on your Linux installation by running "grep hugetlbfs /proc/filesystems". If you do not have huge page support in your kernel, you can build a kernel with the appropriate options switched on: "CONFIG_HUGETLB_PAGE" and "CONFIG_HUGETLBFS".
- libhugetlbfs. This is used for allocating large (huge) pages if the –large_pages runtime option is selected. You can obtain it at http://libhugetlbfs.sourceforge.net. On Ubuntu systems, you can install using "sudo apt-get install libhugetlbfs-dev".

### DOCUMENTATION BUILD PREREQUISITES

The following tools are only needed for automatically regenerating source code documentation with HTML and PDF.

WINDOWS:

- doxygen tool. You can obtain it at http://www.stack.nl/~dimitri/doxygen.

- LaTeX distribution. You can get a Windows distribution at http://www.miktex.org.
- make for Windows. You can obtain it at http://gnuwin32.sourceforge.net/packages/make.htm. You will have to manually add it to your Windows path.

GNU/LINUX:

- doxygen tool. You can obtain it at http://www.stack.nl/~dimitri/doxygen. On Ubuntu systems, you can install with "sudo apt-get install doxygen".
- LaTeX distribution. On Ubuntu systems, LaTeX distributed with doxygen should actually be sufficient. You can install with "sudo apt-get install doxygen-latex".
- make. This should be included on any GNU/Linux system.

### SOURCE CODE DOCUMENTATION

The tool comes with built-in Doxygen comments in the source code, which can be used to generate both HTML and LaTeX –> PDF documentation. Documentation is maintained under the doc/ subdirectory. To build documentation after modifying the source, run build-docs-win.bat on Windows, or build-docs-linux.sh on GNU/Linux systems. Note that Doxygen and a LaTeX distribution must be installed on the system.

## CONTACT, FEEDBACK, AND BUG REPORTS

For questions, comments, criticism, bug reports, and other feedback for this software, please contact Mark Gottscho via email at mgottscho@ucla.edu or via web at http://www.seas.ucla.edu/~gottscho.

## ACKNOWLEDGMENT

Mark Gottscho would like to thank Dr. Mohammed Shoaib of Microsoft Research and Dr. Sriram Govindan of Microsoft for their mentorship in the creation of this software. Further thanks to Dr. Bikash Sharma, Mark Santaniello, Mike Andrewartha, and Laura Caulfield of Microsoft for their contributions, feedback, and assistance. Thank you as well to Dr. Jie Liu of Microsoft Research, Dr. Badriddine Khessib and Dr. Kushagra Vaid of Microsoft, and Prof. Puneet Gupta of UCLA for giving Mark the opportunity to create this work. Finally, Mark would like to thank Dr. Fedor Pikus of Mentor Graphics for teaching him some useful HPC programming techniques.

# Chapter 2

# Hierarchical Index

## 2.1 Class Hierarchy

This inheritance list is sorted roughly, but not completely, alphabetically:


xmem::Parser::Action . . . 13
- xmem::Parser::StoreOptionAction . . . 83
- xmem::Stats::CountOptionsAction . . . 33

xmem::Arg . . . 14
- xmem::ExampleArg . . . 37
  - xmem::MyArg . . . 55

xmem::Benchmark . . . 15
- xmem::LatencyBenchmark . . . 39
- xmem::ThroughputBenchmark . . . 90

xmem::BenchmarkManager . . . 24
xmem::Configurator . . . 26
xmem::Descriptor . . . 34
xmem::PrintUsageImplementation::IStringWriter . . . 39
- xmem::PrintUsageImplementation::FunctionWriter< Function > . . . 38
- xmem::PrintUsageImplementation::OStreamWriter< OStream > . . . 64
- xmem::PrintUsageImplementation::StreamWriter< Function, Stream > . . . 85
- xmem::PrintUsageImplementation::SyscallWriter< Syscall > . . . 86
- xmem::PrintUsageImplementation::TemporaryWriter< Temporary > . . . 87

xmem::PrintUsageImplementation::LinePartIterator . . . 44
xmem::PrintUsageImplementation::LineWrapper . . . 45
xmem::Option . . . 56
xmem::Parser . . . 65
xmem::PrintUsageImplementation . . . 78
xmem::Runnable . . . 79
- xmem::MemoryWorker . . . 48
  - xmem::LatencyWorker . . . 42
  - xmem::LoadWorker . . . 46
- xmem::PowerReader . . . 69

xmem::Stats . . . 80
xmem::Thread . . . 88
xmem::Timer . . . 92

# Chapter 3

# Class Index

## 3.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:


xmem::Parser::Action . . . 13
xmem::Arg
Functions for checking the validity of option arguments . . . 14
xmem::Benchmark
Flexible abstract class for any memory benchmark . . . 15
xmem::BenchmarkManager
Manages running all benchmarks at a high level . . . 24
xmem::Configurator
Handles all user input interpretation and generates the necessary flags for running benchmarks 26
xmem::Stats::CountOptionsAction . . . 33
xmem::Descriptor
Describes an option, its help text (usage) and how it should be parsed . . . 34
xmem::ExampleArg . . . 37
xmem::PrintUsageImplementation::FunctionWriter< Function > . . . 38
xmem::PrintUsageImplementation::IStringWriter . . . 39
xmem::LatencyBenchmark
A type of benchmark that measures memory latency via random pointer chasing. Loading may be provided with separate threads which access memory as quickly as possible using given access patterns . . . 39
xmem::LatencyWorker
Multithreading-friendly class to do memory loading . . . 42
xmem::PrintUsageImplementation::LinePartIterator . . . 44
xmem::PrintUsageImplementation::LineWrapper . . . 45
xmem::LoadWorker
Multithreading-friendly class to do memory loading . . . 46
xmem::MemoryWorker
Multithreading-friendly class to run memory access kernels . . . 48
xmem::MyArg . . . 55
xmem::Option
A parsed option from the command line together with its argument if it has one . . . 56
xmem::PrintUsageImplementation::OStreamWriter< OStream > . . . 64
xmem::Parser
Checks argument vectors for validity and parses them into data structures that are easier to work with . . . 65
xmem::PowerReader
An abstract base class for measuring power from an arbitrary source. This class is runnable using a worker thread . . . 69
xmem::PrintUsageImplementation . . . 78

xmem::Runnable
A base class for any object that implements a thread-safe run() function for use by Thread objects 79
xmem::Stats
Determines the minimum lengths of the buffer and options arrays used for Parser . . . . . 80
xmem::Parser::StoreOptionAction . . . . . 83
xmem::PrintUsageImplementation::StreamWriter< Function, Stream > . . . . . 85
xmem::PrintUsageImplementation::SyscallWriter< Syscall > . . . . . 86
xmem::PrintUsageImplementation::TemporaryWriter< Temporary > . . . . . 87
xmem::Thread
Nice wrapped thread interface independent of particular OS API . . . . . 88
xmem::ThroughputBenchmark
A type of benchmark that measures memory throughput . . . . . 90
xmem::Timer
This class abstracts some characteristics of simple high resolution stopwatch timer. However, due to the inability or complexity of abstracting shared hardware timers, this class does not actually provide start and stop functions . . . . . 92

# Chapter 4

# File Index

## 4.1 File List

Here is a list of all documented files with brief descriptions:


src/Benchmark.cpp
: Implementation file for the Benchmark class . . . 95

src/benchmark_kernels.cpp
: Implementation file for benchmark kernel functions for doing the actual work we care about. :) . 95

src/BenchmarkManager.cpp
: Implementation file for the BenchmarkManager class . . . 96

src/common.cpp
: Implementation file for common preprocessor definitions, macros, functions, and global constants . . . 97

src/Configurator.cpp
: Implementation file for the Configurator class and some helper data structures . . . 98

src/LatencyBenchmark.cpp
: Implementation file for the LatencyBenchmark class . . . 130

src/LatencyWorker.cpp
: Implementation file for the LatencyWorker class . . . 131

src/LoadWorker.cpp
: Implementation file for the LoadWorker class . . . 131

src/main.cpp
: Main entry point to the tool . . . 132

src/MemoryWorker.cpp
: Implementation file for the MemoryWorker class . . . 133

src/PowerReader.cpp
: Implementation file for the PowerReader class . . . 133

src/Runnable.cpp
: Implementation file for the Runnable class . . . 134

src/Thread.cpp
: Implementation file for the Thread class . . . 135

src/ThroughputBenchmark.cpp
: Implementation file for the ThroughputBenchmark class . . . 135

src/Timer.cpp
: Implementation file for the Timer class . . . 136

src/ext/DelayInjectedLoadedLatencyBenchmark/delay_injected_benchmark_kernels.cpp
: Implementation file for benchmark kernel functions for the delay-injected loaded latency benchmark . . . 99

src/ext/DelayInjectedLoadedLatencyBenchmark/DelayInjectedLoadedLatencyBenchmark.cpp
: Implementation file for the DelayInjectedLatencyBenchmark class . . . 99

src/include/Benchmark.h
: Header file for the Benchmark class . . . 100

src/include/benchmark_kernels.h
Header file for benchmark kernel functions for doing the actual work we care about. :) . . . . . 101
src/include/BenchmarkManager.h
Header file for the BenchmarkManager class . . . . . 104
src/include/**build_datetime.h** . . . . . **??**
src/include/common.h
Header file for common preprocessor definitions, macros, functions, and global constants . . . 105
src/include/Configurator.h
Header file for the Configurator class and some helper data structures . . . . . 110
src/include/ExampleArg.h
Slightly-modified third-party code related to OptionParser . . . . . 111
src/include/LatencyBenchmark.h
Header file for the LatencyBenchmark class . . . . . 115
src/include/LatencyWorker.h
Header file for the LatencyWorker class . . . . . 117
src/include/LoadWorker.h
Header file for the LoadWorker class . . . . . 118
src/include/MemoryWorker.h
Header file for the MemoryWorker class . . . . . 119
src/include/MyArg.h
Extensions to third-party optionparser-related code . . . . . 119
src/include/optionparser.h
This is the only file required to use The Lean Mean C++ Option Parser. Just #include it and you're set . . . . . 121
src/include/PowerReader.h
Header file for the PowerReader class . . . . . 125
src/include/Runnable.h
Header file for the Runnable class . . . . . 126
src/include/Thread.h
Header file for the Thread class . . . . . 126
src/include/ThroughputBenchmark.h
Header file for the ThroughputBenchmark class . . . . . 127
src/include/Timer.h
Header file for the Timer class . . . . . 128
src/include/ext/DelayInjectedLoadedLatencyBenchmark/delay_injected_benchmark_kernels.h
Header file for benchmark kernel functions with integrated delays for doing the actual work we care about. :) . . . . . 112
src/include/ext/DelayInjectedLoadedLatencyBenchmark/DelayInjectedLoadedLatencyBenchmark.h
Header file for the DelayInjectedLoadedLatencyBenchmark class . . . . . 115
src/include/win/win_common_third_party.h
Header file for some third-party helper code for working with Windows APIs . . . . . 129
src/include/win/win_CPdhQuery.h
Header and implementation file for some third-party code for measuring Windows OS-exposed performance counters . . . . . 130
src/include/win/WindowsDRAMPowerReader.h
Header file for the WindowsDRAMPowerReader class . . . . . 130
src/win/win_common_third_party.cpp
Implementation file for some third-party helper code for working with Windows APIs . . . . . 137
src/win/WindowsDRAMPowerReader.cpp
Implementation file for the WindowsDRAMPowerReader class . . . . . 137

# Chapter 5

# Class Documentation

## 5.1 xmem::Parser::Action Struct Reference

Inheritance diagram for xmem::Parser::Action:

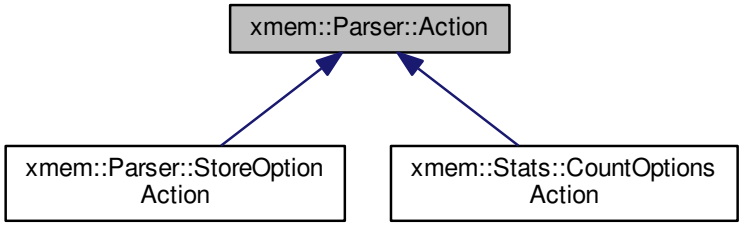


### Public Member Functions

- virtual bool perform (Option &)

  *Called by Parser::workhorse() for each Option that has been successfully parsed (including unknown options if they have a Descriptor whose Descriptor::check_arg does not return ARG_ILLEGAL.*
- virtual bool finished (int numargs, const char **args)

  *Called by Parser::workhorse() after finishing the parse.*

### 5.1.1 Member Function Documentation

#### 5.1.1.1 virtual bool xmem::Parser::Action::finished ( int *numargs*, const char ** *args* ) [inline], [virtual]

Called by Parser::workhorse() after finishing the parse.

**Parameters**

| | |
|---|---|
| *numargs* | the number of non-option arguments remaining |

| | |
|---|---|
| *args* | pointer to the first remaining non-option argument (if numargs > 0). |

**Returns**

`false` iff a fatal error has occurred.

Reimplemented in xmem::Parser::StoreOptionAction.

#### 5.1.1.2 virtual bool xmem::Parser::Action::perform ( Option & ) [inline], [virtual]

Called by Parser::workhorse() for each Option that has been successfully parsed (including unknown options if they have a Descriptor whose Descriptor::check_arg does not return ARG_ILLEGAL.

Returns `false` iff a fatal error has occured and the parse should be aborted.

Reimplemented in xmem::Parser::StoreOptionAction, and xmem::Stats::CountOptionsAction.

The documentation for this struct was generated from the following file:

- src/include/optionparser.h

## 5.2 xmem::Arg Struct Reference

Functions for checking the validity of option arguments.

`#include <optionparser.h>`

Inheritance diagram for xmem::Arg:

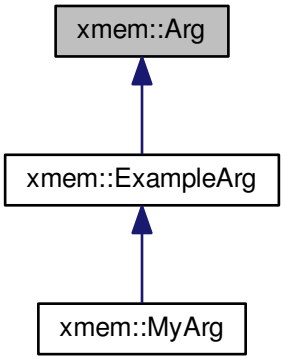


### Static Public Member Functions

- static ArgStatus None (const Option &, bool)

  *For options that don't take an argument: Returns ARG_NONE.*
- static ArgStatus Optional (const Option &option, bool)

  *Returns ARG_OK if the argument is attached and ARG_IGNORE otherwise.*

### 5.2.1 Detailed Description

Functions for checking the validity of option arguments.

The following example code can serve as starting place for writing your own more complex CheckArg functions:

```
struct Arg: public option::Arg
{
  static void printError(const char* msg1, const option::Option& opt, const char* msg2)
  {
    fprintf(stderr, "ERROR: %s", msg1);
    fwrite(opt.name, opt.namelen, 1, stderr);
    fprintf(stderr, "%s", msg2);
  }

  static option::ArgStatus Unknown(const option::Option& option, bool msg)
  {
    if (msg) printError("Unknown option '", option, "'\n");
    return option::ARG_ILLEGAL;
  }

  static option::ArgStatus Required(const option::Option& option, bool msg)
  {
    if (option.arg != 0)
      return option::ARG_OK;

    if (msg) printError("Option '", option, "' requires an argument\n");
    return option::ARG_ILLEGAL;
  }

  static option::ArgStatus NonEmpty(const option::Option& option, bool msg)
  {
    if (option.arg != 0 && option.arg[0] != 0)
      return option::ARG_OK;

    if (msg) printError("Option '", option, "' requires a non-empty argument\n");
    return option::ARG_ILLEGAL;
  }

  static option::ArgStatus Numeric(const option::Option& option, bool msg)
  {
    char* endptr = 0;
    if (option.arg != 0 && strtol(option.arg, &endptr, 10)){};
    if (endptr != option.arg && *endptr == 0)
      return option::ARG_OK;

    if (msg) printError("Option '", option, "' requires a numeric argument\n");
    return option::ARG_ILLEGAL;
  }
};
```

The documentation for this struct was generated from the following file:

- src/include/optionparser.h

## 5.3 xmem::Benchmark Class Reference

Flexible abstract class for any memory benchmark.

`#include <Benchmark.h>`

Inheritance diagram for xmem::Benchmark:

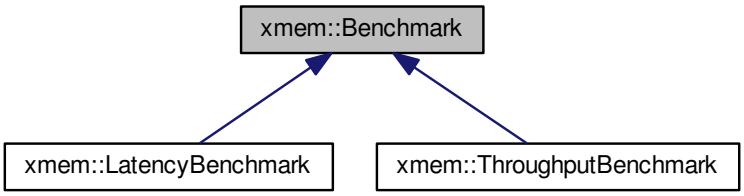


## Public Member Functions

- Benchmark (void *mem_array, size_t len, uint32_t iterations, uint32_t num_worker_threads, uint32_t mem_node, uint32_t cpu_node, pattern_mode_t pattern_mode, rw_mode_t rw_mode, chunk_size_t chunk_size, int32_t stride_size, std::vector< PowerReader * > dram_power_readers, std::string metricUnits, std::string name)

  *Constructor.*
- virtual ~Benchmark ()

  *Destructor.*
- bool run ()

  *Runs the benchmark.*
- void print_benchmark_header () const

  *Prints a header piece of information describing the benchmark to the console.*
- virtual void report_benchmark_info () const

  *Reports benchmark configuration details to the console.*
- virtual void report_results () const

  *Reports results to the console.*
- bool isValid () const

  *Checks to see that the object is in a valid state.*
- bool hasRun () const

  *Checks to see if the benchmark has run.*
- double getMetricOnIter (uint32_t iter) const

  *Extracts the metric of interest for a given iteration. Units are interpreted by the inheriting class.*
- double getAverageMetric () const

  *Gets the average benchmark metric across all iterations.*
- std::string getMetricUnits () const

  *Gets the units of the metric for this benchmark.*
- double getAverageDRAMPower (uint32_t socket_id) const

  *Gets the average DRAM power over the benchmark.*
- double getPeakDRAMPower (uint32_t socket_id) const

  *Gets the peak DRAM power over the benchmark.*
- size_t getLen () const

  *Gets the length of the memory region in bytes. This is not necessarily the "working set size" depending on multi-threading configuration.*
- uint32_t getIterations () const

  *Gets the number of iterations for this benchmark.*
- chunk_size_t getChunkSize () const

*Gets the width of memory access used in this benchmark.*

- int32_t getStrideSize () const

  *Gets the stride size for this benchmark.*

- uint32_t getCPUNode () const

  *Gets the CPU NUMA node used in this benchmark.*

- uint32_t getMemNode () const

  *Gets the memory NUMA node used in this benchmark.*

- uint32_t getNumThreads () const

  *Gets the number of worker threads used in this benchmark.*

- std::string getName () const

  *Gets the human-friendly name of this benchmark.*

- pattern_mode_t getPatternMode () const

  *Gets the pattern mode for this benchmark.*

- rw_mode_t getRWMode () const

  *Gets the read/write mode for this benchmark.*

### Protected Member Functions

- virtual bool _run_core ()=0

  *The core benchmark function.*

- bool _start_power_threads ()

  *Starts the DRAM power measurement threads.*

- bool _stop_power_threads ()

  *Stops the DRAM power measurement threads. This is a blocking call.*

### Protected Attributes

- void * _mem_array
- size_t _len
- uint32_t _iterations
- uint32_t _num_worker_threads
- uint32_t _mem_node
- uint32_t _cpu_node
- pattern_mode_t _pattern_mode
- rw_mode_t _rw_mode
- chunk_size_t _chunk_size
- int32_t _stride_size
- std::vector< PowerReader * > _dram_power_readers
- std::vector< Thread * > _dram_power_threads
- std::vector< double > _metricOnIter
- double _averageMetric
- std::string _metricUnits
- std::vector< double > _average_dram_power_socket
- std::vector< double > _peak_dram_power_socket
- std::string _name
- bool _obj_valid
- bool _hasRun
- bool _warning

### 5.3.1 Detailed Description

Flexible abstract class for any memory benchmark.

This class provides a generic interface for interacting with a benchmark. All benchmarks should be derived from this class.

### 5.3.2 Constructor & Destructor Documentation

#### 5.3.2.1 Benchmark::Benchmark ( void * *mem_array,* size_t *len,* uint32_t *iterations,* uint32_t *num_worker_threads,* uint32_t *mem_node,* uint32_t *cpu_node,* pattern_mode_t *pattern_mode,* rw_mode_t *rw_mode,* chunk_size_t *chunk_size,* int32_t *stride_size,* std::vector< **PowerReader** * > *dram_power_readers,* std::string *metricUnits,* std::string *name* )

Constructor.

**Parameters**

| | |
|---|---|
| *mem_array* | A pointer to a contiguous chunk of memory that has been allocated for benchmarking among potentially several worker threads. This should be aligned to a 256-bit boundary. |
| *len* | Length of mem_array in bytes. This must be a multiple of 4 KB and should be at least the per-thread working set size times the number of worker threads. |
| *iterations* | Number of iterations of the complete benchmark. Used to average results and provide a measure of consistency and reproducibility. |
| *num_worker_threads* | The number of worker threads to use in the benchmark. |
| *mem_node* | The logical memory NUMA node used in the benchmark. |
| *cpu_node* | The logical CPU NUMA node to use for the benchmark. |
| *pattern_mode* | This indicates the general type of access pattern used, e.g. sequential or random. |
| *rw_mode* | This indicates the general type of read/write mix used, e.g. pure reads or pure writes. |
| *chunk_size* | Size of an individual memory access for load-generating worker threads. |
| *stride_size* | For sequential access patterns, this is the address distance between successive accesses, counted in chunks. Negative values indicate a reversed access pattern. A stride of +/-1 is purely sequential. |
| *dram_power_readers* | A group of PowerReader objects for measuring DRAM power. |
| *name* | The name of the benchmark to use when reporting to console. |

### 5.3.3 Member Function Documentation

#### 5.3.3.1 virtual bool xmem::Benchmark::_run_core ( ) [protected], [pure virtual]

The core benchmark function.

**Returns**

True on success.

Implemented in xmem::LatencyBenchmark, and xmem::ThroughputBenchmark.

#### 5.3.3.2 bool Benchmark::_start_power_threads ( ) [protected]

Starts the DRAM power measurement threads.

**Returns**

True on success.

#### 5.3.3.3 bool Benchmark::_stop_power_threads ( ) [protected]

Stops the DRAM power measurement threads. This is a blocking call.

**Returns**

True on success.

#### 5.3.3.4 double Benchmark::getAverageDRAMPower ( uint32_t *socket_id* ) const

Gets the average DRAM power over the benchmark.

**Returns**

The average DRAM power for a given socket in watts, or 0 if the data does not exist (power was unable to be collected or the benchmark has not run).

#### 5.3.3.5 double Benchmark::getAverageMetric ( ) const

Gets the average benchmark metric across all iterations.

**Returns**

The average metric.

#### 5.3.3.6 chunk_size_t Benchmark::getChunkSize ( ) const

Gets the width of memory access used in this benchmark.

**Returns**

The chunk size for this benchmark.

#### 5.3.3.7 uint32_t Benchmark::getCPUNode ( ) const

Gets the CPU NUMA node used in this benchmark.

**Returns**

The NUMA CPU node used in this benchmark.

#### 5.3.3.8 uint32_t Benchmark::getIterations ( ) const

Gets the number of iterations for this benchmark.

**Returns**

The number of iterations for this benchmark.

#### 5.3.3.9 size_t Benchmark::getLen ( ) const

Gets the length of the memory region in bytes. This is not necessarily the "working set size" depending on multi-threading configuration.

**Returns**

Length of the memory region in bytes.

#### 5.3.3.10 uint32_t Benchmark::getMemNode ( ) const

Gets the memory NUMA node used in this benchmark.

**Returns**

The NUMA memory node used in this benchmark.

#### 5.3.3.11 double Benchmark::getMetricOnIter ( uint32_t *iter* ) const

Extracts the metric of interest for a given iteration. Units are interpreted by the inheriting class.

**Parameters**

| | |
|---|---|
| *iter* | Iteration to extract. |

**Returns**

The metric on the iteration specified by the input.

#### 5.3.3.12 std::string Benchmark::getMetricUnits ( ) const

Gets the units of the metric for this benchmark.

**Returns**

A string representing the units for printing to console and file.

#### 5.3.3.13 std::string Benchmark::getName ( ) const

Gets the human-friendly name of this benchmark.

**Returns**

The benchmark test name.

#### 5.3.3.14 uint32_t Benchmark::getNumThreads ( ) const

Gets the number of worker threads used in this benchmark.

**Returns**

The number of worker threads used in this benchmark.

#### 5.3.3.15 pattern_mode_t Benchmark::getPatternMode ( ) const

Gets the pattern mode for this benchmark.

**Returns**

The pattern mode enumerator.

#### 5.3.3.16 double Benchmark::getPeakDRAMPower ( uint32_t *socket_id* ) const

Gets the peak DRAM power over the benchmark.

**Returns**

The peak DRAM power for a given socket in watts, or 0 if the data does not exist (power was unable to be collected or the benchmark has not run).

#### 5.3.3.17 rw_mode_t Benchmark::getRWMode ( ) const

Gets the read/write mode for this benchmark.

**Returns**

The read/write mix mode.

#### 5.3.3.18 int32_t Benchmark::getStrideSize ( ) const

Gets the stride size for this benchmark.

**Returns**

The stride size in chunks.

#### 5.3.3.19 bool Benchmark::hasRun ( ) const

Checks to see if the benchmark has run.

**Returns**

True if run() has already completed successfully.

#### 5.3.3.20 bool Benchmark::isValid ( ) const

Checks to see that the object is in a valid state.

**Returns**

True if the object was constructed correctly and can be used.

#### 5.3.3.21 bool Benchmark::run ( )

Runs the benchmark.

**Returns**

True on benchmark success

Here is the call graph for this function:

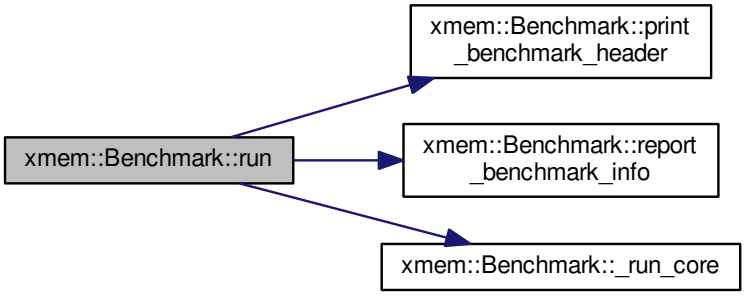


### 5.3.4 Member Data Documentation

#### 5.3.4.1 std::vector<double> xmem::Benchmark::_average_dram_power_socket [protected]

The average DRAM power in this benchmark, per socket.

#### 5.3.4.2 double xmem::Benchmark::_averageMetric [protected]

Average metric over all iterations. Unit-less because any benchmark can set this metric as needed. It is up to the descendant class to interpret units.

#### 5.3.4.3 chunk_size_t xmem::Benchmark::_chunk_size [protected]

Chunk size of memory accesses in this benchmark.

#### 5.3.4.4 uint32_t xmem::Benchmark::_cpu_node [protected]

The CPU NUMA node used in this benchmark.

#### 5.3.4.5 std::vector<PowerReader*> xmem::Benchmark::_dram_power_readers [protected]

The power reading objects for measuring DRAM power on a per-socket basis during the benchmark.

#### 5.3.4.6 std::vector<Thread*> xmem::Benchmark::_dram_power_threads [protected]

The power reading threads for measuring DRAM power on a per-socket basis during the benchmark. These work with the DRAM power readers. Although they are worker threads, they are not counted as the "official" benchmarking worker threads.

#### 5.3.4.7 bool xmem::Benchmark::_hasRun [protected]

Indicates whether the benchmark has run.

#### 5.3.4.8 uint32_t xmem::Benchmark::_iterations [protected]

Number of iterations used in this benchmark.

#### 5.3.4.9 size_t xmem::Benchmark::_len [protected]

Length of the memory region in bytes. This is not the working set size per thread!

#### 5.3.4.10 void* xmem::Benchmark::_mem_array [protected]

Pointer to the memory region to use in this benchmark.

#### 5.3.4.11 uint32_t xmem::Benchmark::_mem_node [protected]

The memory NUMA node used in this benchmark.

#### 5.3.4.12 std::vector<double> xmem::Benchmark::_metricOnIter [protected]

Metrics for each iteration of the benchmark. Unit-less because any benchmark can set this metric as needed. It is up to the descendant class to interpret units.

#### 5.3.4.13 std::string xmem::Benchmark::_metricUnits [protected]

String representing the units of measurement for the metric.

#### 5.3.4.14 std::string xmem::Benchmark::_name [protected]

Name of this benchmark.

#### 5.3.4.15 uint32_t xmem::Benchmark::_num_worker_threads [protected]

The number of worker threads used in this benchmark.

#### 5.3.4.16 bool xmem::Benchmark::_obj_valid [protected]

Indicates whether this benchmark object is valid.

#### 5.3.4.17 pattern_mode_t xmem::Benchmark::_pattern_mode [protected]

Access pattern mode.

#### 5.3.4.18 std::vector<double> xmem::Benchmark::_peak_dram_power_socket [protected]

The peak DRAM power in this benchmark, per socket.

#### 5.3.4.19 rw_mode_t xmem::Benchmark::_rw_mode [protected]

Read/write mode.

#### 5.3.4.20 int32_t xmem::Benchmark::_stride_size [protected]

Stride size in chunks for sequential pattern mode only.

#### 5.3.4.21 bool xmem::Benchmark::_warning [protected]

Indicates whether the benchmarks results might be clearly questionable/inaccurate/incorrect due to a variety of factors.

The documentation for this class was generated from the following files:

- src/include/Benchmark.h
- src/Benchmark.cpp

## 5.4 xmem::BenchmarkManager Class Reference

Manages running all benchmarks at a high level.

`#include <BenchmarkManager.h>`

### Public Member Functions

- BenchmarkManager (Configurator &config)

  *Constructor.*
- ~BenchmarkManager ()

  *Destructor.*
- bool runAll ()

  *Runs all benchmark configurations (does not include extensions).*
- bool runThroughputBenchmarks ()

  *Runs the throughput benchmarks.*
- bool runLatencyBenchmarks ()

  *Runs the latency benchmark.*

### 5.4.1 Detailed Description

Manages running all benchmarks at a high level.

### 5.4.2 Constructor & Destructor Documentation

#### 5.4.2.1 BenchmarkManager::BenchmarkManager ( Configurator & *config* )

Constructor.

**Parameters**

| | |
|---|---|
| *config* | The configuration object containing run-time options for this X-Mem execution instance. |

Here is the call graph for this function:

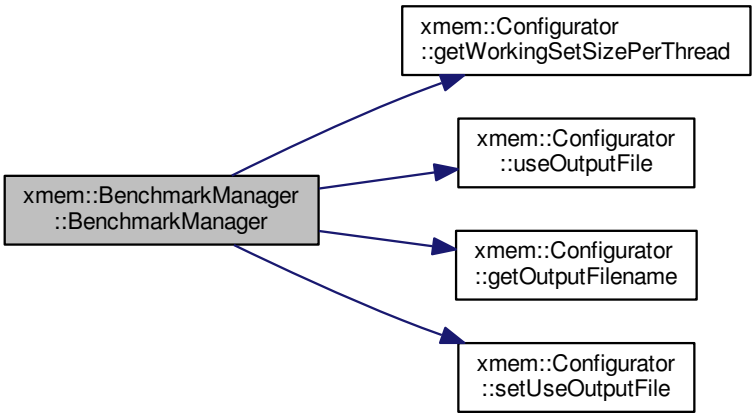


### 5.4.3 Member Function Documentation

#### 5.4.3.1 bool BenchmarkManager::runAll ( )

Runs all benchmark configurations (does not include extensions).

**Returns**

True on success.

Here is the call graph for this function:

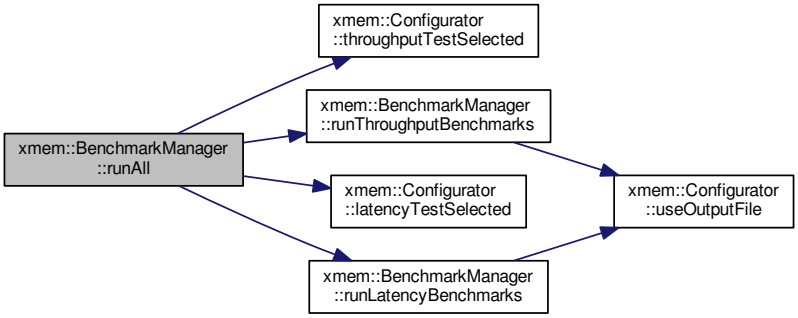


#### 5.4.3.2 bool BenchmarkManager::runLatencyBenchmarks ( )

Runs the latency benchmark.

**Returns**

True on benchmarking success.

Here is the call graph for this function:

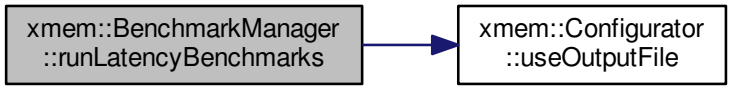


#### 5.4.3.3 bool BenchmarkManager::runThroughputBenchmarks ( )

Runs the throughput benchmarks.

**Returns**

True on benchmarking success.

Here is the call graph for this function:

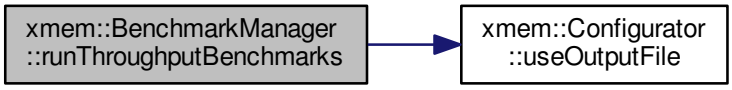


The documentation for this class was generated from the following files:

- src/include/BenchmarkManager.h
- src/BenchmarkManager.cpp

## 5.5 xmem::Configurator Class Reference

Handles all user input interpretation and generates the necessary flags for running benchmarks.

`#include <Configurator.h>`

### Public Member Functions

- Configurator ()

  *Default constructor. A default configuration is set. You will want to run configureFromInput() most likely.*
- int32_t configureFromInput (int argc, char *argv[])

*Configures the tool based on user's command-line inputs.*

- bool extensionsEnabled () const

  *Determines whether user extensions are enabled.*
- bool runExtDelayInjectedLoadedLatencyBenchmark () const

  *If included at compile-time, determines whether the delay-injected loaded latency benchmark extension should be run.*
- bool latencyTestSelected () const

  *Indicates if the latency test has been selected.*
- bool throughputTestSelected () const

  *Indicates if the throughput test has been selected.*
- size_t getWorkingSetSizePerThread () const

  *Gets the working set size in bytes for each worker thread, if applicable.*
- bool useChunk32b () const

  *Determines if chunk size of 32 bits should be used in relevant benchmarks.*
- bool isNUMAEnabled () const

  *Determines if the benchmarks should test for all CPU/memory NUMA combinations.*
- uint32_t getIterationsPerTest () const

  *Gets the number of iterations that should be run of each benchmark.*
- bool useRandomAccessPattern () const

  *Determines if throughput benchmarks should use a random access pattern.*
- bool useSequentialAccessPattern () const

  *Determines if throughput benchmarks should use a sequential access pattern.*
- uint32_t getNumWorkerThreads () const

  *Gets the number of worker threads to use.*
- uint32_t getStartingTestIndex () const

  *Gets the numerical index of the first benchmark for CSV output purposes.*
- std::string getOutputFilename () const

  *Gets the output filename to use, if applicable.*
- bool useOutputFile () const

  *Determines whether to generate an output CSV file.*
- void setUseOutputFile (bool use)

  *Changes whether an output file should be used.*
- bool verboseMode () const

  *Determines whether X-Mem is in verbose mode.*
- bool useLargePages () const

  *Determines whether X-Mem should use large pages.*
- bool useReads () const

  *Determines whether reads should be used in throughput benchmarks.*
- bool useWrites () const

  *Determines whether writes should be used in throughput benchmarks.*
- bool useStrideP1 () const

  *Determines if a stride of +1 should be used in relevant benchmarks.*
- bool useStrideN1 () const

  *Determines if a stride of -1 should be used in relevant benchmarks.*
- bool useStrideP2 () const

  *Determines if a stride of +2 should be used in relevant benchmarks.*
- bool useStrideN2 () const

  *Determines if a stride of -2 should be used in relevant benchmarks.*
- bool useStrideP4 () const

  *Determines if a stride of +4 should be used in relevant benchmarks.*
- bool useStrideN4 () const

*Determines if a stride of -4 should be used in relevant benchmarks.*

- bool useStrideP8 () const

  *Determines if a stride of +8 should be used in relevant benchmarks.*

- bool useStrideN8 () const

  *Determines if a stride of -8 should be used in relevant benchmarks.*

- bool useStrideP16 () const

  *Determines if a stride of +16 should be used in relevant benchmarks.*

- bool useStrideN16 () const

  *Determines if a stride of -16 should be used in relevant benchmarks.*

### 5.5.1 Detailed Description

Handles all user input interpretation and generates the necessary flags for running benchmarks.

### 5.5.2 Member Function Documentation

#### 5.5.2.1 int32_t Configurator::configureFromInput ( int *argc*, char * *argv[]* )

Configures the tool based on user's command-line inputs.

**Parameters**

| | |
|---|---|
| *argc* | The argc from main(). |
| *argv* | The argv from main(). |

**Returns**

0 on success.

Here is the call graph for this function:

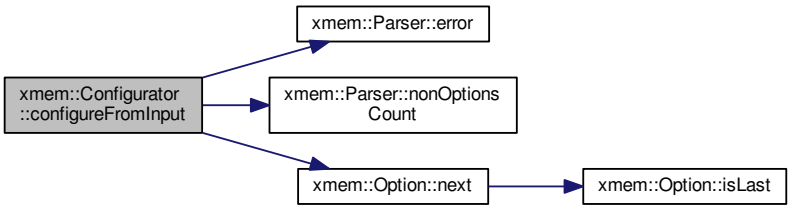


#### 5.5.2.2 bool xmem::Configurator::extensionsEnabled ( ) const [inline]

Determines whether user extensions are enabled.

**Returns**

True if extensions are enabled.

#### 5.5.2.3 uint32_t xmem::Configurator::getIterationsPerTest ( ) const [inline]

Gets the number of iterations that should be run of each benchmark.

**Returns**

The iterations for each test.

#### 5.5.2.4 uint32_t xmem::Configurator::getNumWorkerThreads ( ) const [inline]

Gets the number of worker threads to use.

**Returns**

The number of worker threads.

#### 5.5.2.5 std::string xmem::Configurator::getOutputFilename ( ) const [inline]

Gets the output filename to use, if applicable.

**Returns**

The output filename to use if useOutputFile() returns true. Otherwise return value is "".

#### 5.5.2.6 uint32_t xmem::Configurator::getStartingTestIndex ( ) const [inline]

Gets the numerical index of the first benchmark for CSV output purposes.

**Returns**

The starting benchmark index.

#### 5.5.2.7 size_t xmem::Configurator::getWorkingSetSizePerThread ( ) const [inline]

Gets the working set size in bytes for each worker thread, if applicable.

**Returns**

The working set size in bytes.

#### 5.5.2.8 bool xmem::Configurator::isNUMAEnabled ( ) const [inline]

Determines if the benchmarks should test for all CPU/memory NUMA combinations.

**Returns**

True if all NUMA nodes should be tested.

#### 5.5.2.9 bool xmem::Configurator::latencyTestSelected ( ) const [inline]

Indicates if the latency test has been selected.

**Returns**

True if the latency test has been selected to run.

#### 5.5.2.10 bool xmem::Configurator::runExtDelayInjectedLoadedLatencyBenchmark ( ) const [inline]

If included at compile-time, determines whether the delay-injected loaded latency benchmark extension should be run.

**Returns**

True if it should be run.

#### 5.5.2.11 void xmem::Configurator::setUseOutputFile ( bool *use* ) [inline]

Changes whether an output file should be used.

**Parameters**

| | |
|---|---|
| *use* | If true, then use the output file. |

#### 5.5.2.12 bool xmem::Configurator::throughputTestSelected ( ) const [inline]

Indicates if the throughput test has been selected.

**Returns**

True if the throughput test has been selected to run.

#### 5.5.2.13 bool xmem::Configurator::useChunk32b ( ) const [inline]

Determines if chunk size of 32 bits should be used in relevant benchmarks.

**Returns**

True if 32-bit chunks should be used.

#### 5.5.2.14 bool xmem::Configurator::useLargePages ( ) const [inline]

Determines whether X-Mem should use large pages.

**Parameters**

| | |
|---|---|
| *True* | if large pages should be used. |

#### 5.5.2.15 bool xmem::Configurator::useOutputFile ( ) const [inline]

Determines whether to generate an output CSV file.

**Returns**

True if an output file should be used.

#### 5.5.2.16 bool xmem::Configurator::useRandomAccessPattern ( ) const [inline]

Determines if throughput benchmarks should use a random access pattern.

**Returns**

True if random access should be used.

#### 5.5.2.17 bool xmem::Configurator::useReads ( ) const [inline]

Determines whether reads should be used in throughput benchmarks.

**Returns**

True if reads should be used.

#### 5.5.2.18 bool xmem::Configurator::useSequentialAccessPattern ( ) const [inline]

Determines if throughput benchmarks should use a sequential access pattern.

**Returns**

True if sequential access should be used.

#### 5.5.2.19 bool xmem::Configurator::useStrideN1 ( ) const [inline]

Determines if a stride of -1 should be used in relevant benchmarks.

**Returns**

True if a stride of -1 should be used.

#### 5.5.2.20 bool xmem::Configurator::useStrideN16 ( ) const [inline]

Determines if a stride of -16 should be used in relevant benchmarks.

**Returns**

True if a stride of -16 should be used.

#### 5.5.2.21 bool xmem::Configurator::useStrideN2 ( ) const [inline]

Determines if a stride of -2 should be used in relevant benchmarks.

**Returns**

True if a stride of -2 should be used.

#### 5.5.2.22 bool xmem::Configurator::useStrideN4 ( ) const [inline]

Determines if a stride of -4 should be used in relevant benchmarks.

**Returns**

True if a stride of -4 should be used.

#### 5.5.2.23 bool xmem::Configurator::useStrideN8 ( ) const [inline]

Determines if a stride of -8 should be used in relevant benchmarks.

**Returns**

True if a stride of -8 should be used.

#### 5.5.2.24 bool xmem::Configurator::useStrideP1 ( ) const [inline]

Determines if a stride of +1 should be used in relevant benchmarks.

**Returns**

True if a stride of +1 should be used.

#### 5.5.2.25 bool xmem::Configurator::useStrideP16 ( ) const [inline]

Determines if a stride of +16 should be used in relevant benchmarks.

**Returns**

True if a stride of +16 should be used.

#### 5.5.2.26 bool xmem::Configurator::useStrideP2 ( ) const [inline]

Determines if a stride of +2 should be used in relevant benchmarks.

**Returns**

True if a stride of +2 should be used.

#### 5.5.2.27 bool xmem::Configurator::useStrideP4 ( ) const [inline]

Determines if a stride of +4 should be used in relevant benchmarks.

**Returns**

True if a stride of +4 should be used.

#### 5.5.2.28 bool xmem::Configurator::useStrideP8 ( ) const [inline]

Determines if a stride of +8 should be used in relevant benchmarks.

**Returns**

True if a stride of +8 should be used.

#### 5.5.2.29 bool xmem::Configurator::useWrites ( ) const [inline]

Determines whether writes should be used in throughput benchmarks.

**Returns**

True if writes should be used.

#### 5.5.2.30 bool xmem::Configurator::verboseMode ( ) const [inline]

Determines whether X-Mem is in verbose mode.

**Returns**

True if verbose mode is enabled.

The documentation for this class was generated from the following files:

- src/include/Configurator.h
- src/Configurator.cpp

## 5.6 xmem::Stats::CountOptionsAction Class Reference

Inheritance diagram for xmem::Stats::CountOptionsAction:

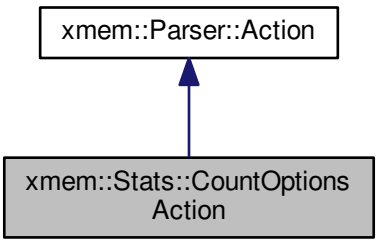


Collaboration diagram for xmem::Stats::CountOptionsAction:

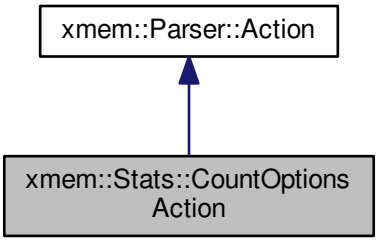


### Public Member Functions

- CountOptionsAction (unsigned *buffer_max_)
- bool perform (Option &)

  *Called by Parser::workhorse() for each Option that has been successfully parsed (including unknown options if they have a Descriptor whose Descriptor::check_arg does not return ARG_ILLEGAL.*

### 5.6.1 Constructor & Destructor Documentation

#### 5.6.1.1 xmem::Stats::CountOptionsAction::CountOptionsAction ( unsigned * *buffer_max_* ) [inline]

Creates a new CountOptionsAction that will increase `*buffer_max_` for each parsed Option.

### 5.6.2 Member Function Documentation

#### 5.6.2.1 bool xmem::Stats::CountOptionsAction::perform ( Option & ) [inline], [virtual]

Called by Parser::workhorse() for each Option that has been successfully parsed (including unknown options if they have a Descriptor whose Descriptor::check_arg does not return ARG_ILLEGAL.

Returns `false` iff a fatal error has occured and the parse should be aborted.

Reimplemented from xmem::Parser::Action.

The documentation for this class was generated from the following file:

- src/include/optionparser.h

## 5.7 xmem::Descriptor Struct Reference

Describes an option, its help text (usage) and how it should be parsed.

`#include <optionparser.h>`

### Public Attributes

- const unsigned index

  *Index of this option's linked list in the array filled in by the parser.*
- const int type

  *Used to distinguish between options with the same index. See index for details.*
- const char *const shortopt

  *Each char in this string will be accepted as a short option character.*
- const char *const longopt

  *The long option name (without the leading - ).*
- const CheckArg check_arg

  *For each option that matches shortopt or longopt this function will be called to check a potential argument to the option.*
- const char * help

  *The usage text associated with the options in this Descriptor.*

### 5.7.1 Detailed Description

Describes an option, its help text (usage) and how it should be parsed.

The main input when constructing an option::Parser is an array of Descriptors.

**Example:**

```
enum OptionIndex {CREATE, ...};
enum OptionType {DISABLE, ENABLE, OTHER};

const option::Descriptor usage[] = {
  { CREATE,                                             // index
```

```
        OTHER,                                                  // type
        "c",                                                    // shortopt
        "create",                                               // longopt
        Arg::None,                                              // check_arg
        "--create  Tells the program to create something." // help
      }
     , ...
    };
```

### 5.7.2 Member Data Documentation

#### 5.7.2.1 const CheckArg xmem::Descriptor::check_arg

For each option that matches shortopt or longopt this function will be called to check a potential argument to the option.

This function will be called even if there is no potential argument. In that case it will be passed `NULL` as `arg` parameter. Do not confuse this with the empty string.

See CheckArg for more information.

#### 5.7.2.2 const char* xmem::Descriptor::help

The usage text associated with the options in this Descriptor.

You can use option::printUsage() to format your usage message based on the `help` texts. You can use dummy Descriptors where shortopt and longopt are both the empty string to add text to the usage that is not related to a specific option.

See option::printUsage() for special formatting characters you can use in `help` to get a column layout.

**Attention**

Must be UTF-8-encoded. If your compiler supports C++11 you can use the "u8" prefix to make sure string literals are properly encoded.

#### 5.7.2.3 const unsigned xmem::Descriptor::index

Index of this option's linked list in the array filled in by the parser.

Command line options whose Descriptors have the same index will end up in the same linked list in the order in which they appear on the command line. If you have multiple long option aliases that refer to the same option, give their descriptors the same `index`.

If you have options that mean exactly opposite things (e.g. `-enable-foo` and `-disable-foo` ), you should also give them the same `index`, but distinguish them through different values for type. That way they end up in the same list and you can just take the last element of the list and use its type. This way you get the usual behaviour where switches later on the command line override earlier ones without having to code it manually.

**Tip:**

Use an enum rather than plain ints for better readability, as shown in the example at Descriptor.

#### 5.7.2.4 const char* const xmem::Descriptor::longopt

The long option name (without the leading - ).

If this Descriptor should not have a long option name, use the empty string "". NULL is not permitted here!

While shortopt allows multiple short option characters, each Descriptor can have only a single long option name. If you have multiple long option names referring to the same option use separate Descriptors that have the same index and type. You may repeat short option characters in such an alias Descriptor but there's no need to.

**Dummy Descriptors:**

You can use dummy Descriptors with an empty string for both shortopt and longopt to add text to the usage that is not related to a specific option. See help. The first dummy Descriptor will be used for unknown options (see below).

**Unknown Option Descriptor:**

The first dummy Descriptor in the list of Descriptors, whose shortopt and longopt are both the empty string, will be used as the Descriptor for unknown options. An unknown option is a string in the argument vector that is not a lone minus `'-'` but starts with a minus character and does not match any Descriptor's shortopt or longopt.
Note that the dummy descriptor's check_arg function *will* be called and its return value will be evaluated as usual. I.e. if it returns ARG_ILLEGAL the parsing will be aborted with `Parser::error()==true`.
if `check_arg` does not return ARG_ILLEGAL the descriptor's index *will* be used to pick the linked list into which to put the unknown option.
If there is no dummy descriptor, unknown options will be dropped silently.

#### 5.7.2.5 const char* const xmem::Descriptor::shortopt

Each char in this string will be accepted as a short option character.

The string must not include the minus character `'-'` or you'll get undefined behaviour.

If this Descriptor should not have short option characters, use the empty string "". NULL is not permitted here!

See longopt for more information.

#### 5.7.2.6 const int xmem::Descriptor::type

Used to distinguish between options with the same index. See index for details.

It is recommended that you use an enum rather than a plain int to make your code more readable.

The documentation for this struct was generated from the following file:

- src/include/optionparser.h

## 5.8 xmem::ExampleArg Class Reference

Inheritance diagram for xmem::ExampleArg:

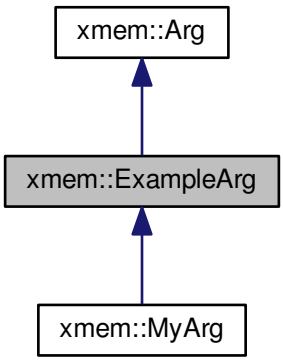


Collaboration diagram for xmem::ExampleArg:

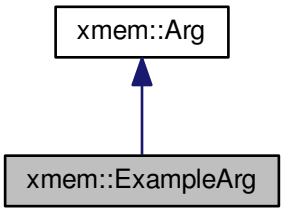


### Static Public Member Functions

- static void **printError** (const char ∗msg1, const Option &opt, const char ∗msg2)
- static ArgStatus **Unknown** (const Option &option, bool msg)
- static ArgStatus **Required** (const Option &option, bool msg)
- static ArgStatus **NonEmpty** (const Option &option, bool msg)

The documentation for this class was generated from the following file:

- src/include/ExampleArg.h

## 5.9 xmem::PrintUsageImplementation::FunctionWriter< Function > Struct Template Reference

Inheritance diagram for xmem::PrintUsageImplementation::FunctionWriter< Function >:

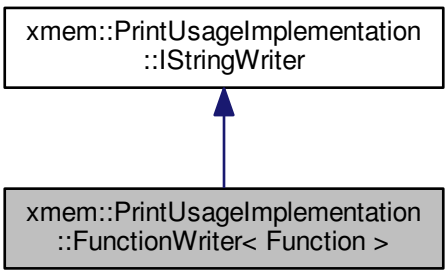


Collaboration diagram for xmem::PrintUsageImplementation::FunctionWriter< Function >:

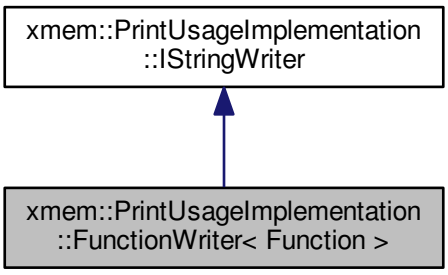


### Public Member Functions

- virtual void operator() (const char *str, int size)

  *Writes the given number of chars beginning at the given pointer somewhere.*
- **FunctionWriter** (Function *w)

### Public Attributes

- Function * **write**

The documentation for this struct was generated from the following file:

- src/include/optionparser.h

## 5.10 xmem::PrintUsageImplementation::IStringWriter Struct Reference

Inheritance diagram for xmem::PrintUsageImplementation::IStringWriter:

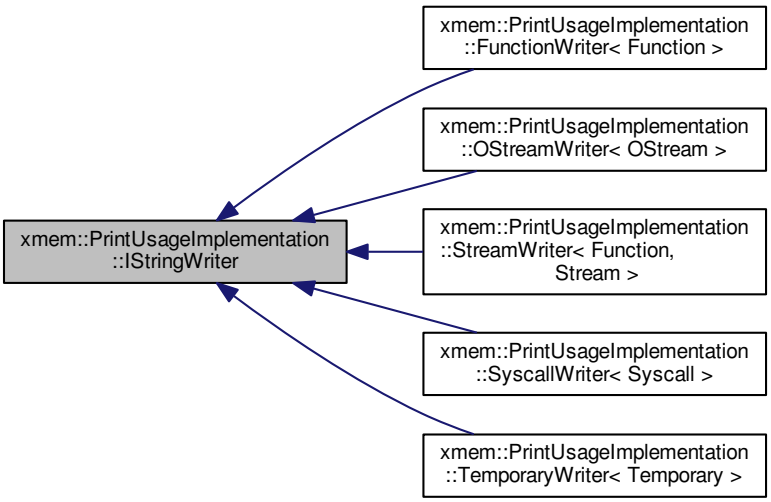


### Public Member Functions

- virtual void operator() (const char *, int)

  *Writes the given number of chars beginning at the given pointer somewhere.*

The documentation for this struct was generated from the following file:

- src/include/optionparser.h

## 5.11 xmem::LatencyBenchmark Class Reference

A type of benchmark that measures memory latency via random pointer chasing. Loading may be provided with separate threads which access memory as quickly as possible using given access patterns.

`#include <LatencyBenchmark.h>`

Inheritance diagram for xmem::LatencyBenchmark:

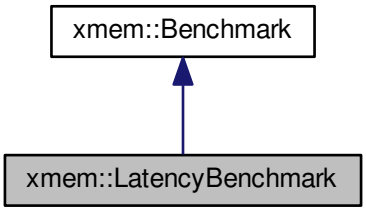


Collaboration diagram for xmem::LatencyBenchmark:

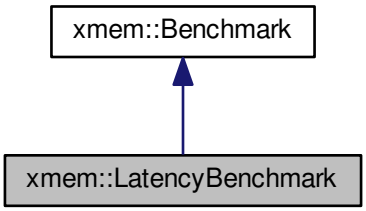


## Public Member Functions

- LatencyBenchmark (void *mem_array, size_t len, uint32_t iterations, uint32_t num_worker_threads, uint32_t mem_node, uint32_t cpu_node, pattern_mode_t pattern_mode, rw_mode_t rw_mode, chunk_size_t chunk_size, int32_t stride_size, std::vector< PowerReader * > dram_power_readers, std::string name)

  *Constructor. Parameters are passed directly to the Benchmark constructor. See Benchmark class documentation for parameter semantics.*
- virtual ~LatencyBenchmark ()

  *Destructor.*
- double getLoadMetricOnIter (uint32_t iter) const

  *Get the average load throughput in MB/sec that was imposed on the latency measurement during the given iteration.*
- double getAvgLoadMetric () const

  *Get the overall average load throughput in MB/sec that was imposed on the latency measurement.*
- virtual void report_benchmark_info () const

  *Reports benchmark configuration details to the console.*
- virtual void report_results () const

  *Reports results to the console.*

## Protected Member Functions

- virtual bool _run_core ()

  *The core benchmark function.*

### Protected Attributes

- std::vector< double > _loadMetricOnIter
- double _averageLoadMetric

## 5.11.1 Detailed Description

A type of benchmark that measures memory latency via random pointer chasing. Loading may be provided with separate threads which access memory as quickly as possible using given access patterns.

## 5.11.2 Member Function Documentation

#### 5.11.2.1 bool LatencyBenchmark::_run_core ( ) [protected], [virtual]

The core benchmark function.

**Returns**

True on success.

Implements xmem::Benchmark.

Here is the call graph for this function:

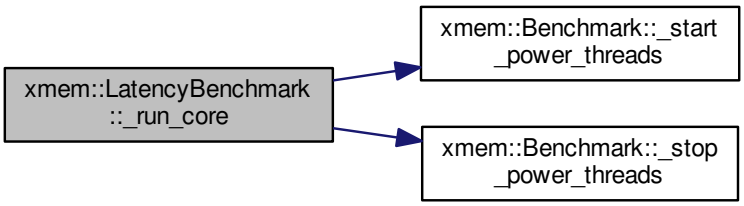


#### 5.11.2.2 double LatencyBenchmark::getAvgLoadMetric ( ) const

Get the overall average load throughput in MB/sec that was imposed on the latency measurement.

**Returns**

The average throughput in MB/sec.

#### 5.11.2.3 double LatencyBenchmark::getLoadMetricOnIter ( uint32_t *iter* ) const

Get the average load throughput in MB/sec that was imposed on the latency measurement during the given iteration.

iter The iteration of interest.

**Returns**

The average throughput in MB/sec.

### 5.11.3 Member Data Documentation

#### 5.11.3.1 double xmem::LatencyBenchmark::_averageLoadMetric [protected]

The average load throughput in MB/sec that was imposed on the latency measurement.

#### 5.11.3.2 std::vector<double> xmem::LatencyBenchmark::_loadMetricOnIter [protected]

Load metrics for each iteration of the benchmark. This is in MB/s.

The documentation for this class was generated from the following files:

- src/include/LatencyBenchmark.h
- src/LatencyBenchmark.cpp

## 5.12 xmem::LatencyWorker Class Reference

Multithreading-friendly class to do memory loading.

`#include <LatencyWorker.h>`

Inheritance diagram for xmem::LatencyWorker:

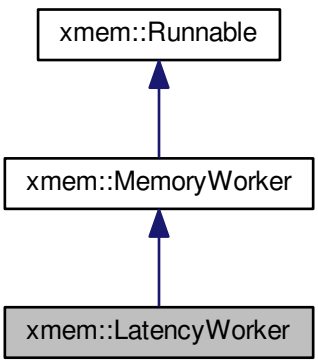

Collaboration diagram for xmem::LatencyWorker:

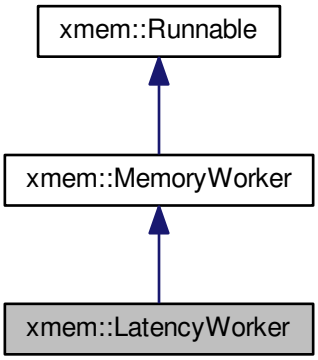


## Public Member Functions

- LatencyWorker (void ∗mem_array, size_t len, RandomFunction kernel_fptr, RandomFunction kernel_dummy_fptr, int32_t cpu_affinity)

  *Constructor for sequential-access patterns.*
- virtual ∼LatencyWorker ()

  *Destructor.*
- virtual void run ()

  *Thread-safe worker method.*

## Additional Inherited Members

### 5.12.1 Detailed Description

Multithreading-friendly class to do memory loading.

### 5.12.2 Constructor & Destructor Documentation

#### 5.12.2.1 LatencyWorker::LatencyWorker ( void ∗ *mem_array,* size_t *len,* RandomFunction *kernel_fptr,* RandomFunction *kernel_dummy_fptr,* int32_t *cpu_affinity* )

Constructor for sequential-access patterns.

**Parameters**

| | |
|---|---|
| *mem_array* | Pointer to the memory region to use by this worker. |
| *len* | Length of the memory region to use by this worker. |
| *kernel_fptr* | Pointer to the sequential core benchmark kernel to use. |
| *kernel_dummy_fptr* | Pointer to the sequential dummy version of the core benchmark kernel to use. |

| | |
|---|---|
| *cpu_affinity* | Logical CPU identifier to lock this worker's thread to. |

The documentation for this class was generated from the following files:

- src/include/LatencyWorker.h
- src/LatencyWorker.cpp

## 5.13 xmem::PrintUsageImplementation::LinePartIterator Class Reference

### Public Member Functions

- LinePartIterator (const Descriptor usage[])

  *Creates an iterator for* `usage`*.*
- bool nextTable ()

  *Moves iteration to the next table (if any). Has to be called once on a new LinePartIterator to move to the 1st table.*
- void restartTable ()

  *Reset iteration to the beginning of the current table.*
- bool nextRow ()

  *Moves iteration to the next row (if any). Has to be called once after each call to nextTable() to move to the 1st row of the table.*
- void restartRow ()

  *Reset iteration to the beginning of the current row.*
- bool next ()

  *Moves iteration to the next part (if any). Has to be called once after each call to nextRow() to move to the 1st part of the row.*
- int column ()

  *Returns the index (counting from 0) of the column in which the part pointed to by data() is located.*
- int line ()

  *Returns the index (counting from 0) of the line within the current column this part belongs to.*
- int length ()

  *Returns the length of the part pointed to by data() in raw chars (not UTF-8 characters).*
- int screenLength ()

  *Returns the width in screen columns of the part pointed to by data(). Takes multi-byte UTF-8 sequences and wide characters into account.*
- const char * data ()

  *Returns the current part of the iteration.*

### 5.13.1 Member Function Documentation

#### 5.13.1.1 bool xmem::PrintUsageImplementation::LinePartIterator::next ( ) [inline]

Moves iteration to the next part (if any). Has to be called once after each call to nextRow() to move to the 1st part of the row.

**Return values**

| | |
|---|---|
| *false* | if moving to next part failed because no further part exists. |

See LinePartIterator for details about the iteration.

#### 5.13.1.2 bool xmem::PrintUsageImplementation::LinePartIterator::nextRow ( ) [inline]

Moves iteration to the next row (if any). Has to be called once after each call to nextTable() to move to the 1st row of the table.

**Return values**

| | |
|---|---|
| *false* | if moving to next row failed because no further row exists. |

Here is the call graph for this function:

#### 5.13.1.3 bool xmem::PrintUsageImplementation::LinePartIterator::nextTable ( ) [inline]

Moves iteration to the next table (if any). Has to be called once on a new LinePartIterator to move to the 1st table.

**Return values**

| | |
|---|---|
| *false* | if moving to next table failed because no further table exists. |

Here is the call graph for this function:

The documentation for this class was generated from the following file:

- src/include/optionparser.h

## 5.14 xmem::PrintUsageImplementation::LineWrapper Class Reference

### Public Member Functions

- void flush (IStringWriter &write)

  *Writes out all remaining data from the LineWrapper using `write`. Unlike process() this method indents all lines including the first and will output a \n at the end (but only if something has been written).*
- void process (IStringWriter &write, const char *data, int len)

  *Process, wrap and output the next piece of data.*
- LineWrapper (int x1, int x2)

  *Constructs a LineWrapper that wraps its output to fit into screen columns `x1` (incl.) to `x2` (excl.).*

### 5.14.1 Constructor & Destructor Documentation

#### 5.14.1.1 xmem::PrintUsageImplementation::LineWrapper::LineWrapper ( int *x1*, int *x2* ) [inline]

Constructs a LineWrapper that wraps its output to fit into screen columns `x1` (incl.) to `x2` (excl.).

`x1` gives the indentation LineWrapper uses if it needs to indent.

### 5.14.2 Member Function Documentation

#### 5.14.2.1 void xmem::PrintUsageImplementation::LineWrapper::process ( IStringWriter & *write*, const char * *data*, int *len* ) [inline]

Process, wrap and output the next piece of data.

process() will output at least one line of output. This is not necessarily the `data` passed in. It may be data queued from a prior call to process(). If the internal buffer is full, more than 1 line will be output.

process() assumes that the a proper amount of indentation has already been output. It won't write any further indentation before the 1st line. If more than 1 line is written due to buffer constraints, the lines following the first will be indented by this method, though.

No \n is written by this method after the last line that is written.

**Parameters**

| | |
|---|---|
| *write* | where to write the data. |
| *data* | the new chunk of data to write. |
| *len* | the length of the chunk of data to write. |

Here is the call graph for this function:

The documentation for this class was generated from the following file:

- src/include/optionparser.h

## 5.15 xmem::LoadWorker Class Reference

Multithreading-friendly class to do memory loading.

`#include <LoadWorker.h>`

Inheritance diagram for xmem::LoadWorker:

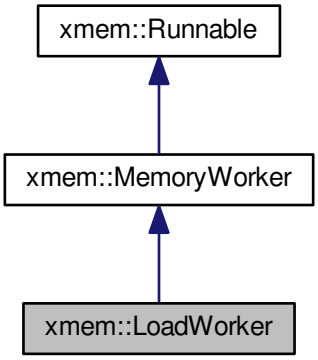


Collaboration diagram for xmem::LoadWorker:

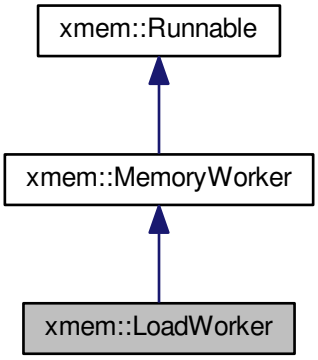


## Public Member Functions

- LoadWorker (void *mem_array, size_t len, SequentialFunction kernel_fptr, SequentialFunction kernel_dummy_fptr, int32_t cpu_affinity)

  *Constructor for sequential-access patterns.*
- LoadWorker (void *mem_array, size_t len, RandomFunction kernel_fptr, RandomFunction kernel_dummy_fptr, int32_t cpu_affinity)

  *Constructor for random-access patterns.*
- virtual ~LoadWorker ()

  *Destructor.*
- virtual void run ()

  *Thread-safe worker method.*

### Additional Inherited Members

#### 5.15.1 Detailed Description

Multithreading-friendly class to do memory loading.

#### 5.15.2 Constructor & Destructor Documentation

##### 5.15.2.1 LoadWorker::LoadWorker ( void * *mem_array,* size_t *len,* SequentialFunction *kernel_fptr,* SequentialFunction *kernel_dummy_fptr,* int32_t *cpu_affinity* )

Constructor for sequential-access patterns.

**Parameters**

| | |
|---|---|
| *mem_array* | Pointer to the memory region to use by this worker. |
| *len* | Length of the memory region to use by this worker. |
| *kernel_fptr* | Pointer to the sequential core benchmark kernel to use. |
| *kernel_dummy_fptr* | Pointer to the sequential dummy version of the core benchmark kernel to use. |
| *cpu_affinity* | Logical CPU identifier to lock this worker's thread to. |

##### 5.15.2.2 LoadWorker::LoadWorker ( void * *mem_array,* size_t *len,* RandomFunction *kernel_fptr,* RandomFunction *kernel_dummy_fptr,* int32_t *cpu_affinity* )

Constructor for random-access patterns.

**Parameters**

| | |
|---|---|
| *mem_array* | Pointer to the memory region to use by this worker. |
| *len* | Length of the memory region to use by this worker. |
| *kernel_fptr* | Pointer to the random core benchmark kernel to use. |
| *kernel_dummy_fptr* | Pointer to the random dummy version of the core benchmark kernel to use. |
| *cpu_affinity* | Logical CPU identifier to lock this worker's thread to. |

The documentation for this class was generated from the following files:

- src/include/LoadWorker.h
- src/LoadWorker.cpp

## 5.16 xmem::MemoryWorker Class Reference

Multithreading-friendly class to run memory access kernels.

`#include <MemoryWorker.h>`

Inheritance diagram for xmem::MemoryWorker:

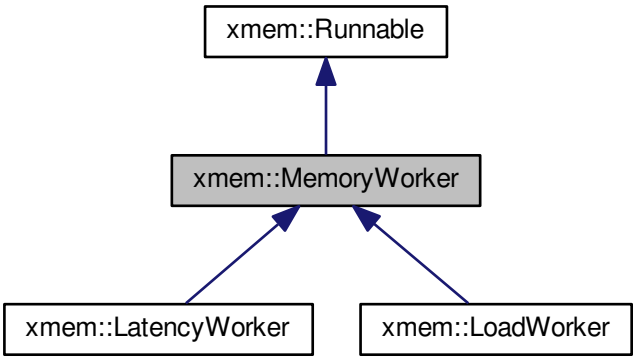


Collaboration diagram for xmem::MemoryWorker:

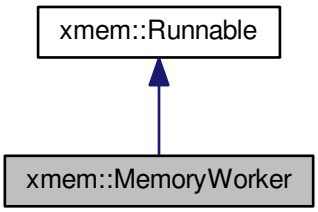


## Public Member Functions

- MemoryWorker (void *mem_array, size_t len, int32_t cpu_affinity)

  *Constructor.*
- virtual ~MemoryWorker ()

  *Destructor.*
- virtual void run ()=0

  *Thread-safe worker method.*
- size_t getLen ()

  *Gets the length of the memory region used by this worker.*
- uint32_t getBytesPerPass ()

  *Gets the number of bytes used in each pass of the benchmark by this worker.*
- uint32_t getPasses ()

  *Gets the number of passes for this worker.*
- tick_t getElapsedTicks ()

  *Gets the elapsed ticks for this worker on the core benchmark kernel.*

- tick_t getElapsedDummyTicks ()

  *Gets the elapsed ticks for this worker on the dummy version of the core benchmark kernel.*

- tick_t getAdjustedTicks ()

  *Gets the adjusted ticks for this worker. This is elapsed ticks minus elapsed dummy ticks.*

- bool hadWarning ()

  *Indicates whether worker's results may be questionable/inaccurate/invalid.*

### Protected Attributes

- void * _mem_array
- size_t _len
- int32_t _cpu_affinity
- uint32_t _bytes_per_pass
- uint32_t _passes
- tick_t _elapsed_ticks
- tick_t _elapsed_dummy_ticks
- tick_t _adjusted_ticks
- bool _warning
- bool _completed

### Additional Inherited Members

## 5.16.1 Detailed Description

Multithreading-friendly class to run memory access kernels.

## 5.16.2 Constructor & Destructor Documentation

#### 5.16.2.1 MemoryWorker::MemoryWorker ( void * *mem_array*, size_t *len*, int32_t *cpu_affinity* )

Constructor.

**Parameters**

| | |
|---|---|
| *mem_array* | Pointer to the memory region to use by this worker. |
| *len* | Length of the memory region to use by this worker. |
| *cpu_affinity* | Logical CPU identifier to lock this worker's thread to. |

## 5.16.3 Member Function Documentation

#### 5.16.3.1 tick_t MemoryWorker::getAdjustedTicks ( )

Gets the adjusted ticks for this worker. This is elapsed ticks minus elapsed dummy ticks.

**Returns**

The adjusted ticks for this worker.

Here is the call graph for this function:

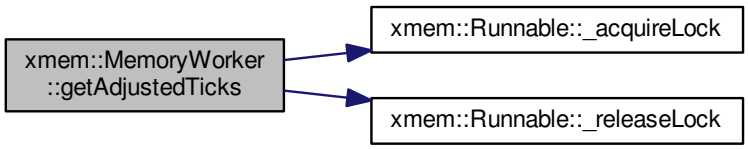


#### 5.16.3.2 uint32_t MemoryWorker::getBytesPerPass ( )

Gets the number of bytes used in each pass of the benchmark by this worker.

**Returns**

Number of bytes in each pass.

Here is the call graph for this function:

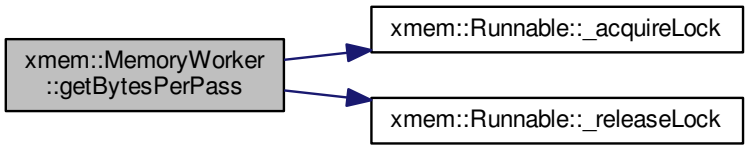


#### 5.16.3.3 tick_t MemoryWorker::getElapsedDummyTicks ( )

Gets the elapsed ticks for this worker on the dummy version of the core benchmark kernel.

**Returns**

The number of elapsed dummy ticks.

Here is the call graph for this function:

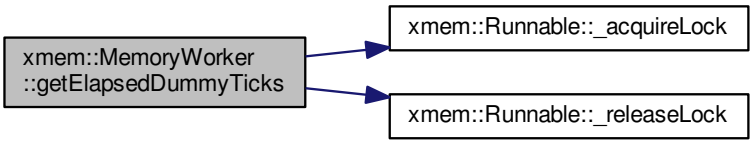


#### 5.16.3.4 tick_t MemoryWorker::getElapsedTicks ( )

Gets the elapsed ticks for this worker on the core benchmark kernel.

**Returns**

The number of elapsed ticks.

Here is the call graph for this function:

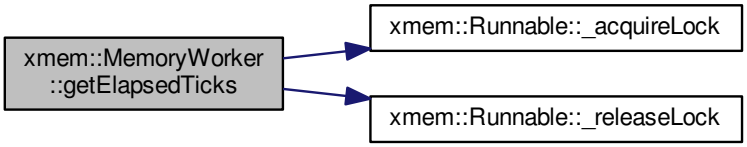


#### 5.16.3.5 size_t MemoryWorker::getLen ( )

Gets the length of the memory region used by this worker.

**Returns**

Length of memory region in bytes.

Here is the call graph for this function:

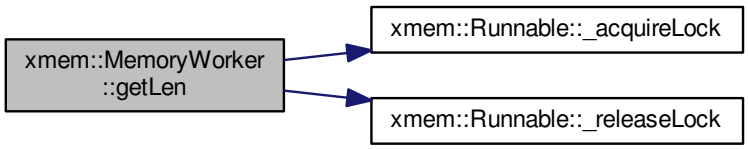


#### 5.16.3.6 uint32_t MemoryWorker::getPasses ( )

Gets the number of passes for this worker.

**Returns**

The number of passes.

Here is the call graph for this function:

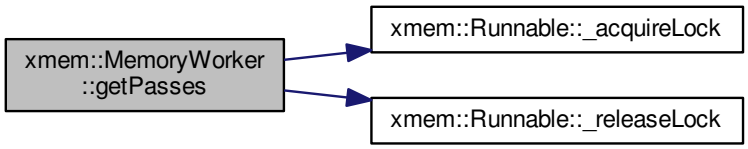


#### 5.16.3.7 bool MemoryWorker::hadWarning ( )

Indicates whether worker's results may be questionable/inaccurate/invalid.

**Returns**

True if the worker's results had a warning.

Here is the call graph for this function:

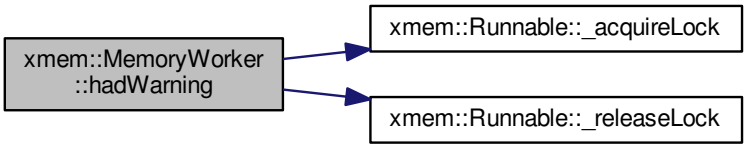


### 5.16.4 Member Data Documentation

#### 5.16.4.1 tick_t xmem::MemoryWorker::_adjusted_ticks [protected]

Elapsed ticks minus dummy elapsed ticks.

#### 5.16.4.2 uint32_t xmem::MemoryWorker::_bytes_per_pass [protected]

Number of bytes accessed in each kernel pass.

#### 5.16.4.3 bool xmem::MemoryWorker::_completed [protected]

If true, worker completed.

#### 5.16.4.4 int32_t xmem::MemoryWorker::_cpu_affinity [protected]

The logical CPU affinity for this worker.

#### 5.16.4.5 tick_t xmem::MemoryWorker::_elapsed_dummy_ticks [protected]

Total elapsed ticks on the dummy kernel routine.

#### 5.16.4.6 tick_t xmem::MemoryWorker::_elapsed_ticks [protected]

Total elapsed ticks on the kernel routine.

#### 5.16.4.7 size_t xmem::MemoryWorker::_len [protected]

The length of the memory region for this worker.

#### 5.16.4.8 void* xmem::MemoryWorker::_mem_array [protected]

The memory region for this worker.

#### 5.16.4.9 uint32_t xmem::MemoryWorker::_passes [protected]

Number of passes.

#### 5.16.4.10 bool xmem::MemoryWorker::_warning [protected]

If true, results may be suspect.

The documentation for this class was generated from the following files:

- src/include/MemoryWorker.h
- src/MemoryWorker.cpp

## 5.17 xmem::MyArg Class Reference

Inheritance diagram for xmem::MyArg:

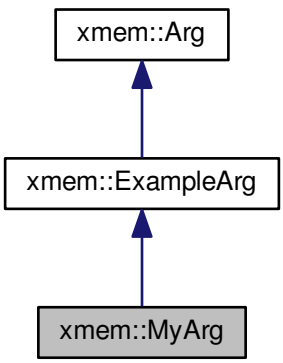

Collaboration diagram for xmem::MyArg:

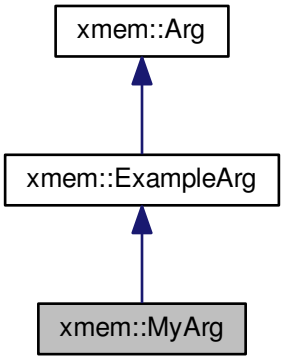


### Static Public Member Functions

- static ArgStatus Integer (const Option &option, bool msg)

  *Checks an option that it is an integer.*

- static ArgStatus NonnegativeInteger (const Option &option, bool msg)

  *Checks an option that it is a nonnegative integer.*

- static ArgStatus PositiveInteger (const Option &option, bool msg)

  *Checks an option that it is a positive integer.*

The documentation for this class was generated from the following file:

- src/include/MyArg.h

## 5.18 xmem::Option Class Reference

A parsed option from the command line together with its argument if it has one.

`#include <optionparser.h>`

Collaboration diagram for xmem::Option:

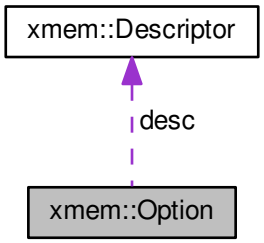


## Public Member Functions

- int type () const

  *Returns Descriptor::type of this Option's Descriptor, or 0 if this Option is invalid (unused).*
- int index () const

  *Returns Descriptor::index of this Option's Descriptor, or -1 if this Option is invalid (unused).*
- int count ()

  *Returns the number of times this Option (or others with the same Descriptor::index) occurs in the argument vector.*
- bool isFirst () const

  *Returns true iff this is the first element of the linked list.*
- bool isLast () const

  *Returns true iff this is the last element of the linked list.*
- Option * first ()

  *Returns a pointer to the first element of the linked list.*
- Option * last ()

  *Returns a pointer to the last element of the linked list.*
- Option * prev ()

  *Returns a pointer to the previous element of the linked list or NULL if called on first().*
- Option * prevwrap ()

  *Returns a pointer to the previous element of the linked list with wrap-around from first() to last().*
- Option * next ()

  *Returns a pointer to the next element of the linked list or NULL if called on last().*
- Option * nextwrap ()

  *Returns a pointer to the next element of the linked list with wrap-around from last() to first().*
- void append (Option *new_last)

  *Makes `new_last` the new last() by chaining it into the list after last().*
- operator const Option * () const

  *Casts from Option to const Option* but only if this Option is valid.*
- operator Option * ()

  *Casts from Option to Option* but only if this Option is valid.*
- Option ()

  *Creates a new Option that is a one-element linked list and has NULL desc, name, arg and namelen.*
- Option (const Descriptor *desc_, const char *name_, const char *arg_)

  *Creates a new Option that is a one-element linked list and has the given values for desc, name and arg.*

- void operator= (const Option &orig)

  *Makes `*this` a copy of `orig` except for the linked list pointers.*
- Option (const Option &orig)

  *Makes `*this` a copy of `orig` except for the linked list pointers.*

### Public Attributes

- const Descriptor * desc

  *Pointer to this Option's Descriptor.*
- const char * name

  *The name of the option as used on the command line.*
- const char * arg

  *Pointer to this Option's argument (if any).*
- int namelen

  *The length of the option name.*

## 5.18.1 Detailed Description

A parsed option from the command line together with its argument if it has one.

The Parser chains all parsed options with the same Descriptor::index together to form a linked list. This allows you to easily implement all of the common ways of handling repeated options and enable/disable pairs.

- Test for presence of a switch in the argument vector:

```
if ( options[QUIET] ) ...
```

- Evaluate –enable-foo/–disable-foo pair where the last one used wins:

```
if ( options[FOO].last()->type() == DISABLE ) ...
```

- Cumulative option (-v verbose, -vv more verbose, -vvv even more verbose):

```
int verbosity = options[VERBOSE].count();
```

- Iterate over all –file=<fname> arguments:

```
for (Option* opt = options[FILE]; opt; opt = opt->next())
 fname = opt->arg; ...
```

## 5.18.2 Constructor & Destructor Documentation

### 5.18.2.1 xmem::Option::Option ( const Descriptor * *desc_*, const char * *name_*, const char * *arg_* ) [inline]

Creates a new Option that is a one-element linked list and has the given values for desc, name and arg.

If `name_` points at a character other than '-' it will be assumed to refer to a short option and namelen will be set to 1. Otherwise the length will extend to the first '=' character or the string's 0-terminator.

### 5.18.2.2 xmem::Option::Option ( const Option & *orig* ) [inline]

Makes `*this` a copy of `orig` except for the linked list pointers.

After this operation `*this` will be a one-element linked list.

### 5.18.3 Member Function Documentation

#### 5.18.3.1 void xmem::Option::append ( Option * *new_last* ) [inline]

Makes `new_last` the new last() by chaining it into the list after last().

It doesn't matter which element you call append() on. The new element will always be appended to last().

**Attention**

`new_last` must not yet be part of a list, or that list will become corrupted, because this method does not unchain `new_last` from an existing list.

Here is the call graph for this function:

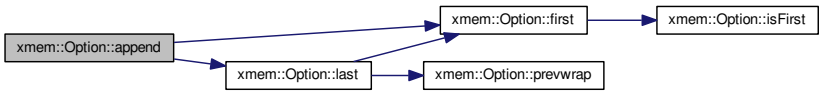


#### 5.18.3.2 int xmem::Option::count ( ) [inline]

Returns the number of times this Option (or others with the same Descriptor::index) occurs in the argument vector.

This corresponds to the number of elements in the linked list this Option is part of. It doesn't matter on which element you call count(). The return value is always the same.

Use this to implement cumulative options, such as -v, -vv, -vvv for different verbosity levels.

Returns 0 when called for an unused/invalid option.

Here is the call graph for this function:

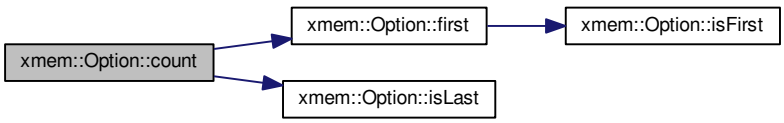


#### 5.18.3.3 Option* xmem::Option::first ( ) [inline]

Returns a pointer to the first element of the linked list.

Use this when you want the first occurrence of an option on the command line to take precedence. Note that this is not the way most programs handle options. You should probably be using last() instead.

**Note**

This method may be called on an unused/invalid option and will return a pointer to the option itself.

Here is the call graph for this function:

#### 5.18.3.4 bool xmem::Option::isFirst ( ) const [inline]

Returns true iff this is the first element of the linked list.

The first element in the linked list is the first option on the command line that has the respective Descriptor::index value.

Returns true for an unused/invalid option.

#### 5.18.3.5 bool xmem::Option::isLast ( ) const [inline]

Returns true iff this is the last element of the linked list.

The last element in the linked list is the last option on the command line that has the respective Descriptor::index value.

Returns true for an unused/invalid option.

#### 5.18.3.6 Option* xmem::Option::last ( ) [inline]

Returns a pointer to the last element of the linked list.

Use this when you want the last occurrence of an option on the command line to take precedence. This is the most common way of handling conflicting options.

**Note**

This method may be called on an unused/invalid option and will return a pointer to the option itself.

**Tip:**

If you have options with opposite meanings (e.g. `-enable-foo` and `-disable-foo`), you can assign them the same Descriptor::index to get them into the same list. Distinguish them by Descriptor::type and all you have to do is check `last()->type()` to get the state listed last on the command line.

Here is the call graph for this function:

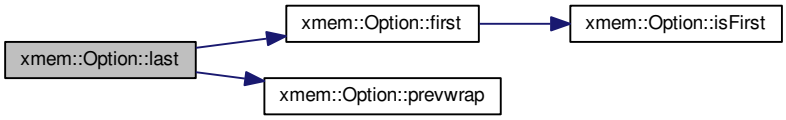


#### 5.18.3.7 Option* xmem::Option::next ( ) [inline]

Returns a pointer to the next element of the linked list or NULL if called on last().

If called on last() this method returns NULL. Otherwise it will return the option with the same Descriptor::index that follows this option on the command line.

Here is the call graph for this function:

#### 5.18.3.8 Option* xmem::Option::nextwrap ( ) [inline]

Returns a pointer to the next element of the linked list with wrap-around from last() to first().

If called on last() this method returns first(). Otherwise it will return the option with the same Descriptor::index that follows this option on the command line.

#### 5.18.3.9 xmem::Option::operator const Option * ( ) const [inline]

Casts from Option to const Option* but only if this Option is valid.

If this Option is valid (i.e. `desc!=NULL`), returns this. Otherwise returns NULL. This allows testing an Option directly in an if-clause to see if it is used:

```
if (options[CREATE])
{
  ...
}
```

It also allows you to write loops like this:

```
for (Option* opt = options[FILE]; opt; opt = opt->next())
 fname = opt->arg; ...
```

#### 5.18.3.10 xmem::Option::operator Option * ( ) [inline]

Casts from Option to Option* but only if this Option is valid.

If this Option is valid (i.e. `desc!=NULL`), returns this. Otherwise returns NULL. This allows testing an Option directly in an if-clause to see if it is used:

```
if (options[CREATE])
{
  ...
}
```

It also allows you to write loops like this:

```
for (Option* opt = options[FILE]; opt; opt = opt->next())
 fname = opt->arg; ...
```

#### 5.18.3.11 void xmem::Option::operator= ( const Option & *orig* ) [inline]

Makes `*this` a copy of `orig` except for the linked list pointers.

After this operation `*this` will be a one-element linked list.

#### 5.18.3.12 Option* xmem::Option::prev ( ) [inline]

Returns a pointer to the previous element of the linked list or NULL if called on first().

If called on first() this method returns NULL. Otherwise it will return the option with the same Descriptor::index that precedes this option on the command line.

Here is the call graph for this function:

#### 5.18.3.13 Option* xmem::Option::prevwrap ( ) [inline]

Returns a pointer to the previous element of the linked list with wrap-around from first() to last().

If called on first() this method returns last(). Otherwise it will return the option with the same Descriptor::index that precedes this option on the command line.

#### 5.18.3.14 int xmem::Option::type ( ) const [inline]

Returns Descriptor::type of this Option's Descriptor, or 0 if this Option is invalid (unused).

Because this method (and last(), too) can be used even on unused Options with desc==0, you can (provided you arrange your types properly) switch on type() without testing validity first.

```
enum OptionType { UNUSED=0, DISABLED=0, ENABLED=1 };
enum OptionIndex { FOO };
const Descriptor usage[] = {
  { FOO, ENABLED,  "", "enable-foo",  Arg::None, 0 },
  { FOO, DISABLED, "", "disable-foo", Arg::None, 0 },
  { 0, 0, 0, 0, 0, 0 } };
...
switch(options[FOO].last()->type()) // no validity check required!
{
  case ENABLED: ...
  case DISABLED: ...  // UNUSED==DISABLED !
}
```

### 5.18.4 Member Data Documentation

#### 5.18.4.1 const char* xmem::Option::arg

Pointer to this Option's argument (if any).

NULL if this option has no argument. Do not confuse this with the empty string which is a valid argument.

#### 5.18.4.2 const Descriptor* xmem::Option::desc

Pointer to this Option's Descriptor.

Remember that the first dummy descriptor (see Descriptor::longopt) is used for unknown options.

**Attention**

`desc==NULL` signals that this Option is unused. This is the default state of elements in the result array. You don't need to test `desc` explicitly. You can simply write something like this:

```
if (options[CREATE])
{
  ...
}
```

This works because of `operator const Option*()` .

#### 5.18.4.3 const char* xmem::Option::name

The name of the option as used on the command line.

The main purpose of this string is to be presented to the user in messages.

In the case of a long option, this is the actual `argv` pointer, i.e. the first character is a '-'. In the case of a short option this points to the option character within the `argv` string.

Note that in the case of a short option group or an attached option argument, this string will contain additional characters following the actual name. Use namelen to filter out the actual option name only.

#### 5.18.4.4 int xmem::Option::namelen

The length of the option name.

Because name points into the actual `argv` string, the option name may be followed by more characters (e.g. other short options in the same short option group). This value is the number of bytes (not characters!) that are part of the actual name.

For a short option, this length is always 1. For a long option this length is always at least 2 if single minus long options are permitted and at least 3 if they are disabled.

**Note**

In the pathological case of a minus within a short option group (e.g. `-xf-z`), this length is incorrect, because this case will be misinterpreted as a long option and the name will therefore extend to the string's 0-terminator or a following '=" character if there is one. This is irrelevant for most uses of name and `namelen`. If you really need to distinguish the case of a long and a short option, compare name to the `argv` pointers. A long option's `name` is always identical to one of them, whereas a short option's is never.

The documentation for this class was generated from the following file:

- src/include/optionparser.h

## 5.19 xmem::PrintUsageImplementation::OStreamWriter< OStream > Struct Template Reference

Inheritance diagram for xmem::PrintUsageImplementation::OStreamWriter< OStream >:

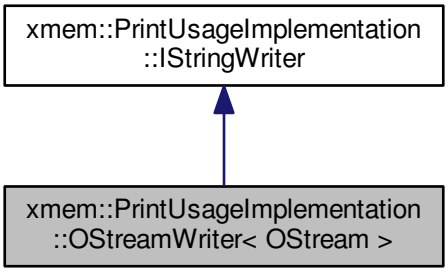


Collaboration diagram for xmem::PrintUsageImplementation::OStreamWriter< OStream >:

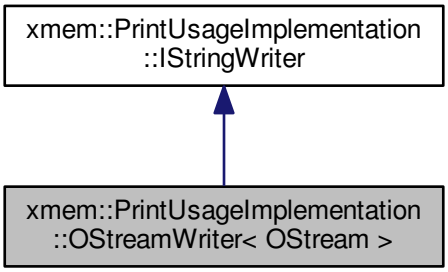

### Public Member Functions

- virtual void operator() (const char *str, int size)

  *Writes the given number of chars beginning at the given pointer somewhere.*
- **OStreamWriter** (OStream &o)

### Public Attributes

- OStream & **ostream**

The documentation for this struct was generated from the following file:

- src/include/optionparser.h

## 5.20 xmem::Parser Class Reference

Checks argument vectors for validity and parses them into data structures that are easier to work with.

`#include <optionparser.h>`

### Classes

- struct Action
- class StoreOptionAction

### Public Member Functions

- Parser ()

  *Creates a new Parser.*
- Parser (bool gnu, const Descriptor usage[], int argc, const char **argv, Option options[], Option buffer[], int min_abbr_len=0, bool single_minus_longopt=false, int bufmax=-1)

  *Creates a new Parser and immediately parses the given argument vector.*
- Parser (bool gnu, const Descriptor usage[], int argc, char **argv, Option options[], Option buffer[], int min_abbr_len=0, bool single_minus_longopt=false, int bufmax=-1)

  *Parser(...) with non-const argv.*
- Parser (const Descriptor usage[], int argc, const char **argv, Option options[], Option buffer[], int min_abbr_len=0, bool single_minus_longopt=false, int bufmax=-1)

  *POSIX Parser(...) (gnu==false).*
- Parser (const Descriptor usage[], int argc, char **argv, Option options[], Option buffer[], int min_abbr_len=0, bool single_minus_longopt=false, int bufmax=-1)

  *POSIX Parser(...) (gnu==false) with non-const argv.*
- void parse (bool gnu, const Descriptor usage[], int argc, const char **argv, Option options[], Option buffer[], int min_abbr_len=0, bool single_minus_longopt=false, int bufmax=-1)

  *Parses the given argument vector.*
- void parse (bool gnu, const Descriptor usage[], int argc, char **argv, Option options[], Option buffer[], int min_abbr_len=0, bool single_minus_longopt=false, int bufmax=-1)

  *parse() with non-const argv.*
- void parse (const Descriptor usage[], int argc, const char **argv, Option options[], Option buffer[], int min_abbr_len=0, bool single_minus_longopt=false, int bufmax=-1)

  *POSIX parse() (gnu==false).*
- void parse (const Descriptor usage[], int argc, char **argv, Option options[], Option buffer[], int min_abbr_len=0, bool single_minus_longopt=false, int bufmax=-1)

*POSIX parse() (gnu==false) with non-const argv.*

- int optionsCount ()

  *Returns the number of valid Option objects in `buffer[]`.*

- int nonOptionsCount ()

  *Returns the number of non-option arguments that remained at the end of the most recent parse() that actually encountered non-option arguments.*

- const char ** nonOptions ()

  *Returns a pointer to an array of non-option arguments (only valid if `nonOptionsCount() >0` ).*

- const char * nonOption (int i)

  *Returns **`nonOptions()[i]`** (without checking if i is in range!).*

- bool error ()

  *Returns `true` if an unrecoverable error occurred while parsing options.*

## Friends

- struct **Stats**

### 5.20.1 Detailed Description

Checks argument vectors for validity and parses them into data structures that are easier to work with.

**Example:**

```
int main(int argc, char* argv[])
{
  argc-=(argc>0); argv+=(argc>0); // skip program name argv[0] if present
  option::Stats  stats(usage, argc, argv);
  option::Option options[stats.options_max], buffer[stats.buffer_max];
  option::Parser parse(usage, argc, argv, options, buffer);

  if (parse.error())
    return 1;

  if (options[HELP])
  ...
```

### 5.20.2 Constructor & Destructor Documentation

#### 5.20.2.1 xmem::Parser::Parser ( bool *gnu*, const Descriptor *usage[]*, int *argc*, const char ** *argv*, Option *options[]*, Option *buffer[]*, int *min_abbr_len* = `0`, bool *single_minus_longopt* = `false`, int *bufmax* = `-1` ) `[inline]`

Creates a new Parser and immediately parses the given argument vector.

**Parameters**

| | |
|---|---|
| *gnu* | if true, parse() will not stop at the first non-option argument. Instead it will reorder arguments so that all non-options are at the end. This is the default behaviour of GNU getopt() but is not conforming to POSIX.<br>Note, that once the argument vector has been reordered, the `gnu` flag will have no further effect on this argument vector. So it is enough to pass `gnu==true` when creating Stats. |

| | |
|---|---|
| *usage* | Array of Descriptor objects that describe the options to support. The last entry of this array must have 0 in all fields. |
| *argc* | The number of elements from `argv` that are to be parsed. If you pass -1, the number will be determined automatically. In that case the `argv` list must end with a NULL pointer. |
| *argv* | The arguments to be parsed. If you pass -1 as `argc` the last pointer in the `argv` list must be NULL to mark the end. |
| *options* | Each entry is the first element of a linked list of Options. Each new option that is parsed will be appended to the list specified by that Option's Descriptor::index. If an entry is not yet used (i.e. the Option is invalid), it will be replaced rather than appended to.<br>The minimum length of this array is the greatest Descriptor::index value that occurs in `usage` *PLUS* ONE. |
| *buffer* | Each argument that is successfully parsed (including unknown arguments, if they have a Descriptor whose CheckArg does not return ARG_ILLEGAL) will be stored in this array. parse() scans the array for the first invalid entry and begins writing at that index. You can pass `bufmax` to limit the number of options stored. |
| *min_abbr_len* | Passing a value `min_abbr_len > 0` enables abbreviated long options. The parser will match a prefix of a long option as if it was the full long option (e.g. `-foob=10` will be interpreted as if it was `-foobar=10` ), as long as the prefix has at least `min_abbr_len` characters (not counting the `-` ) and is unambiguous.<br>Be careful if combining `min_abbr_len=1` with `single_minus_longopt=true` because the ambiguity check does not consider short options and abbreviated single minus long options will take precedence over short options. |
| *single_minus_longopt* | Passing `true` for this option allows long options to begin with a single minus. The double minus form will still be recognized. Note that single minus long options take precedence over short options and short option groups. E.g. `-file` would be interpreted as `-file` and not as `-f -i -l -e` (assuming a long option named "`file`" exists). |
| *bufmax* | The greatest index in the `buffer`[] array that parse() will write to is `bufmax-1`. If there are more options, they will be processed (in particular their CheckArg will be called) but not stored.<br>If you used Stats::buffer_max to dimension this array, you can pass -1 (or not pass `bufmax` at all) which tells parse() that the buffer is "large enough". |

**Attention**

Remember that `options` and `buffer` store Option *objects*, not pointers. Therefore it is not possible for the same object to be in both arrays. For those options that are found in both `buffer`[] and `options`[] the respective objects are independent copies. And only the objects in `options`[] are properly linked via Option::next() and Option::prev(). You can iterate over `buffer`[] to process all options in the order they appear in the argument vector, but if you want access to the other Options with the same Descriptor::index, then you *must* access the linked list via `options`[]. You can get the linked list in options from a buffer object via something like `options`[buffer[i].index()].

Here is the call graph for this function:

### 5.20.3 Member Function Documentation

#### 5.20.3.1 bool xmem::Parser::error ( ) [inline]

Returns `true` if an unrecoverable error occurred while parsing options.

An illegal argument to an option (i.e. CheckArg returns ARG_ILLEGAL) is an unrecoverable error that aborts the parse. Unknown options are only an error if their CheckArg function returns ARG_ILLEGAL. Otherwise they are collected. In that case if you want to exit the program if either an illegal argument or an unknown option has been passed, use code like this

```
if (parser.error() || options[UNKNOWN])
  exit(1);
```

#### 5.20.3.2 const char** xmem::Parser::nonOptions ( ) [inline]

Returns a pointer to an array of non-option arguments (only valid if `nonOptionsCount() >0` ).

**Note**

- parse() does not copy arguments, so this pointer points into the actual argument vector as passed to parse().
- As explained at nonOptionsCount() this pointer is only changed by parse() calls that actually encounter non-option arguments. A parse() call that encounters only options, will not change nonOptions().

#### 5.20.3.3 int xmem::Parser::nonOptionsCount ( ) [inline]

Returns the number of non-option arguments that remained at the end of the most recent parse() that actually encountered non-option arguments.

**Note**

A parse() that does not encounter non-option arguments will leave this value as well as nonOptions() undisturbed. This means you can feed the Parser a default argument vector that contains non-option arguments (e.g. a default filename). Then you feed it the actual arguments from the user. If the user has supplied at least one non-option argument, all of the non-option arguments from the default disappear and are replaced by the user's non-option arguments. However, if the user does not supply any non-option arguments the defaults will still be in effect.

#### 5.20.3.4 int xmem::Parser::optionsCount ( ) [inline]

Returns the number of valid Option objects in `buffer`[].

**Note**

- The returned value always reflects the number of Options in the buffer[] array used for the most recent call to parse().
- The count (and the buffer[]) includes unknown options if they are collected (see Descriptor::longopt).

#### 5.20.3.5 void xmem::Parser::parse ( bool *gnu*, const Descriptor *usage[]*, int *argc*, const char ** *argv*, Option *options[]*, Option *buffer[]*, int *min_abbr_len* = `0`, bool *single_minus_longopt* = `false`, int *bufmax* = `-1` ) [inline]

Parses the given argument vector.

**Parameters**

| | |
|---|---|
| *gnu* | if true, parse() will not stop at the first non-option argument. Instead it will reorder arguments so that all non-options are at the end. This is the default behaviour of GNU getopt() but is not conforming to POSIX.<br>Note, that once the argument vector has been reordered, the `gnu` flag will have no further effect on this argument vector. So it is enough to pass `gnu==true` when creating Stats. |
| *usage* | Array of Descriptor objects that describe the options to support. The last entry of this array must have 0 in all fields. |
| *argc* | The number of elements from `argv` that are to be parsed. If you pass -1, the number will be determined automatically. In that case the `argv` list must end with a NULL pointer. |
| *argv* | The arguments to be parsed. If you pass -1 as `argc` the last pointer in the `argv` list must be NULL to mark the end. |
| *options* | Each entry is the first element of a linked list of Options. Each new option that is parsed will be appended to the list specified by that Option's Descriptor::index. If an entry is not yet used (i.e. the Option is invalid), it will be replaced rather than appended to.<br>The minimum length of this array is the greatest Descriptor::index value that occurs in `usage` *PLUS* ONE. |
| *buffer* | Each argument that is successfully parsed (including unknown arguments, if they have a Descriptor whose CheckArg does not return ARG_ILLEGAL) will be stored in this array. parse() scans the array for the first invalid entry and begins writing at that index. You can pass `bufmax` to limit the number of options stored. |
| *min_abbr_len* | Passing a value `min_abbr_len > 0` enables abbreviated long options. The parser will match a prefix of a long option as if it was the full long option (e.g. `-foob=10` will be interpreted as if it was `-foobar=10` ), as long as the prefix has at least `min_abbr_len` characters (not counting the `-` ) and is unambiguous.<br>Be careful if combining `min_abbr_len=1` with `single_minus_longopt=true` because the ambiguity check does not consider short options and abbreviated single minus long options will take precedence over short options. |
| *single_minus_longopt* | Passing `true` for this option allows long options to begin with a single minus. The double minus form will still be recognized. Note that single minus long options take precedence over short options and short option groups. E.g. `-file` would be interpreted as `-file` and not as `-f -i -l -e` (assuming a long option named "`file`" exists). |
| *bufmax* | The greatest index in the `buffer`[] array that parse() will write to is `bufmax-1`. If there are more options, they will be processed (in particular their CheckArg will be called) but not stored.<br>If you used Stats::buffer_max to dimension this array, you can pass -1 (or not pass `bufmax` at all) which tells parse() that the buffer is "large enough". |

**Attention**

Remember that `options` and `buffer` store Option *objects*, not pointers. Therefore it is not possible for the same object to be in both arrays. For those options that are found in both `buffer`[] and `options`[] the respective objects are independent copies. And only the objects in `options`[] are properly linked via Option::next() and Option::prev(). You can iterate over `buffer`[] to process all options in the order they appear in the argument vector, but if you want access to the other Options with the same Descriptor::index, then you *must* access the linked list via `options`[]. You can get the linked list in options from a buffer object via something like `options`[buffer[i].index()].

The documentation for this class was generated from the following file:

- src/include/optionparser.h

## 5.21 xmem::PowerReader Class Reference

An abstract base class for measuring power from an arbitrary source. This class is runnable using a worker thread.

`#include <PowerReader.h>`

Inheritance diagram for xmem::PowerReader:

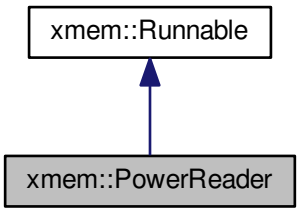


Collaboration diagram for xmem::PowerReader:

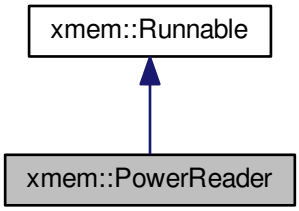


## Public Member Functions

- PowerReader (uint32_t sampling_period, double power_units, std::string name, int32_t cpu_affinity)

  *Constructor.*
- ~PowerReader ()

  *Destructor.*
- virtual void run ()=0

  *Starts measuring power at the rate implied by the sampling_period passed in the constructor. Call stop() to indicate to stop measuring.*
- bool stop ()

  *Signals to stop measuring power. This is a non-blocking call and return does not indicate the measurement has actually stopped.*
- bool calculateMetrics ()

  *Calculates the relevant metrics.*
- bool clear ()

  *Clears the stored power data.*
- bool clear_and_reset ()

  *Clears the stored power data and resets state so that a new thread can be used with this object.*
- std::vector< double > getPowerTrace ()

*Gets the power trace.*

- double getAveragePower ()

  *Gets the average power.*

- double getPeakPower ()

  *Gets the peak power.*

- double getLastSample ()

  *Gets the last sample.*

- uint32_t getSamplingPeriod ()

  *Gets the sampling period.*

- double getPowerUnits ()

  *Gets the units of samples in watts.*

- size_t getNumSamples ()

  *Gets the number of samples collected.*

- std::string name ()

  *Gets the name of this object.*

### Protected Attributes

- bool _stop_signal
- double _power_units
- std::string _name
- int32_t _cpu_affinity
- std::vector< double > _power_trace
- double _average_power
- double _peak_power
- size_t _num_samples
- uint32_t _sampling_period

### Additional Inherited Members

## 5.21.1 Detailed Description

An abstract base class for measuring power from an arbitrary source. This class is runnable using a worker thread.

## 5.21.2 Constructor & Destructor Documentation

### 5.21.2.1 PowerReader::PowerReader ( uint32_t *sampling_period*, double *power_units*, std::string *name*, int32_t *cpu_affinity* )

Constructor.

**Parameters**

| | |
|---|---|
| *sampling_period* | The time between power samples in milliseconds. |
| *power_units* | The power units for each sample in watts. |
| *name* | The human-friendly name of this object. |
| *cpu_affinity* | The logical CPU to be used by the thread calling this object's run() method. If negative, any CPU is OK (no affinity). |

### 5.21.3 Member Function Documentation

#### 5.21.3.1 bool PowerReader::calculateMetrics ( )

Calculates the relevant metrics.

**Returns**

True on success.

Here is the call graph for this function:

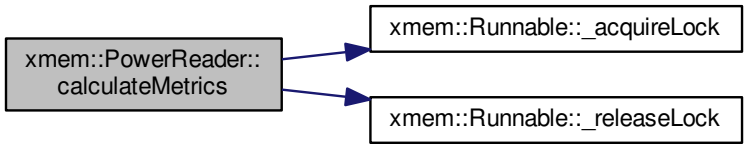


#### 5.21.3.2 bool PowerReader::clear ( )

Clears the stored power data.

**Returns**

True on success.

Here is the call graph for this function:

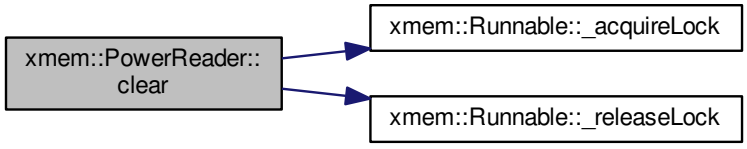


#### 5.21.3.3 bool PowerReader::clear_and_reset ( )

Clears the stored power data and resets state so that a new thread can be used with this object.

**Returns**

True on success.

Here is the call graph for this function:

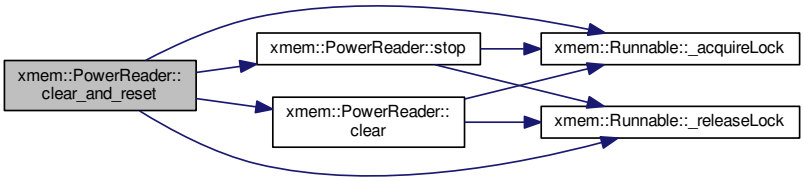


#### 5.21.3.4 double PowerReader::getAveragePower ( )

Gets the average power.

**Returns**

The average power from the measurements. If no data was collected, returns 0.

Here is the call graph for this function:

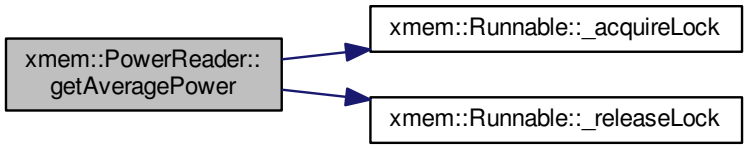


#### 5.21.3.5 double PowerReader::getLastSample ( )

Gets the last sample.

**Returns**

The last power sample measured.

Here is the call graph for this function:

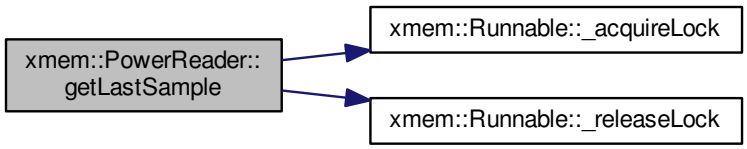


#### 5.21.3.6 size_t PowerReader::getNumSamples ( )

Gets the number of samples collected.

**Returns**

Number of samples collected.

Here is the call graph for this function:

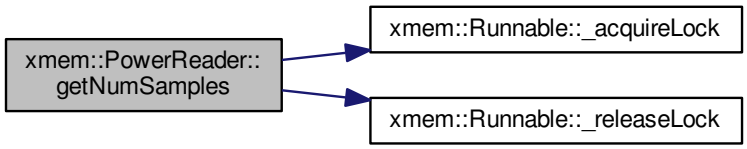


#### 5.21.3.7 double PowerReader::getPeakPower ( )

Gets the peak power.

**Returns**

The peak power sample from the measurements. If no data was collected, returns 0.

Here is the call graph for this function:

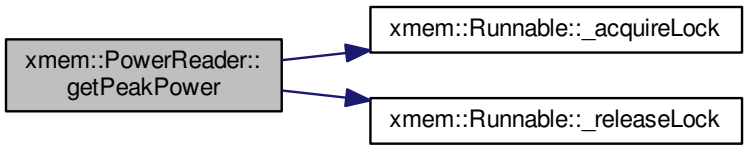


#### 5.21.3.8 std::vector< double > PowerReader::getPowerTrace ( )

Gets the power trace.

**Returns**

The measured power trace in a vector. If no data was collected, the vector will be empty.

Here is the call graph for this function:

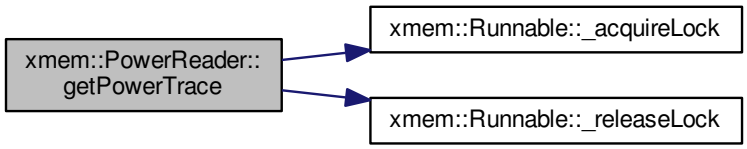


#### 5.21.3.9 double PowerReader::getPowerUnits ( )

Gets the units of samples in watts.

**Returns**

The power units for each measurement sample in watts. For example, if each measurement is in milliwatts, then this returns 1e-3.

Here is the call graph for this function:

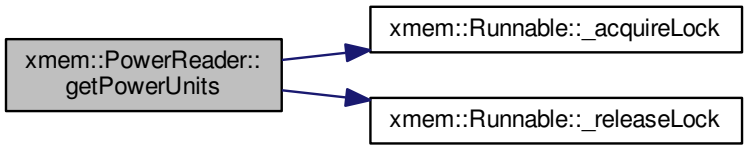


#### 5.21.3.10 uint32_t PowerReader::getSamplingPeriod ( )

Gets the sampling period.

**Returns**

The sampling period of the measurements in milliseconds.

Here is the call graph for this function:

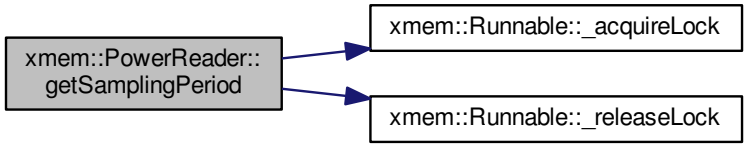


#### 5.21.3.11 std::string PowerReader::name ( )

Gets the name of this object.

**Returns**

The human-friendly name of this PowerReader.

Here is the call graph for this function:

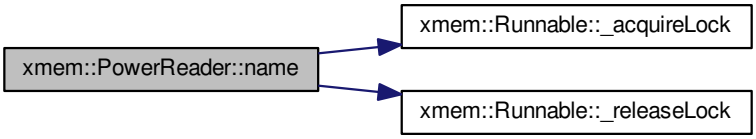


#### 5.21.3.12 bool PowerReader::stop ( )

Signals to stop measuring power. This is a non-blocking call and return does not indicate the measurement has actually stopped.

**Returns**

True if it successfully signaled a stop.

Here is the call graph for this function:

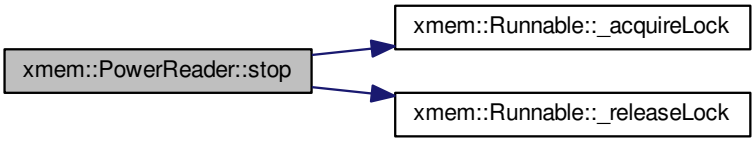


### 5.21.4 Member Data Documentation

#### 5.21.4.1 double xmem::PowerReader::_average_power [protected]

The average power.

#### 5.21.4.2 int32_t xmem::PowerReader::_cpu_affinity [protected]

CPU affinity for any thread using this object's run() method. If negative, no affinity preference.

#### 5.21.4.3 std::string xmem::PowerReader::_name [protected]

Name of this object.

#### 5.21.4.4 size_t xmem::PowerReader::_num_samples [protected]

The number of samples collected.

#### 5.21.4.5 double xmem::PowerReader::_peak_power [protected]

The peak power observed.

#### 5.21.4.6 std::vector<double> xmem::PowerReader::_power_trace [protected]

The time-ordered list of power samples. The first index is the oldest measurement.

#### 5.21.4.7 double xmem::PowerReader::_power_units [protected]

Power units in watts.

#### 5.21.4.8 uint32_t xmem::PowerReader::_sampling_period [protected]

Power sampling period in milliseconds.

#### 5.21.4.9 bool xmem::PowerReader::_stop_signal [protected]

When true, the run() function should finish after the current sample iteration it is working on.

The documentation for this class was generated from the following files:

- src/include/PowerReader.h
- src/PowerReader.cpp

## 5.22 xmem::PrintUsageImplementation Struct Reference

### Classes

- struct FunctionWriter
- struct IStringWriter
- class LinePartIterator
- class LineWrapper
- struct OStreamWriter
- struct StreamWriter
- struct SyscallWriter
- struct TemporaryWriter

### Static Public Member Functions

- static void **upmax** (int &i1, int i2)
- static void **indent** (IStringWriter &write, int &x, int want_x)
- static bool isWideChar (unsigned ch)

  *Returns true if ch is the unicode code point of a wide character.*
- static void **printUsage** (IStringWriter &write, const Descriptor usage[], int width=80, int last_column_min_-percent=50, int last_column_own_line_max_percent=75)

### 5.22.1 Member Function Documentation

#### 5.22.1.1 static bool xmem::PrintUsageImplementation::isWideChar ( unsigned *ch* ) [inline], [static]

Returns true if ch is the unicode code point of a wide character.

**Note**

The following character ranges are treated as wide

```
1100..115F
2329..232A  (just 2 characters!)
2E80..A4C6  except for 303F
A960..A97C
AC00..D7FB
F900..FAFF
FE10..FE6B
FF01..FF60
FFE0..FFE6
1B000......
```

The documentation for this struct was generated from the following file:

- src/include/optionparser.h

## 5.23 xmem::Runnable Class Reference

A base class for any object that implements a thread-safe run() function for use by Thread objects.

`#include <Runnable.h>`

Inheritance diagram for xmem::Runnable:

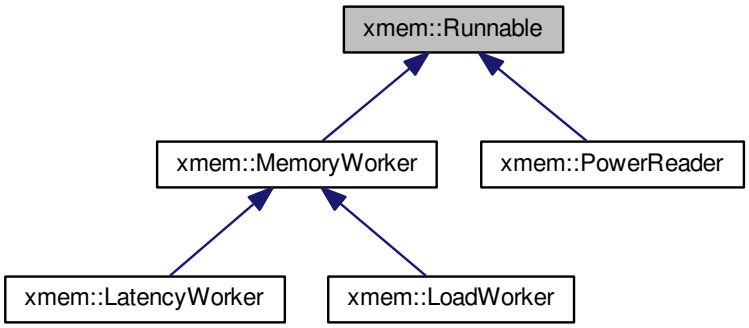


### Public Member Functions

- Runnable ()

  *Constructor.*
- ~Runnable ()

  *Destructor.*
- virtual void run ()=0

  *Does some "work". Pure virtual method that any derived class must implement in a thread-safe manner.*

### Protected Member Functions

- bool _acquireLock (int32_t timeout)

  *Acquires the object lock to access all object state in thread-safe manner.*
- bool _releaseLock ()

  *Releases the object lock to access all object state in thread-safe manner.*

### 5.23.1 Detailed Description

A base class for any object that implements a thread-safe run() function for use by Thread objects.

### 5.23.2 Member Function Documentation

#### 5.23.2.1 bool Runnable::_acquireLock ( int32_t *timeout* ) [protected]

Acquires the object lock to access all object state in thread-safe manner.

**Parameters**

| | |
|---|---|
| *timeout* | timeout in milliseconds to acquire the lock. If 0, does not wait at all. If negative, waits indefinitely. |

**Returns**

true on success. If not successful, the lock was not acquired, possibly due to a timeout, or the lock might already be held.

#### 5.23.2.2 bool Runnable::_releaseLock ( ) [protected]

Releases the object lock to access all object state in thread-safe manner.

**Returns**

true on success. If not successful, the lock is either still held or the call was illegal (e.g., releasing a lock that was never acquired).

The documentation for this class was generated from the following files:

- src/include/Runnable.h
- src/Runnable.cpp

## 5.24 xmem::Stats Struct Reference

Determines the minimum lengths of the buffer and options arrays used for Parser.

`#include <optionparser.h>`

### Classes

- class CountOptionsAction

### Public Member Functions

- Stats ()

  *Creates a Stats object with counts set to 1 (for the sentinel element).*
- Stats (bool gnu, const Descriptor usage[], int argc, const char **argv, int min_abbr_len=0, bool single_minus_longopt=false)

  *Creates a new Stats object and immediately updates it for the given `usage` and argument vector. You may pass 0 for `argc` and/or `argv`, if you just want to update options_max.*
- Stats (bool gnu, const Descriptor usage[], int argc, char **argv, int min_abbr_len=0, bool single_minus_longopt=false)

  *Stats(...) with non-const argv.*
- Stats (const Descriptor usage[], int argc, const char **argv, int min_abbr_len=0, bool single_minus_longopt=false)

  *POSIX Stats(...) (gnu==false).*
- Stats (const Descriptor usage[], int argc, char **argv, int min_abbr_len=0, bool single_minus_longopt=false)

  *POSIX Stats(...) (gnu==false) with non-const argv.*
- void add (bool gnu, const Descriptor usage[], int argc, const char **argv, int min_abbr_len=0, bool single_minus_longopt=false)

  *Updates this Stats object for the given `usage` and argument vector. You may pass 0 for `argc` and/or `argv`, if you just want to update options_max.*
- void add (bool gnu, const Descriptor usage[], int argc, char **argv, int min_abbr_len=0, bool single_minus_longopt=false)

  *add() with non-const argv.*
- void add (const Descriptor usage[], int argc, const char **argv, int min_abbr_len=0, bool single_minus_longopt=false)

  *POSIX add() (gnu==false).*
- void add (const Descriptor usage[], int argc, char **argv, int min_abbr_len=0, bool single_minus_longopt=false)

  *POSIX add() (gnu==false) with non-const argv.*

### Public Attributes

- unsigned buffer_max

  *Number of elements needed for a `buffer[]` array to be used for parsing the same argument vectors that were fed into this Stats object.*
- unsigned options_max

  *Number of elements needed for an `options[]` array to be used for parsing the same argument vectors that were fed into this Stats object.*

## 5.24.1 Detailed Description

Determines the minimum lengths of the buffer and options arrays used for Parser.

Because Parser doesn't use dynamic memory its output arrays have to be pre-allocated. If you don't want to use fixed size arrays (which may turn out too small, causing command line arguments to be dropped), you can use Stats to determine the correct sizes. Stats work cumulative. You can first pass in your default options and then the real options and afterwards the counts will reflect the union.

## 5.24.2 Constructor & Destructor Documentation

### 5.24.2.1 xmem::Stats::Stats ( bool *gnu*, const Descriptor *usage[]*, int *argc*, const char ** *argv*, int *min_abbr_len* = `0`, bool *single_minus_longopt* = `false` ) [inline]

Creates a new Stats object and immediately updates it for the given `usage` and argument vector. You may pass 0 for `argc` and/or `argv`, if you just want to update options_max.

**Note**

The calls to Stats methods must match the later calls to Parser methods. See Parser::parse() for the meaning of the arguments.

Here is the call graph for this function:

### 5.24.3 Member Function Documentation

#### 5.24.3.1 void xmem::Stats::add ( bool *gnu*, const Descriptor *usage[]*, int *argc*, const char ** *argv*, int *min_abbr_len* = 0, bool *single_minus_longopt* = false ) [inline]

Updates this Stats object for the given `usage` and argument vector. You may pass 0 for `argc` and/or `argv`, if you just want to update options_max.

**Note**

The calls to Stats methods must match the later calls to Parser methods. See Parser::parse() for the meaning of the arguments.

### 5.24.4 Member Data Documentation

#### 5.24.4.1 unsigned xmem::Stats::buffer_max

Number of elements needed for a `buffer`[] array to be used for parsing the same argument vectors that were fed into this Stats object.

**Note**

This number is always 1 greater than the actual number needed, to give you a sentinel element.

#### 5.24.4.2 unsigned xmem::Stats::options_max

Number of elements needed for an `options`[] array to be used for parsing the same argument vectors that were fed into this Stats object.

**Note**

- This number is always 1 greater than the actual number needed, to give you a sentinel element.
- This number depends only on the `usage`, not the argument vectors, because the `options` array needs exactly one slot for each possible Descriptor::index.

The documentation for this struct was generated from the following file:

- src/include/optionparser.h

## 5.25 xmem::Parser::StoreOptionAction Class Reference

Inheritance diagram for xmem::Parser::StoreOptionAction:

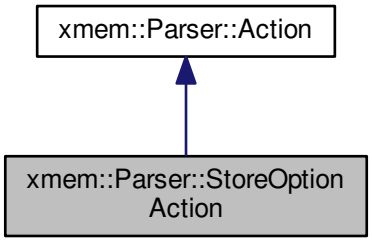


Collaboration diagram for xmem::Parser::StoreOptionAction:

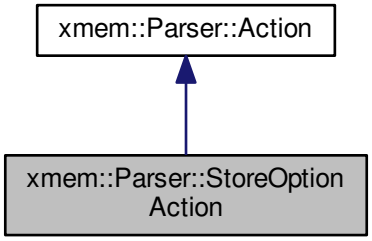


### Public Member Functions

- StoreOptionAction (Parser &parser_, Option options_[], Option buffer_[], int bufmax_)

  *Number of slots in `buffer`. `-1` means "large enough".*
- bool perform (Option &option)

  *Called by Parser::workhorse() for each Option that has been successfully parsed (including unknown options if they have a Descriptor whose Descriptor::check_arg does not return ARG_ILLEGAL.*
- bool finished (int numargs, const char **args)

  *Called by Parser::workhorse() after finishing the parse.*

### 5.25.1 Constructor & Destructor Documentation

#### 5.25.1.1 xmem::Parser::StoreOptionAction::StoreOptionAction ( Parser & *parser_*, Option *options_[]*, Option *buffer_[]*, int *bufmax_* ) `[inline]`

Number of slots in `buffer`. `-1` means "large enough".

Creates a new StoreOption action.

**Parameters**

| | |
|---|---|
| *parser_* | the parser whose op_count should be updated. |
| *options_* | each Option `o` is chained into the linked list `options_`[o.desc->index] |
| *buffer_* | each Option is appended to this array as long as there's a free slot. |
| *bufmax_* | number of slots in `buffer_`. `-1` means "large enough". |

### 5.25.2 Member Function Documentation

#### 5.25.2.1 bool xmem::Parser::StoreOptionAction::finished ( int *numargs*, const char ** *args* ) [inline], [virtual]

Called by Parser::workhorse() after finishing the parse.

**Parameters**

| | |
|---|---|
| *numargs* | the number of non-option arguments remaining |
| *args* | pointer to the first remaining non-option argument (if numargs > 0). |

**Returns**

`false` iff a fatal error has occurred.

Reimplemented from xmem::Parser::Action.

#### 5.25.2.2 bool xmem::Parser::StoreOptionAction::perform ( Option & ) [inline], [virtual]

Called by Parser::workhorse() for each Option that has been successfully parsed (including unknown options if they have a Descriptor whose Descriptor::check_arg does not return ARG_ILLEGAL.

Returns `false` iff a fatal error has occured and the parse should be aborted.

Reimplemented from xmem::Parser::Action.

Here is the call graph for this function:

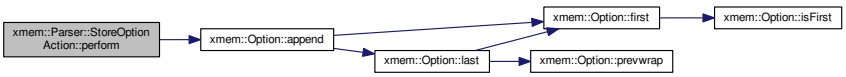


The documentation for this class was generated from the following file:

- src/include/optionparser.h

## 5.26 xmem::PrintUsageImplementation::StreamWriter< Function, Stream > Struct Template Reference

Inheritance diagram for xmem::PrintUsageImplementation::StreamWriter< Function, Stream >:

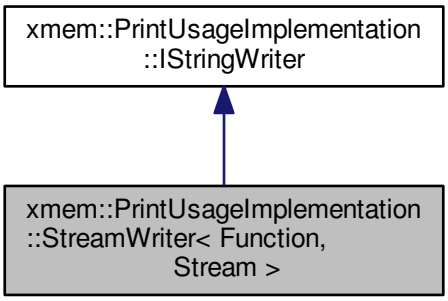


Collaboration diagram for xmem::PrintUsageImplementation::StreamWriter< Function, Stream >:

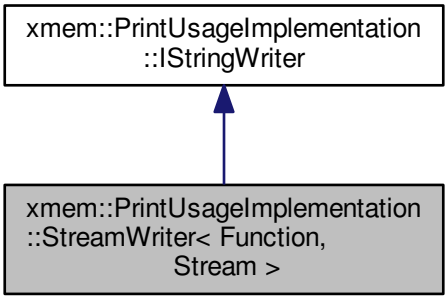


### Public Member Functions

- virtual void operator() (const char *str, int size)

  *Writes the given number of chars beginning at the given pointer somewhere.*
- **StreamWriter** (Function *w, Stream *s)

### Public Attributes

- Function * **fwrite**
- Stream * **stream**

The documentation for this struct was generated from the following file:

- src/include/optionparser.h

## 5.27 xmem::PrintUsageImplementation::SyscallWriter< Syscall > Struct Template Reference

Inheritance diagram for xmem::PrintUsageImplementation::SyscallWriter< Syscall >:

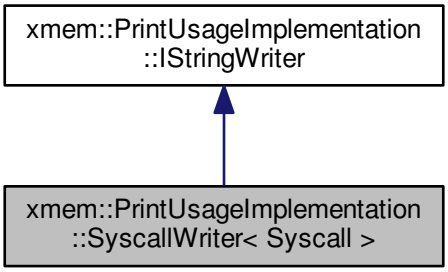


Collaboration diagram for xmem::PrintUsageImplementation::SyscallWriter< Syscall >:

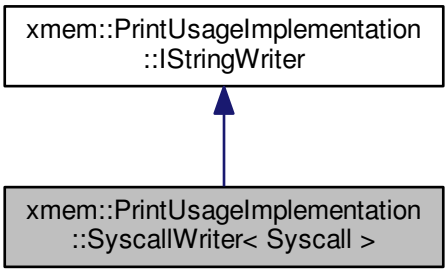


### Public Member Functions

- virtual void operator() (const char *str, int size)

  *Writes the given number of chars beginning at the given pointer somewhere.*
- **SyscallWriter** (Syscall *w, int f)

### Public Attributes

- Syscall * **write**
- int **fd**

The documentation for this struct was generated from the following file:

- src/include/optionparser.h

## 5.28 xmem::PrintUsageImplementation::TemporaryWriter< Temporary > Struct Template Reference

Inheritance diagram for xmem::PrintUsageImplementation::TemporaryWriter< Temporary >:

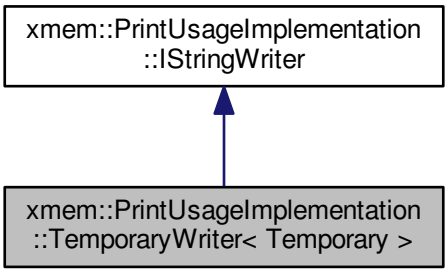


Collaboration diagram for xmem::PrintUsageImplementation::TemporaryWriter< Temporary >:

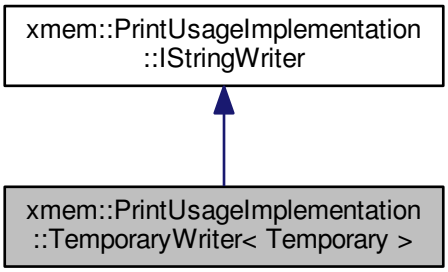


### Public Member Functions

- virtual void operator() (const char *str, int size)

  *Writes the given number of chars beginning at the given pointer somewhere.*
- **TemporaryWriter** (const Temporary &u)

### Public Attributes

- const Temporary & **userstream**

The documentation for this struct was generated from the following file:

- src/include/optionparser.h

## 5.29 xmem::Thread Class Reference

a nice wrapped thread interface independent of particular OS API

`#include <Thread.h>`

### Public Member Functions

- Thread (Runnable *target)
- ~Thread ()
- bool create_and_start ()
- bool join ()
- bool cancel ()
- int32_t getExitCode ()
- bool started ()
- bool completed ()
- bool validTarget ()
- bool created ()
- bool isThreadSuspended ()
- bool isThreadRunning ()
- Runnable * getTarget ()

### 5.29.1 Detailed Description

a nice wrapped thread interface independent of particular OS API

### 5.29.2 Constructor & Destructor Documentation

#### 5.29.2.1 Thread::Thread ( Runnable * *target* )

Constructor. Does not actually create the real thread or run it.

**Parameters**

| | |
|---|---|
| *target* | The target object to do some work with in a new thread. |

#### 5.29.2.2 Thread::~Thread ( )

Destructor. Immediately cancels the thread if it exists. This can be unsafe!

Here is the call graph for this function:

### 5.29.3 Member Function Documentation

#### 5.29.3.1 bool Thread::cancel ( )

Cancels the worker thread immediately. This should only be done in emergencies, as it is effectively killed and undefined behavior might occur.

**Returns**

true if the worker thread was successfully killed.

#### 5.29.3.2 bool Thread::completed ( )

**Returns**

true if the thread completed, regardless of the manner in which it terminated. Returns false if it has not been started.

#### 5.29.3.3 bool Thread::create_and_start ( )

Creates and starts the thread immediately if the target Runnable is valid. This invokes the run() method in the Runnable target that was passed in the constructor.

**Returns**

true if the thread was successfully created and started.

#### 5.29.3.4 bool Thread::created ( )

**Returns**

true if the thread has been created successfully.

#### 5.29.3.5 int32_t Thread::getExitCode ( )

**Returns**

the exit code of the worker thread if it completed. If it did not complete or has not started, returns 0.

#### 5.29.3.6 Runnable * Thread::getTarget ( )

**Returns**

a pointer to the target Runnable object

#### 5.29.3.7 bool Thread::isThreadRunning ( )

**Returns**

true if the thread is running. Returns false if the thread has not been created.

#### 5.29.3.8 bool Thread::isThreadSuspended ( )

**Returns**

true if the thread is suspended. Returns false if the thread has not been created.

#### 5.29.3.9 bool Thread::join ( )

Blocks the calling thread until the worker thread managed by this object terminates. For simplicity, this does not support a timeout due to pthreads incompatibility with the Windows threading API. If the worker thread has already terminated, returns immediately. If the worker has not yet started, returns immediately.

**Returns**

true if the worker thread terminated successfully, false otherwise.

#### 5.29.3.10 bool Thread::started ( )

**Returns**

true if the thread has been started, regardless if has completed or not.

#### 5.29.3.11 bool Thread::validTarget ( )

**Returns**

true if the Runnable target is valid.

The documentation for this class was generated from the following files:

- src/include/Thread.h
- src/Thread.cpp

## 5.30 xmem::ThroughputBenchmark Class Reference

A type of benchmark that measures memory throughput.

`#include <ThroughputBenchmark.h>`

Inheritance diagram for xmem::ThroughputBenchmark:

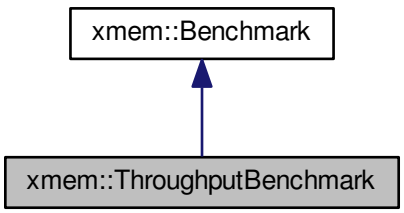

Collaboration diagram for xmem::ThroughputBenchmark:

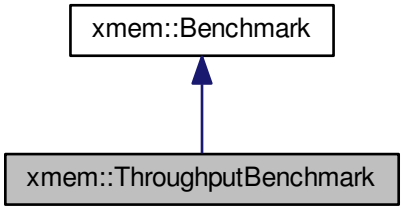


### Public Member Functions

- ThroughputBenchmark (void *mem_array, size_t len, uint32_t iterations, uint32_t num_worker_threads, uint32_t mem_node, uint32_t cpu_node, pattern_mode_t pattern_mode, rw_mode_t rw_mode, chunk_size_t chunk_size, int32_t stride_size, std::vector< PowerReader * > dram_power_readers, std::string name)

  *Constructor. Parameters are passed directly to the Benchmark constructor. See Benchmark class documentation for parameter semantics.*

- virtual ~ThroughputBenchmark ()

  *Destructor.*

### Protected Member Functions

- virtual bool _run_core ()

  *The core benchmark function.*

### Additional Inherited Members

## 5.30.1 Detailed Description

A type of benchmark that measures memory throughput.

## 5.30.2 Member Function Documentation

#### 5.30.2.1 bool ThroughputBenchmark::_run_core ( ) [protected], [virtual]

The core benchmark function.

**Returns**

True on success.

Implements xmem::Benchmark.

Here is the call graph for this function:

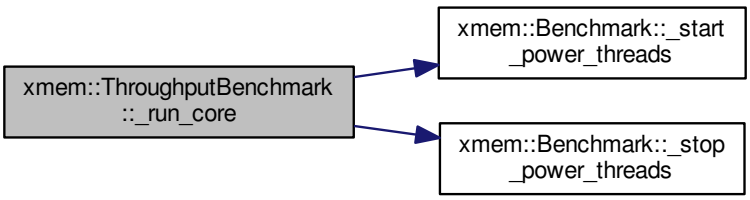


The documentation for this class was generated from the following files:

- src/include/ThroughputBenchmark.h
- src/ThroughputBenchmark.cpp

## 5.31 xmem::Timer Class Reference

This class abstracts some characteristics of simple high resolution stopwatch timer. However, due to the inability or complexity of abstracting shared hardware timers, this class does not actually provide start and stop functions.

`#include <Timer.h>`

### Public Member Functions

- Timer ()

  *Constructor. This may take a noticeable amount of time.*
- tick_t get_ticks_per_ms ()

  *Gets ticks per ms for this timer.*
- float get_ns_per_tick ()

  *Gets nanoseconds per tick for this timer.*

### Protected Attributes

- tick_t _ticks_per_ms
- float _ns_per_tick

### 5.31.1 Detailed Description

This class abstracts some characteristics of simple high resolution stopwatch timer. However, due to the inability or complexity of abstracting shared hardware timers, this class does not actually provide start and stop functions.

### 5.31.2 Member Function Documentation

#### 5.31.2.1 float Timer::get_ns_per_tick ( )

Gets nanoseconds per tick for this timer.

**Returns**

the number of nanoseconds per tick

#### 5.31.2.2 tick_t Timer::get_ticks_per_ms ( )

Gets ticks per ms for this timer.

**Returns**

The reported number of ticks per ms.

### 5.31.3 Member Data Documentation

#### 5.31.3.1 float xmem::Timer::_ns_per_tick [protected]

Nanoseconds per tick for this timer.

#### 5.31.3.2 tick_t xmem::Timer::_ticks_per_ms [protected]

Ticks per ms for this timer.

The documentation for this class was generated from the following files:

- src/include/Timer.h
- src/Timer.cpp

# Chapter 6

# File Documentation

## 6.1 src/Benchmark.cpp File Reference

Implementation file for the Benchmark class.

```
#include <Benchmark.h>
#include <common.h>
#include <benchmark_kernels.h>
#include <PowerReader.h>
#include <cstdint>
#include <iostream>
#include <vector>
#include <time.h>
```

Include dependency graph for Benchmark.cpp:

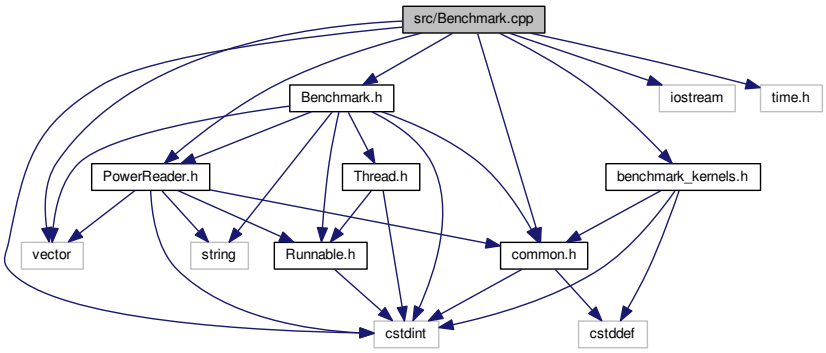


### 6.1.1 Detailed Description

Implementation file for the Benchmark class.

## 6.2 src/benchmark_kernels.cpp File Reference

Implementation file for benchmark kernel functions for doing the actual work we care about. :)

```
#include <benchmark_kernels.h>
#include <common.h>
#include <iostream>
#include <random>
#include <algorithm>
#include <time.h>
```

Include dependency graph for benchmark_kernels.cpp:

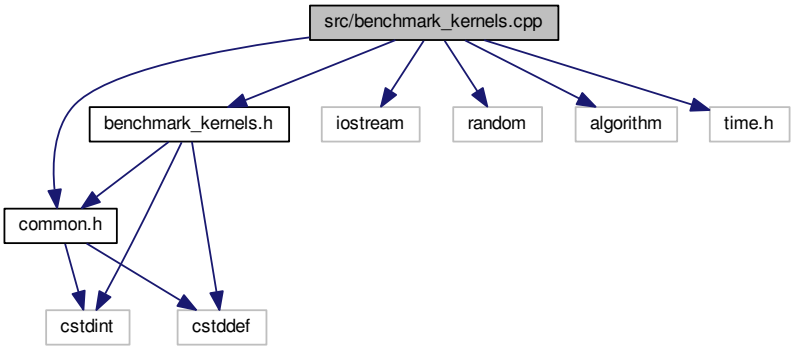


### 6.2.1 Detailed Description

Implementation file for benchmark kernel functions for doing the actual work we care about. :) Optimization tricks include:

- UNROLL macros to manual loop unrolling. This reduces the relative branch overhead of the loop. We don't want to benchmark loops, we want to benchmark memory! But unrolling too much can hurt code size and instruction locality, potentially decreasing I-cache utilization and causing extra overheads. This is why we allow multiple unroll lengths at compile-time.

- volatile keyword to prevent compiler from optimizing the code and removing instructions that we need. The compiler is too smart for its own good!

## 6.3 src/BenchmarkManager.cpp File Reference

Implementation file for the BenchmarkManager class.

```
#include <BenchmarkManager.h>
#include <common.h>
#include <Configurator.h>
#include <cstdint>
#include <stdlib.h>
#include <iostream>
#include <sstream>
#include <assert.h>
```

Include dependency graph for BenchmarkManager.cpp:

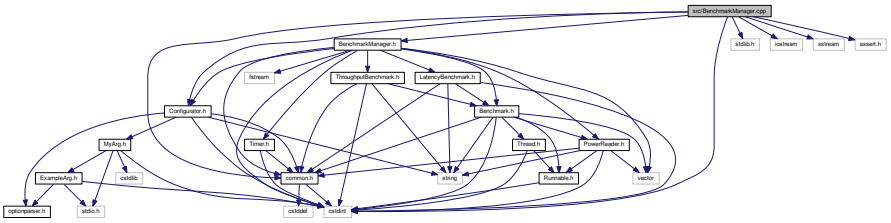


### 6.3.1 Detailed Description

Implementation file for the BenchmarkManager class.

## 6.4 src/common.cpp File Reference

Implementation file for common preprocessor definitions, macros, functions, and global constants.

```
#include <common.h>
#include <Timer.h>
#include <iostream>
```

Include dependency graph for common.cpp:

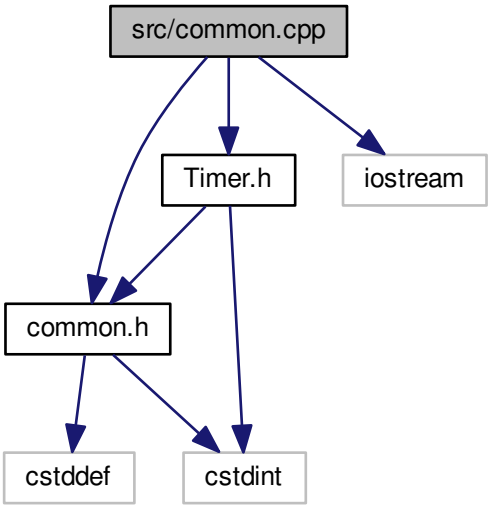


### Variables

- bool **xmem::g_verbose** = false

  ∗*typedef union {*
- size_t **xmem::g_page_size**

- size_t **xmem::g_large_page_size**
- uint32_t **xmem::g_num_nodes**
- uint32_t **xmem::g_num_logical_cpus**
- uint32_t **xmem::g_num_physical_cpus**
- uint32_t **xmem::g_num_physical_packages**
- uint32_t **xmem::g_total_l1_caches**
- uint32_t **xmem::g_total_l2_caches**
- uint32_t **xmem::g_total_l3_caches**
- uint32_t **xmem::g_total_l4_caches**
- uint32_t **xmem::g_starting_test_index**
- uint32_t **xmem::g_test_index**
- tick_t **xmem::g_ticks_per_ms**
- float **xmem::g_ns_per_tick**

### 6.4.1 Detailed Description

Implementation file for common preprocessor definitions, macros, functions, and global constants.

## 6.5 src/Configurator.cpp File Reference

Implementation file for the Configurator class and some helper data structures.

```
#include <Configurator.h>
#include <common.h>
#include <optionparser.h>
#include <MyArg.h>
#include <cstdint>
#include <iostream>
#include <string>
```

Include dependency graph for Configurator.cpp:

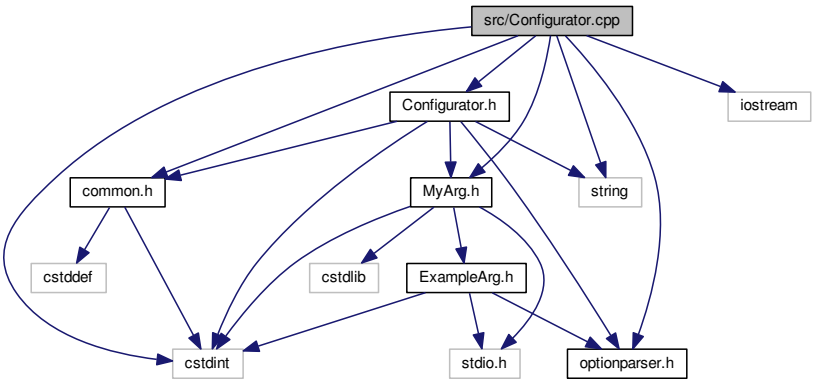


### 6.5.1 Detailed Description

Implementation file for the Configurator class and some helper data structures.

## 6.6 src/ext/DelayInjectedLoadedLatencyBenchmark/delay_injected_benchmark_kernels.cpp File Reference

Implementation file for benchmark kernel functions for the delay-injected loaded latency benchmark.

```
#include <delay_injected_benchmark_kernels.h>
#include <common.h>
#include <iostream>
```

Include dependency graph for delay_injected_benchmark_kernels.cpp:

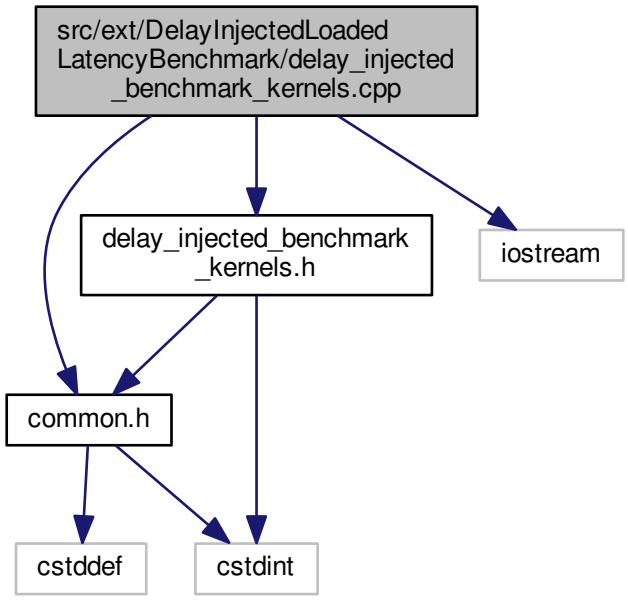


### 6.6.1 Detailed Description

Implementation file for benchmark kernel functions for the delay-injected loaded latency benchmark. Optimization tricks include:

- UNROLL macros to manual loop unrolling. This reduces the relative branch overhead of the loop. We don't want to benchmark loops, we want to benchmark memory! But unrolling too much can hurt code size and instruction locality, potentially decreasing I-cache utilization and causing extra overheads. This is why we allow multiple unroll lengths at compile-time.
- volatile keyword to prevent compiler from optimizing the code and removing instructions that we need. The compiler is too smart for its own good!

## 6.7 src/ext/DelayInjectedLoadedLatencyBenchmark/DelayInjectedLoadedLatencyBenchmark.cpp File Reference

Implementation file for the DelayInjectedLatencyBenchmark class.

```
#include <common.h>
```

Include dependency graph for DelayInjectedLoadedLatencyBenchmark.cpp:

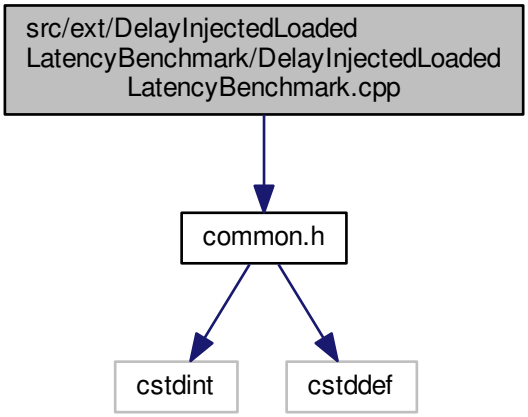


### 6.7.1 Detailed Description

Implementation file for the DelayInjectedLatencyBenchmark class.

## 6.8 src/include/Benchmark.h File Reference

Header file for the Benchmark class.

```
#include <common.h>
#include <PowerReader.h>
#include <Thread.h>
#include <Runnable.h>
#include <cstdint>
#include <string>
#include <vector>
```

Include dependency graph for Benchmark.h:

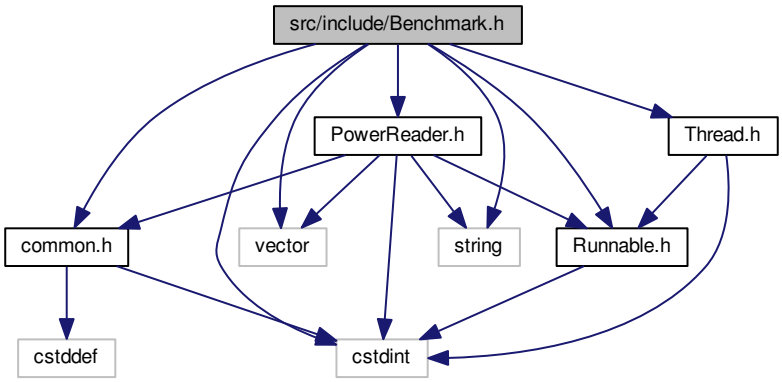


This graph shows which files directly or indirectly include this file:

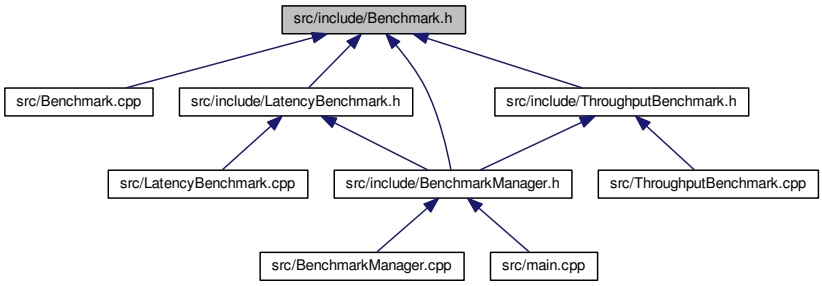


### Classes

- class xmem::Benchmark

  *Flexible abstract class for any memory benchmark.*

### 6.8.1 Detailed Description

Header file for the Benchmark class.

## 6.9 src/include/benchmark_kernels.h File Reference

Header file for benchmark kernel functions for doing the actual work we care about. :)

```
#include <common.h>
#include <cstdint>
#include <cstddef>
```

Include dependency graph for benchmark_kernels.h:

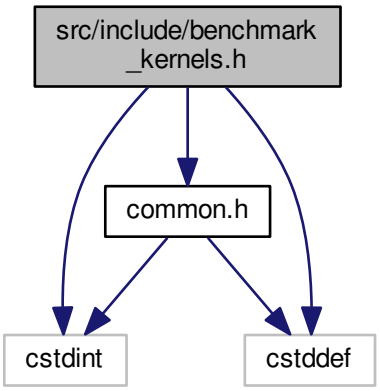


This graph shows which files directly or indirectly include this file:

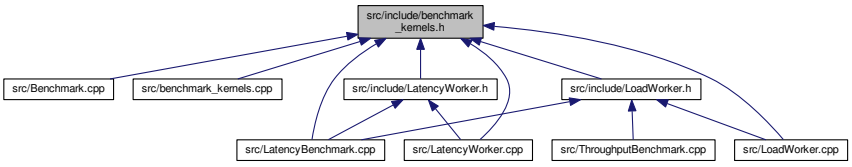


## Typedefs

- typedef int32_t(∗ **xmem::SequentialFunction** )(void ∗, void ∗)
- typedef int32_t(∗ **xmem::RandomFunction** )(uintptr_t ∗, uintptr_t ∗∗, size_t)

## Functions

- bool **xmem::determineSequentialKernel** (rw_mode_t rw_mode, chunk_size_t chunk_size, int32_t stride_size, SequentialFunction ∗kernel_function, SequentialFunction ∗dummy_kernel_function)

  *Determines which sequential memory access kernel to use based on the read/write mode, chunk size, and stride size.*
- bool **xmem::determineRandomKernel** (rw_mode_t rw_mode, chunk_size_t chunk_size, RandomFunction ∗kernel_function, RandomFunction ∗dummy_kernel_function)

  *Determines which random memory access kernel to use based on the read/write mode, chunk size, and stride size.*
- bool **xmem::buildRandomPointerPermutation** (void ∗start_address, void ∗end_address, chunk_size_t chunk_size)

  *Builds a random chain of pointers within the specified memory region.*
- int32_t **xmem::dummy_chasePointers** (uintptr_t ∗, uintptr_t ∗∗, size_t len)

  *Mimics the chasePointers() method but doesn't do the memory accesses.*
- int32_t **xmem::chasePointers** (uintptr_t ∗first_address, uintptr_t ∗∗last_touched_address, size_t len)

*Walks over the allocated memory in random order by chasing pointers.*

- int32_t **xmem::dummy_empty** (void *, void *)

  *Does nothing. Used for measuring the time it takes just to call a benchmark routine via function pointer.*

- int32_t **xmem::dummy_forwSequentialLoop_Word32** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in forward sequential Word 32 loops except for the memory access itself.*

- int32_t **xmem::dummy_revSequentialLoop_Word32** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in reverse sequential Word 32 loops except for the memory access itself.*

- int32_t **xmem::dummy_forwStride2Loop_Word32** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in forward 2-strided Word 32 loops except for the memory access itself.*

- int32_t **xmem::dummy_revStride2Loop_Word32** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in reverse 2-strided Word 32 loops except for the memory access itself.*

- int32_t **xmem::dummy_forwStride4Loop_Word32** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in forward 4-strided Word 32 loops except for the memory access itself.*

- int32_t **xmem::dummy_revStride4Loop_Word32** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in reverse 4-strided Word 32 loops except for the memory access itself.*

- int32_t **xmem::dummy_forwStride8Loop_Word32** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in forward 8-strided Word 32 loops except for the memory access itself.*

- int32_t **xmem::dummy_revStride8Loop_Word32** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in reverse 8-strided Word 32 loops except for the memory access itself.*

- int32_t **xmem::dummy_forwStride16Loop_Word32** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in forward 16-strided Word 32 loops except for the memory access itself.*

- int32_t **xmem::dummy_revStride16Loop_Word32** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in reverse 16-strided Word 32 loops except for the memory access itself.*

- int32_t **xmem::dummy_randomLoop_Word32** (uintptr_t *, uintptr_t **, size_t len)

  *Mimics the randomRead_Word32 and randomWrite_Word32 functions except for the memory accesses.*

- int32_t **xmem::forwSequentialRead_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory forward sequentially, reading in 32-bit chunks.*

- int32_t **xmem::revSequentialRead_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory reverse sequentially, reading in 32-bit chunks.*

- int32_t **xmem::forwSequentialWrite_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory forward sequentially, writing in 32-bit chunks.*

- int32_t **xmem::revSequentialWrite_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory reverse sequentially, writing in 32-bit chunks.*

- int32_t **xmem::forwStride2Read_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in forward strides of size 2, reading in 32-bit chunks.*

- int32_t **xmem::revStride2Read_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in reverse strides of size 2, reading in 32-bit chunks.*

- int32_t **xmem::forwStride2Write_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in forward strides of size 2, writing in 32-bit chunks.*

- int32_t **xmem::revStride2Write_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in reverse strides of size 2, writing in 32-bit chunks.*

- int32_t **xmem::forwStride4Read_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in forward strides of size 4, reading in 32-bit chunks.*

- int32_t **xmem::revStride4Read_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in reverse strides of size 4, reading in 32-bit chunks.*
- int32_t **xmem::forwStride4Write_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in forward strides of size 4, writing in 32-bit chunks.*
- int32_t **xmem::revStride4Write_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in reverse strides of size 4, writing in 32-bit chunks.*
- int32_t **xmem::forwStride8Read_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in forward strides of size 8, reading in 32-bit chunks.*
- int32_t **xmem::revStride8Read_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in reverse strides of size 8, reading in 32-bit chunks.*
- int32_t **xmem::forwStride8Write_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in forward strides of size 8, writing in 32-bit chunks.*
- int32_t **xmem::revStride8Write_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in reverse strides of size 8, writing in 32-bit chunks.*
- int32_t **xmem::forwStride16Read_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in forward strides of size 16, reading in 32-bit chunks.*
- int32_t **xmem::revStride16Read_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in reverse strides of size 16, reading in 32-bit chunks.*
- int32_t **xmem::forwStride16Write_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in forward strides of size 16, writing in 32-bit chunks.*
- int32_t **xmem::revStride16Write_Word32** (void *start_address, void *end_address)

  *Walks over the allocated memory in reverse strides of size 16, writing in 32-bit chunks.*
- int32_t **xmem::randomRead_Word32** (uintptr_t *first_address, uintptr_t **last_touched_address, size_t len)

  *Walks over the allocated memory in random order by chasing 64-bit pointers.*
- int32_t **xmem::randomWrite_Word32** (uintptr_t *first_address, uintptr_t **last_touched_address, size_t len)

  *Walks over the allocated memory in random order by chasing 32-bit pointers. A pointer is read and written back with the same value before chasing to the next pointer. Thus, each memory address is a read followed by immediate write operation.*

### 6.9.1 Detailed Description

Header file for benchmark kernel functions for doing the actual work we care about. :)

## 6.10 src/include/BenchmarkManager.h File Reference

Header file for the BenchmarkManager class.

```
#include <common.h>
#include <Timer.h>
#include <PowerReader.h>
#include <Benchmark.h>
#include <ThroughputBenchmark.h>
#include <LatencyBenchmark.h>
#include <Configurator.h>
#include <cstdint>
#include <vector>
#include <fstream>
```

Include dependency graph for BenchmarkManager.h:

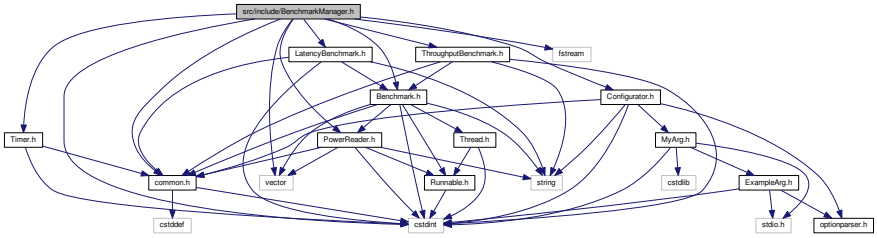


This graph shows which files directly or indirectly include this file:

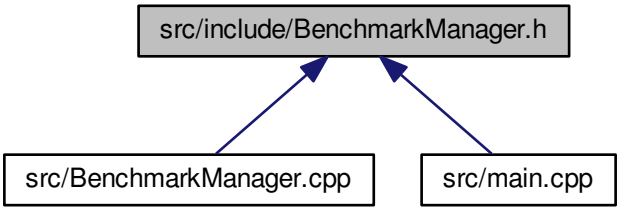


### Classes

- class xmem::BenchmarkManager

  *Manages running all benchmarks at a high level.*

### 6.10.1 Detailed Description

Header file for the BenchmarkManager class.

## 6.11 src/include/common.h File Reference

Header file for common preprocessor definitions, macros, functions, and global constants.

```
#include <cstdint>
#include <cstddef>
```

Include dependency graph for common.h:

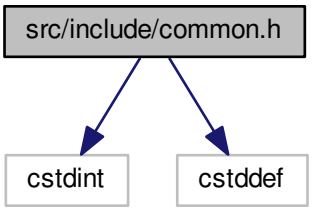


This graph shows which files directly or indirectly include this file:

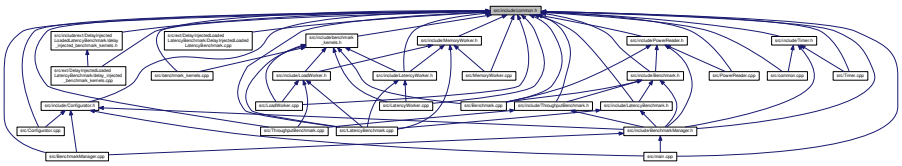

## Macros

- #define **VERSION** "2.2.7"
- #define **KB** 1024
- #define **MB** 1048576
- #define **MB_4** 4194304
- #define **MB_16** 16777216
- #define **MB_64** 67108864
- #define **MB_256** 268435456
- #define **MB_512** 536870912
- #define **GB** 1073741824
- #define DEFAULT_PAGE_SIZE 4∗KB
- #define DEFAULT_LARGE_PAGE_SIZE 2∗MB
- #define DEFAULT_WORKING_SET_SIZE_PER_THREAD DEFAULT_PAGE_SIZE
- #define DEFAULT_NUM_WORKER_THREADS 1
- #define DEFAULT_NUM_NODES 0
- #define DEFAULT_NUM_PHYSICAL_PACKAGES 0
- #define DEFAULT_NUM_PHYSICAL_CPUS 0
- #define DEFAULT_NUM_LOGICAL_CPUS 0
- #define DEFAULT_NUM_L1_CACHES 0
- #define DEFAULT_NUM_L2_CACHES 0
- #define DEFAULT_NUM_L3_CACHES 0
- #define DEFAULT_NUM_L4_CACHES 0
- #define MIN_ELAPSED_TICKS 10000
- #define **UNROLL2**(x) x x
- #define **UNROLL4**(x) UNROLL2(x) UNROLL2(x)
- #define **UNROLL8**(x) UNROLL4(x) UNROLL4(x)
- #define **UNROLL16**(x) UNROLL8(x) UNROLL8(x)

- #define **UNROLL32**(x) UNROLL16(x) UNROLL16(x)
- #define **UNROLL64**(x) UNROLL32(x) UNROLL32(x)
- #define **UNROLL128**(x) UNROLL64(x) UNROLL64(x)
- #define **UNROLL256**(x) UNROLL128(x) UNROLL128(x)
- #define **UNROLL512**(x) UNROLL256(x) UNROLL256(x)
- #define **UNROLL1024**(x) UNROLL512(x) UNROLL512(x)
- #define **UNROLL2048**(x) UNROLL1024(x) UNROLL1024(x)
- #define **UNROLL4096**(x) UNROLL2048(x) UNROLL2048(x)
- #define **UNROLL8192**(x) UNROLL4096(x) UNROLL4096(x)
- #define **UNROLL16384**(x) UNROLL8192(x) UNROLL8192(x)
- #define **UNROLL32768**(x) UNROLL16384(x) UNROLL16384(x)
- #define **UNROLL65536**(x) UNROLL32768(x) UNROLL32768(x)
- #define LATENCY_BENCHMARK_UNROLL_LENGTH 512
- #define USE_OS_TIMER
- #define BENCHMARK_DURATION_MS 250
- #define THROUGHPUT_BENCHMARK_BYTES_PER_PASS 4096
- #define POWER_SAMPLING_PERIOD_MS 1000
- #define EXT_DELAY_INJECTED_LOADED_LATENCY_BENCHMARK

## Typedefs

- typedef uint64_t **xmem::tick_t**
- typedef uint32_t **xmem::Word32_t**

## Enumerations

- enum **pattern_mode_t** { **SEQUENTIAL**, **RANDOM**, **NUM_PATTERN_MODES** }

  *Memory access patterns are broadly categorized by sequential or random-access.*
- enum **rw_mode_t** { **READ**, **WRITE**, **NUM_RW_MODES** }

  *Memory access batterns are broadly categorized by reads and writes.*
- enum **chunk_size_t** { **CHUNK_32b**, **NUM_CHUNK_SIZES** }

  *Legal memory read/write chunk sizes in bits.*
- enum **ext_t** { **EXT_NUM_DELAY_INJECTED_LOADED_LATENCY_BENCHMARK**, **NUM_EXTENSIONS** }

## Functions

- void **xmem::print_welcome_message** ()

  *Prints a basic welcome message to the console with useful information.*
- void **xmem::print_types_report** ()

  *Prints the various C/C++ types to the console for this machine.*
- void **xmem::print_compile_time_options** ()

  *Prints compile-time option information to the console.*
- void **xmem::setup_timer** ()

  *Initializes the timer and outputs results to the console for sanity checking.*
- void **xmem::report_timer** ()

  *Reports timer info to the console.*
- void **xmem::test_thread_affinities** ()

  *Checks to see if the calling thread can be locked to all logical CPUs in the system, and reports to the console the progress.*
- bool **xmem::lock_thread_to_numa_node** (uint32_t numa_node)

  *Sets the affinity of the calling thread to the lowest numbered logical CPU in the given NUMA node. TODO: Improve this functionality, it is quite limiting.*

- bool **xmem::unlock_thread_to_numa_node** ()

  *Clears the affinity of the calling thread to any given NUMA node.*
- bool **xmem::lock_thread_to_cpu** (uint32_t cpu_id)

  *Sets the affinity of the calling thread to a given logical CPU.*
- bool **xmem::unlock_thread_to_cpu** ()

  *Clears the affinity of the calling thread to any given logical CPU.*
- int32_t **xmem::cpu_id_in_numa_node** (uint32_t numa_node, uint32_t cpu_in_node)

  *Gets the CPU ID for a logical CPU of interest in a particular NUMA node. For example, if numa_node is 1 and cpu_in_node is 2, and there are 4 logical CPUs per node, then this will give the answer 6 (6th CPU), assuming CPU IDs start at 0.*
- size_t **xmem::compute_number_of_passes** (size_t working_set_size_KB)

  *Computes the number of passes to use for a given working set size in KB, when size-based benchmarking mode is enabled at compile-time. You may want to change this implementation to suit your needs. See the compile-time options in common.h.*
- bool **xmem::config_page_size** ()

  *Queries the page sizes from the system and sets relevant global variables.*
- void **xmem::init_globals** ()

  *Initializes useful global variables.*
- int32_t **xmem::query_sys_info** ()

  *Sets up global variables based on system information at runtime.*
- void **xmem::report_sys_info** ()

  *Reports the system configuration to the console as indicated by global variables.*
- tick_t **xmem::start_timer** ()

  *Query the timer for the start of a timed section of code.*
- tick_t **xmem::stop_timer** ()

  *Query the timer for the end of a timed section of code.*

### 6.11.1 Detailed Description

Header file for common preprocessor definitions, macros, functions, and global constants.

### 6.11.2 Macro Definition Documentation

#### 6.11.2.1 #define BENCHMARK_DURATION_MS 250

RECOMMENDED VALUE: At least 250. Number of milliseconds to run in each benchmark.

#### 6.11.2.2 #define DEFAULT_LARGE_PAGE_SIZE 2∗MB

Default platform large page size in bytes. This generally should not be relied on, but is a failsafe.

#### 6.11.2.3 #define DEFAULT_NUM_L1_CACHES 0

Default number of L1 caches.

#### 6.11.2.4 #define DEFAULT_NUM_L2_CACHES 0

Default number of L2 caches.

#### 6.11.2.5 #define DEFAULT_NUM_L3_CACHES 0

Default number of L3 caches.

#### 6.11.2.6 #define DEFAULT_NUM_L4_CACHES 0

Default number of L4 caches.

#### 6.11.2.7 #define DEFAULT_NUM_LOGICAL_CPUS 0

Default number of logical CPU cores.

#### 6.11.2.8 #define DEFAULT_NUM_NODES 0

Default number of NUMA nodes.

#### 6.11.2.9 #define DEFAULT_NUM_PHYSICAL_CPUS 0

Default number of physical CPU cores.

#### 6.11.2.10 #define DEFAULT_NUM_PHYSICAL_PACKAGES 0

Default number of physical packages.

#### 6.11.2.11 #define DEFAULT_NUM_WORKER_THREADS 1

Default number of worker threads to use.

#### 6.11.2.12 #define DEFAULT_PAGE_SIZE 4∗KB

Default platform page size in bytes. This generally should not be relied on, but is a failsafe.

#### 6.11.2.13 #define DEFAULT_WORKING_SET_SIZE_PER_THREAD DEFAULT_PAGE_SIZE

Default working set size in bytes.

#### 6.11.2.14 #define EXT_DELAY_INJECTED_LOADED_LATENCY_BENCHMARK

RECOMMENDED ENABLED. This allows for a custom extension to X-Mem that performs latency benchmarking with forward sequential 64-bit and 256-bit read-based load threads with variable delays injected in between memory accesses.

#### 6.11.2.15 #define LATENCY_BENCHMARK_UNROLL_LENGTH 512

Number of unrolls in the latency benchmark pointer chasing core function.

#### 6.11.2.16 #define MIN_ELAPSED_TICKS 10000

If any routine measured fewer than this number of ticks its results should be viewed with suspicion. This is because the latency of the timer itself will matter.

#### 6.11.2.17 #define POWER_SAMPLING_PERIOD_MS 1000

RECOMMENDED VALUE: 1000. Sampling period in milliseconds for all power measurement mechanisms.

#### 6.11.2.18 #define THROUGHPUT_BENCHMARK_BYTES_PER_PASS 4096

RECOMMENDED VALUE: 4096. Number of bytes read or written per pass of any ThroughputBenchmark. This must be less than or equal to the minimum working set size, which is currently 4 KB.

#### 6.11.2.19 #define USE_OS_TIMER

RECOMMENDED ENABLED. If enabled, uses the QPC timer on Windows and the POSIX clock_gettime() on GNU/Linux for all timing purposes.

## 6.12 src/include/Configurator.h File Reference

Header file for the Configurator class and some helper data structures.

```
#include <common.h>
#include <optionparser.h>
#include <MyArg.h>
#include <cstdint>
#include <string>
```

Include dependency graph for Configurator.h:

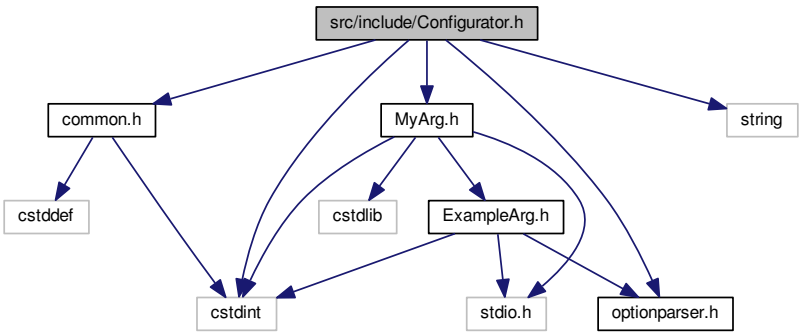

This graph shows which files directly or indirectly include this file:

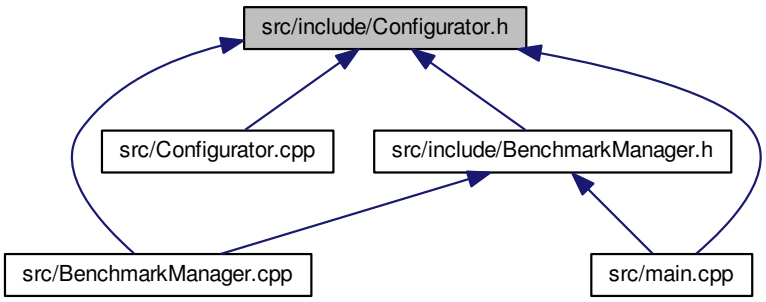


### Classes

- class xmem::Configurator

  *Handles all user input interpretation and generates the necessary flags for running benchmarks.*

### Enumerations

- enum **optionIndex** {
  **UNKNOWN**, **ALL**, **CHUNK_SIZE**, **EXTENSION**,
  **OUTPUT_FILE**, **HELP**, **BASE_TEST_INDEX**, **NUM_WORKER_THREADS**,
  **MEAS_LATENCY**, **ITERATIONS**, **RANDOM_ACCESS_PATTERN**, **SEQUENTIAL_ACCESS_PATTERN**,
  **MEAS_THROUGHPUT**, **NUMA_DISABLE**, **VERBOSE**, **WORKING_SET_SIZE_PER_THREAD**,
  **USE_LARGE_PAGES**, **USE_READS**, **USE_WRITES**, **STRIDE_SIZE** }

  *Enumerates all possible types of command-line options.*

### Variables

- const Descriptor **xmem::usage** []

  *Command-line option descriptors as needed by stuff in optionparser.h>. This is basically the help message content.*

#### 6.12.1 Detailed Description

Header file for the Configurator class and some helper data structures.

## 6.13 src/include/ExampleArg.h File Reference

Slightly-modified third-party code related to OptionParser.

```
#include <optionparser.h>
#include <cstdint>
#include <stdio.h>
```

Include dependency graph for ExampleArg.h:

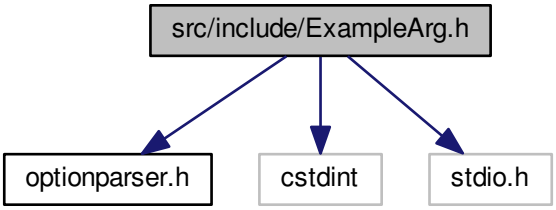


This graph shows which files directly or indirectly include this file:

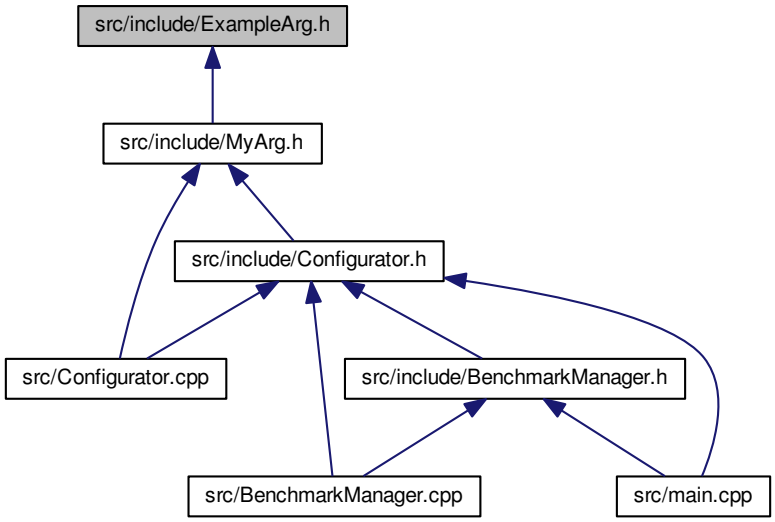


### Classes

- class xmem::ExampleArg

#### 6.13.1 Detailed Description

Slightly-modified third-party code related to OptionParser.

## 6.14 src/include/ext/DelayInjectedLoadedLatencyBenchmark/delay_injected_benchmark_kernels.h File Reference

Header file for benchmark kernel functions with integrated delays for doing the actual work we care about. :)

```
#include <common.h>
#include <cstdint>
```

Include dependency graph for delay_injected_benchmark_kernels.h:

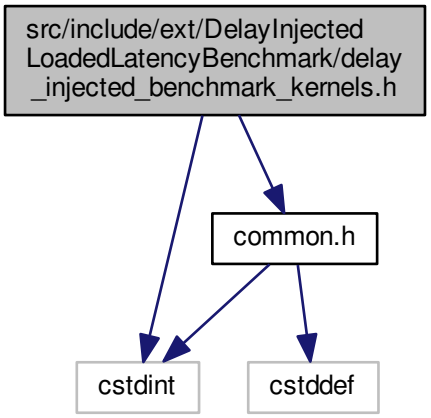


This graph shows which files directly or indirectly include this file:

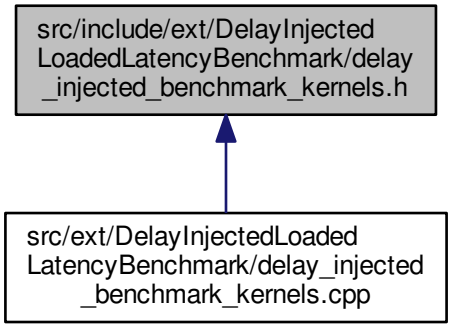


## Macros

- #define **my_nop2**() my_nop(); my_nop()
- #define **my_nop4**() my_nop2(); my_nop2()
- #define **my_nop8**() my_nop4(); my_nop4()
- #define **my_nop16**() my_nop8(); my_nop8()
- #define **my_nop32**() my_nop16(); my_nop16()
- #define **my_nop64**() my_nop32(); my_nop32()
- #define **my_nop128**() my_nop64(); my_nop64()
- #define **my_nop256**() my_nop128(); my_nop128()

- #define **my_nop512**() my_nop256(); my_nop256()
- #define **my_nop1024**() my_nop512(); my_nop512()

## Functions

- int32_t **xmem::dummy_forwSequentialLoop_Word32_Delay1** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in delay-injected forward sequential Word 32 loops except for the memory access and delays themselves.*
- int32_t **xmem::dummy_forwSequentialLoop_Word32_Delay2** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in delay-injected forward sequential Word 32 loops except for the memory access and delays themselves.*
- int32_t **xmem::dummy_forwSequentialLoop_Word32_Delay4** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in delay-injected forward sequential Word 32 loops except for the memory access and delays themselves.*
- int32_t **xmem::dummy_forwSequentialLoop_Word32_Delay8** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in delay-injected forward sequential Word 32 loops except for the memory access and delays themselves.*
- int32_t **xmem::dummy_forwSequentialLoop_Word32_Delay16** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in delay-injected forward sequential Word 32 loops except for the memory access and delays themselves.*
- int32_t **xmem::dummy_forwSequentialLoop_Word32_Delay32** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in delay-injected forward sequential Word 32 loops except for the memory access and delays themselves.*
- int32_t **xmem::dummy_forwSequentialLoop_Word32_Delay64** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in delay-injected forward sequential Word 32 loops except for the memory access and delays themselves.*
- int32_t **xmem::dummy_forwSequentialLoop_Word32_Delay128** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in delay-injected forward sequential Word 32 loops except for the memory access and delays themselves.*
- int32_t **xmem::dummy_forwSequentialLoop_Word32_Delay256** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in delay-injected forward sequential Word 32 loops except for the memory access and delays themselves.*
- int32_t **xmem::dummy_forwSequentialLoop_Word32_Delay512plus** (void *start_address, void *end_address)

  *Used for measuring the time spent doing everything in delay-injected forward sequential Word 32 loops except for the memory access and delays themselves.*
- int32_t **xmem::forwSequentialRead_Word32_Delay1** (void *start_address, void *end_address)

  *Walks over the allocated memory forward sequentially, reading in 32-bit chunks. 1 delays (nops) are inserted between memory instructions.*
- int32_t **xmem::forwSequentialRead_Word32_Delay2** (void *start_address, void *end_address)

  *Walks over the allocated memory forward sequentially, reading in 32-bit chunks. 2 delays (nops) are inserted between memory instructions.*
- int32_t **xmem::forwSequentialRead_Word32_Delay4** (void *start_address, void *end_address)

  *Walks over the allocated memory forward sequentially, reading in 32-bit chunks. 4 delays (nops) are inserted between memory instructions.*
- int32_t **xmem::forwSequentialRead_Word32_Delay8** (void *start_address, void *end_address)

  *Walks over the allocated memory forward sequentially, reading in 32-bit chunks. 8 delays (nops) are inserted between memory instructions.*
- int32_t **xmem::forwSequentialRead_Word32_Delay16** (void *start_address, void *end_address)

  *Walks over the allocated memory forward sequentially, reading in 32-bit chunks. 16 delays (nops) are inserted between memory instructions.*
- int32_t **xmem::forwSequentialRead_Word32_Delay32** (void *start_address, void *end_address)

*Walks over the allocated memory forward sequentially, reading in 32-bit chunks. 32 delays (nops) are inserted between memory instructions.*

- int32_t **xmem::forwSequentialRead_Word32_Delay64** (void *start_address, void *end_address)

  *Walks over the allocated memory forward sequentially, reading in 32-bit chunks. 64 delays (nops) are inserted between memory instructions.*

- int32_t **xmem::forwSequentialRead_Word32_Delay128** (void *start_address, void *end_address)

  *Walks over the allocated memory forward sequentially, reading in 32-bit chunks. 128 delays (nops) are inserted between memory instructions.*

- int32_t **xmem::forwSequentialRead_Word32_Delay256** (void *start_address, void *end_address)

  *Walks over the allocated memory forward sequentially, reading in 32-bit chunks. 256 delays (nops) are inserted between memory instructions.*

- int32_t **xmem::forwSequentialRead_Word32_Delay512** (void *start_address, void *end_address)

  *Walks over the allocated memory forward sequentially, reading in 32-bit chunks. 512 delays (nops) are inserted between memory instructions.*

- int32_t **xmem::forwSequentialRead_Word32_Delay1024** (void *start_address, void *end_address)

  *Walks over the allocated memory forward sequentially, reading in 32-bit chunks. 1024 delays (nops) are inserted between memory instructions.*

### 6.14.1 Detailed Description

Header file for benchmark kernel functions with integrated delays for doing the actual work we care about. :)

## 6.15 src/include/ext/DelayInjectedLoadedLatencyBenchmark/DelayInjectedLoaded-LatencyBenchmark.h File Reference

Header file for the DelayInjectedLoadedLatencyBenchmark class.

### 6.15.1 Detailed Description

Header file for the DelayInjectedLoadedLatencyBenchmark class.

## 6.16 src/include/LatencyBenchmark.h File Reference

Header file for the LatencyBenchmark class.

```
#include <Benchmark.h>
#include <common.h>
#include <cstdint>
#include <string>
```

Include dependency graph for LatencyBenchmark.h:

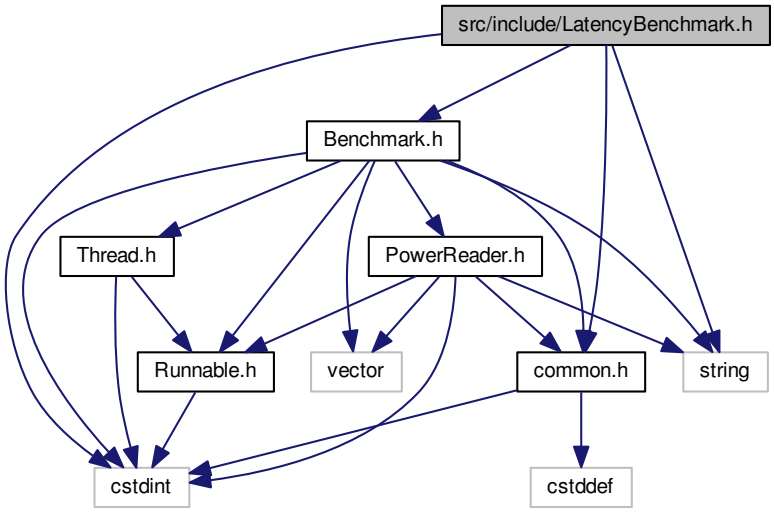


This graph shows which files directly or indirectly include this file:

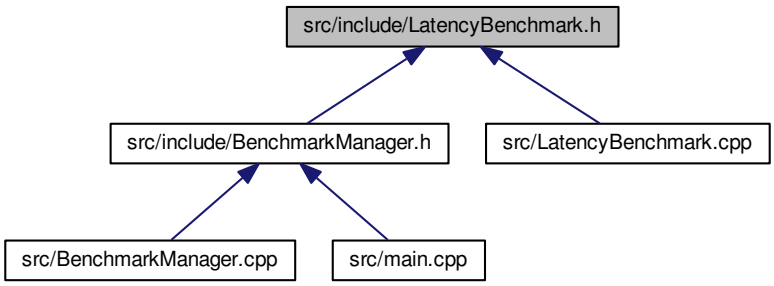


## Classes

- class xmem::LatencyBenchmark

  *A type of benchmark that measures memory latency via random pointer chasing. Loading may be provided with separate threads which access memory as quickly as possible using given access patterns.*

### 6.16.1 Detailed Description

Header file for the LatencyBenchmark class.

## 6.17 src/include/LatencyWorker.h File Reference

Header file for the LatencyWorker class.

```
#include <MemoryWorker.h>
#include <benchmark_kernels.h>
#include <common.h>
```

Include dependency graph for LatencyWorker.h:

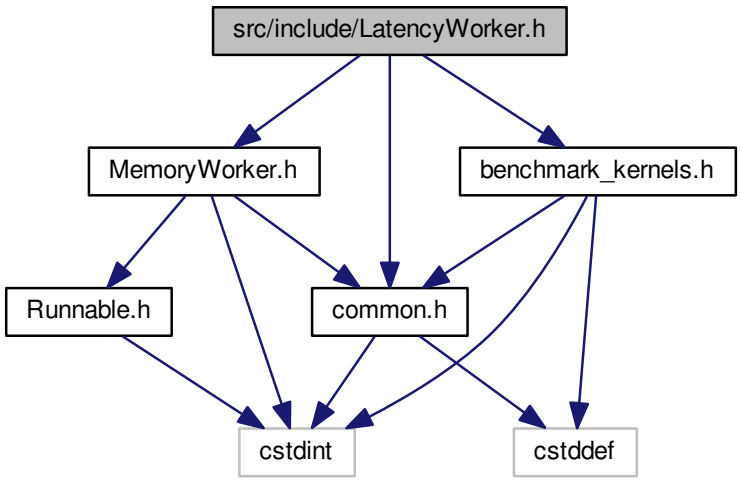


This graph shows which files directly or indirectly include this file:

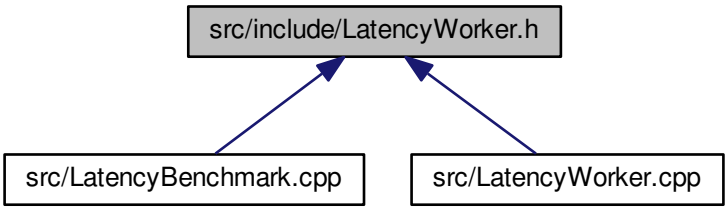


### Classes

- class xmem::LatencyWorker

  *Multithreading-friendly class to do memory loading.*

### 6.17.1 Detailed Description

Header file for the LatencyWorker class.

## 6.18 src/include/LoadWorker.h File Reference

Header file for the LoadWorker class.

```
#include <MemoryWorker.h>
#include <benchmark_kernels.h>
```

Include dependency graph for LoadWorker.h:

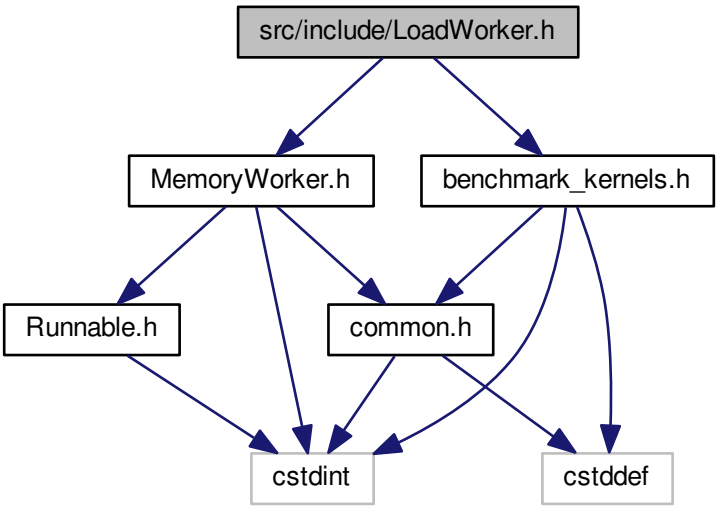


This graph shows which files directly or indirectly include this file:

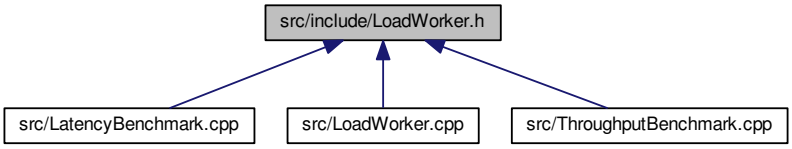


### Classes

- class xmem::LoadWorker

  *Multithreading-friendly class to do memory loading.*

### 6.18.1 Detailed Description

Header file for the LoadWorker class.

## 6.19 src/include/MemoryWorker.h File Reference

Header file for the MemoryWorker class.

```
#include <common.h>
#include <Runnable.h>
#include <cstdint>
```

Include dependency graph for MemoryWorker.h:

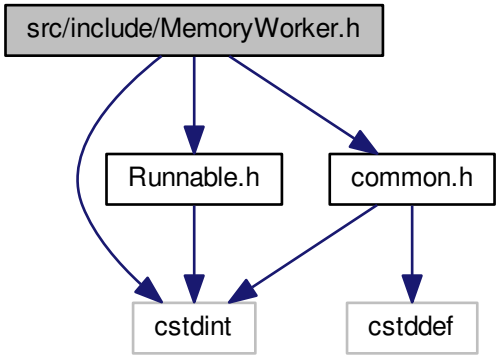


This graph shows which files directly or indirectly include this file:

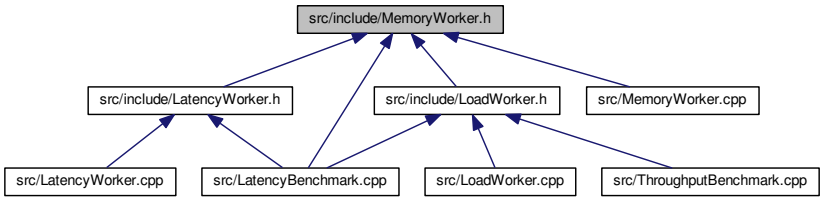


### Classes

- class xmem::MemoryWorker

  *Multithreading-friendly class to run memory access kernels.*

### 6.19.1 Detailed Description

Header file for the MemoryWorker class.

## 6.20 src/include/MyArg.h File Reference

Extensions to third-party optionparser-related code.

```
#include <ExampleArg.h>
#include <cstdint>
#include <stdio.h>
#include <cstdlib>
```

Include dependency graph for MyArg.h:

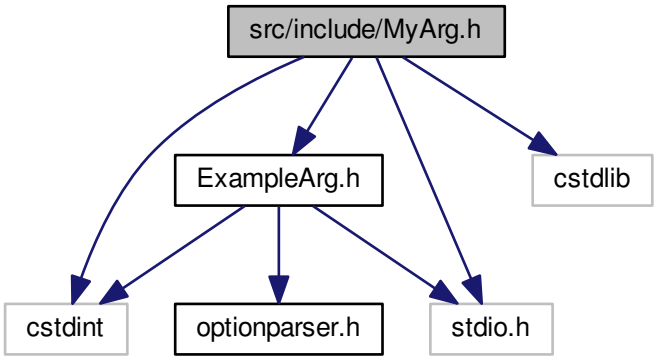


This graph shows which files directly or indirectly include this file:

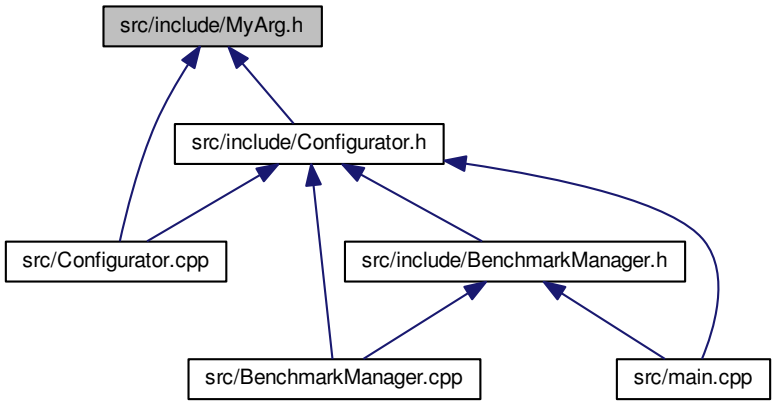


## Classes

- class xmem::MyArg

### 6.20.1 Detailed Description

Extensions to third-party optionparser-related code.

## 6.21 src/include/optionparser.h File Reference

This is the only file required to use The Lean Mean C++ Option Parser. Just #include it and you're set.

This graph shows which files directly or indirectly include this file:

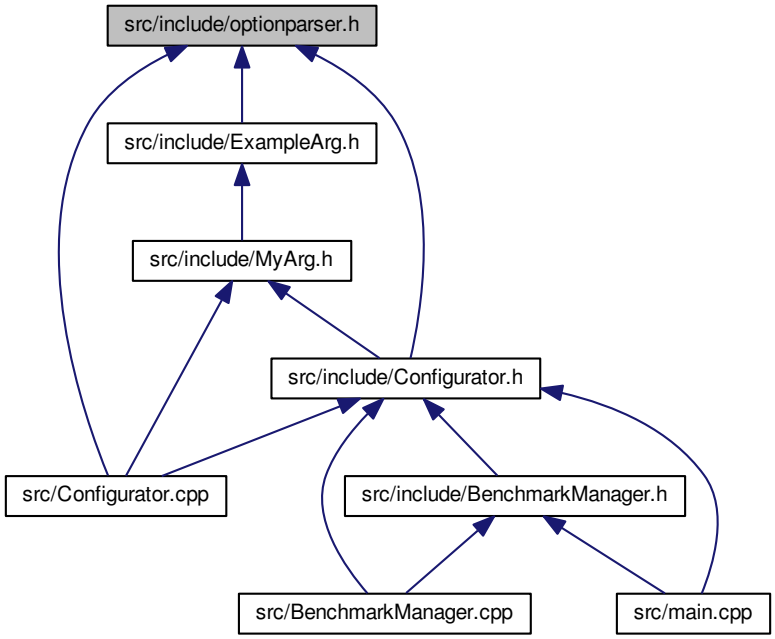


### Classes

- struct xmem::Descriptor

  *Describes an option, its help text (usage) and how it should be parsed.*
- class xmem::Option

  *A parsed option from the command line together with its argument if it has one.*
- struct xmem::Arg

  *Functions for checking the validity of option arguments.*
- struct xmem::Stats

  *Determines the minimum lengths of the buffer and options arrays used for Parser.*
- class xmem::Parser

  *Checks argument vectors for validity and parses them into data structures that are easier to work with.*
- struct xmem::Parser::Action
- class xmem::Stats::CountOptionsAction
- class xmem::Parser::StoreOptionAction
- struct xmem::PrintUsageImplementation
- struct xmem::PrintUsageImplementation::IStringWriter
- struct xmem::PrintUsageImplementation::FunctionWriter< Function >
- struct xmem::PrintUsageImplementation::OStreamWriter< OStream >
- struct xmem::PrintUsageImplementation::TemporaryWriter< Temporary >
- struct xmem::PrintUsageImplementation::SyscallWriter< Syscall >
- struct xmem::PrintUsageImplementation::StreamWriter< Function, Stream >
- class xmem::PrintUsageImplementation::LinePartIterator
- class xmem::PrintUsageImplementation::LineWrapper

### Typedefs

- typedef ArgStatus(∗ **xmem::CheckArg** )(const Option &option, bool msg)

  *Signature of functions that check if an argument is valid for a certain type of option.*

### Enumerations

- enum **ArgStatus** { **xmem::ARG_NONE**, **xmem::ARG_OK**, **xmem::ARG_IGNORE**, **xmem::ARG_ILLEGAL** }

  *Possible results when checking if an argument is valid for a certain option.*

### Functions

- template<typename OStream >
  void **xmem::printUsage** (OStream &prn, const Descriptor usage[], int width=80, int last_column_min_-percent=50, int last_column_own_line_max_percent=75)

  *Outputs a nicely formatted usage string with support for multi-column formatting and line-wrapping.*
- template<typename Function >
  void **xmem::printUsage** (Function ∗prn, const Descriptor usage[], int width=80, int last_column_min_-percent=50, int last_column_own_line_max_percent=75)
- template<typename Temporary >
  void **xmem::printUsage** (const Temporary &prn, const Descriptor usage[], int width=80, int last_column_-min_percent=50, int last_column_own_line_max_percent=75)
- template<typename Syscall >
  void **xmem::printUsage** (Syscall ∗prn, int fd, const Descriptor usage[], int width=80, int last_column_min_-percent=50, int last_column_own_line_max_percent=75)
- template<typename Function , typename Stream >
  void **xmem::printUsage** (Function ∗prn, Stream ∗stream, const Descriptor usage[], int width=80, int last_-column_min_percent=50, int last_column_own_line_max_percent=75)

#### 6.21.1 Detailed Description

This is the only file required to use The Lean Mean C++ Option Parser. Just #include it and you're set. The Lean Mean C++ Option Parser handles the program's command line arguments (argc, argv). It supports the short and long option formats of getopt(), getopt_long() and getopt_long_only() but has a more convenient interface. The following features set it apart from other option parsers:

**Highlights:**

- It is a header-only library. Just `#include "optionparser.h"` and you're set.
- It is freestanding. There are no dependencies whatsoever, not even the C or C++ standard library.
- It has a usage message formatter that supports column alignment and line wrapping. This aids localization because it adapts to translated strings that are shorter or longer (even if they contain Asian wide characters).
- Unlike getopt() and derivatives it doesn't force you to loop through options sequentially. Instead you can access options directly like this:
  - Test for presence of a switch in the argument vector:
    ```
    if ( options[QUIET] ) ...
    ```
  - Evaluate –enable-foo/–disable-foo pair where the last one used wins:
    ```
    if ( options[FOO].last()->type() == DISABLE ) ...
    ```
  - Cumulative option (-v verbose, -vv more verbose, -vvv even more verbose):
    ```
    int verbosity = options[VERBOSE].count();
    ```
  - Iterate over all –file=<fname> arguments:

```
for (Option* opt = options[FILE]; opt; opt = opt->next())
 fname = opt->arg; ...
```

– If you really want to, you can still process all arguments in order:

```
for (int i = 0; i < p.optionsCount(); ++i) {
  Option& opt = buffer[i];
  switch(opt.index()) {
    case HELP:    ...
    case VERBOSE: ...
    case FILE:    fname = opt.arg; ...
    case UNKNOWN: ...
```

Despite these features the code size remains tiny. It is smaller than uClibc's GNU getopt() and just a couple 100 bytes larger than uClibc's SUSv3 getopt().
(This does not include the usage formatter, of course. But you don't have to use that.)

**Download:**

Tarball with examples and test programs: optionparser-1.3.tar.gz
Just the header (this is all you really need): optionparser.h

**Changelog:**

**Version 1.3:** Compatible with Microsoft Visual C++.
**Version 1.2:** Added Option::namelen and removed the extraction of short option characters into a special buffer.
Changed Arg::Optional to accept arguments if they are attached rather than separate. This is what GNU getopt() does and how POSIX recommends utilities should interpret their arguments.
**Version 1.1:** Optional mode with argument reordering as done by GNU getopt(), so that options and non-options can be mixed. See Parser::parse().

**Feedback:**

Send questions, bug reports, feature requests etc. to: **optionparser-feedback (a) lists.sourceforge.net**

**Example program:**

(Note: option::* identifiers are links that take you to their documentation.)

```
#include <iostream>
#include "optionparser.h"

enum  optionIndex { UNKNOWN, HELP, PLUS };
const option::Descriptor usage[] =
{
 {UNKNOWN, 0,"" , ""    ,option::Arg::None, "USAGE: example [options]\n\n"
                                            "Options:" },
 {HELP,    0,"" , "help",option::Arg::None, "  --help  \tPrint usage and exit." },
 {PLUS,    0,"p", "plus",option::Arg::None, "  --plus, -p  \tIncrement count." },
 {UNKNOWN, 0,"" ,  ""   ,option::Arg::None, "\nExamples:\n"
                                            "  example --unknown -- --this_is_no_option\n"
                                            "  example -unk --plus -ppp file1 file2\n" },

 {0,0,0,0,0,0}
};

int main(int argc, char* argv[])
{
  argc-=(argc>0); argv+=(argc>0); // skip program name argv[0] if present
  option::Stats  stats(usage, argc, argv);
  option::Option options[stats.options_max], buffer[stats.buffer_max];
  option::Parser parse(usage, argc, argv, options, buffer);

  if (parse.error())
    return 1;

  if (options[HELP] || argc == 0) {
    option::printUsage(std::cout, usage);
    return 0;
  }

  std::cout << "--plus count: " <<
    options[PLUS].count() << "\n";

  for (option::Option* opt = options[UNKNOWN]; opt; opt = opt->next())
    std::cout << "Unknown option: " << opt->name << "\n";
```

```
  for (int i = 0; i < parse.nonOptionsCount(); ++i)
    std::cout << "Non-option #" << i << ": " << parse.nonOption(i) << "\n";
}
```

**Option syntax:**

- The Lean Mean C++ Option Parser follows POSIX `getopt()` conventions and supports GNU-style `getopt_long()` long options as well as Perl-style single-minus long options (`getopt_long_only()`).
- short options have the format `-X` where `X` is any character that fits in a char.
- short options can be grouped, i.e. `-X -Y` is equivalent to `-XY`.
- a short option may take an argument either separate (`-X foo`) or attached (`-Xfoo`). You can make the parser accept the additional format `-X=foo` by registering `X` as a long option (in addition to being a short option) and enabling single-minus long options.
- an argument-taking short option may be grouped if it is the last in the group, e.g. `-ABCXfoo` or `-ABCX foo` (`foo` is the argument to the `-X` option).
- a lone minus character `'-'` is not treated as an option. It is customarily used where a file name is expected to refer to stdin or stdout.
- long options have the format `-option-name`.
- the option-name of a long option can be anything and include any characters. Even = characters will work, but don't do that.
- [optional] long options may be abbreviated as long as the abbreviation is unambiguous. You can set a minimum length for abbreviations.
- [optional] long options may begin with a single minus. The double minus form is always accepted, too.
- a long option may take an argument either separate (`-option arg` ) or attached (`-option=arg` ). In the attached form the equals sign is mandatory.
- an empty string can be passed as an attached long option argument: `-option-name=` . Note the distinction between an empty string as argument and no argument at all.
- an empty string is permitted as separate argument to both long and short options.
- Arguments to both short and long options may start with a `'-'` character. E.g. `-X-X` , `-X -X` or `-long-X=-X` . If `-X` and `-long-X` take an argument, that argument will be `"-X"` in all 3 cases.
- If using the built-in Arg::Optional, optional arguments must be attached.
- the special option - (i.e. without a name) terminates the list of options. Everything that follows is a non-option argument, even if it starts with a `'-'` character. The - itself will not appear in the parse results.
- the first argument that doesn't start with `'-'` or `'-'` and does not belong to a preceding argument-taking option, will terminate the option list and is the first non-option argument. All following command line arguments are treated as non-option arguments, even if they start with `'-'` .

  NOTE: This behaviour is mandated by POSIX, but GNU getopt() only honours this if it is explicitly requested (e.g. by setting POSIXLY_CORRECT).

  You can enable the GNU behaviour by passing `true` as first argument to e.g. Parser::parse().
- Arguments that look like options (i.e. `'-'` followed by at least 1 character) but aren't, are NOT treated as non-option arguments. They are treated as unknown options and are collected into a list of unknown options for error reporting.

  This means that in order to pass a first non-option argument beginning with the minus character it is required to use the - special option, e.g.

  ```
  program -x -- --strange-filename
  ```

  In this example, `-strange-filename` is a non-option argument. If the - were omitted, it would be treated as an unknown option.

  See option::Descriptor::longopt for information on how to collect unknown options.

## 6.22 src/include/PowerReader.h File Reference

Header file for the PowerReader class.

```
#include <common.h>
#include <Runnable.h>
#include <cstdint>
#include <vector>
#include <string>
```

Include dependency graph for PowerReader.h:

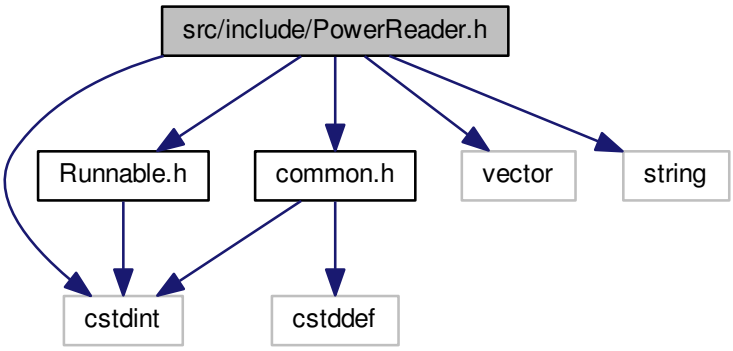


This graph shows which files directly or indirectly include this file:

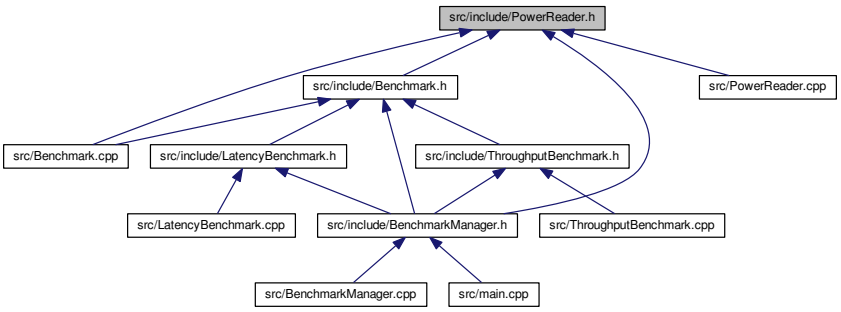


### Classes

- class xmem::PowerReader

  *An abstract base class for measuring power from an arbitrary source. This class is runnable using a worker thread.*

### 6.22.1 Detailed Description

Header file for the PowerReader class.

## 6.23 src/include/Runnable.h File Reference

Header file for the Runnable class.

```
#include <cstdint>
```

Include dependency graph for Runnable.h:

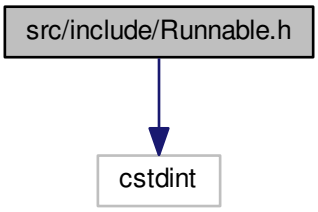


This graph shows which files directly or indirectly include this file:

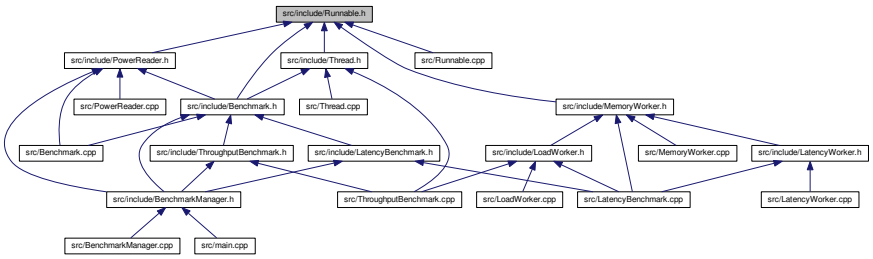


### Classes

- class xmem::Runnable

  *A base class for any object that implements a thread-safe run() function for use by Thread objects.*

### 6.23.1 Detailed Description

Header file for the Runnable class.

## 6.24 src/include/Thread.h File Reference

Header file for the Thread class.

```
#include <Runnable.h>
#include <cstdint>
```

Include dependency graph for Thread.h:

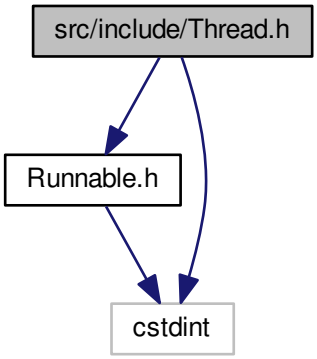


This graph shows which files directly or indirectly include this file:

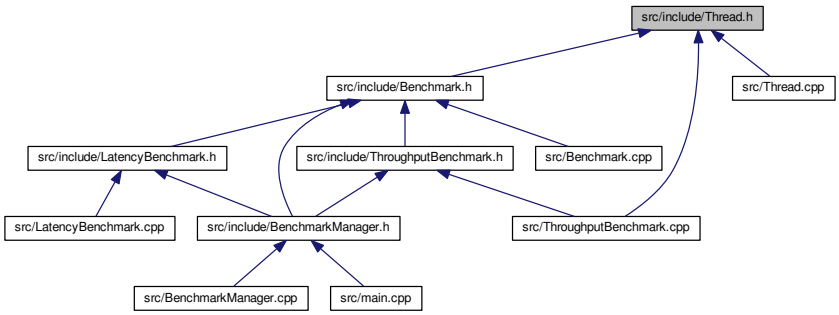


## Classes

- class xmem::Thread

  *a nice wrapped thread interface independent of particular OS API*

### 6.24.1 Detailed Description

Header file for the Thread class.

## 6.25 src/include/ThroughputBenchmark.h File Reference

Header file for the ThroughputBenchmark class.

```
#include <Benchmark.h>
#include <common.h>
#include <cstdint>
#include <string>
```

Include dependency graph for ThroughputBenchmark.h:

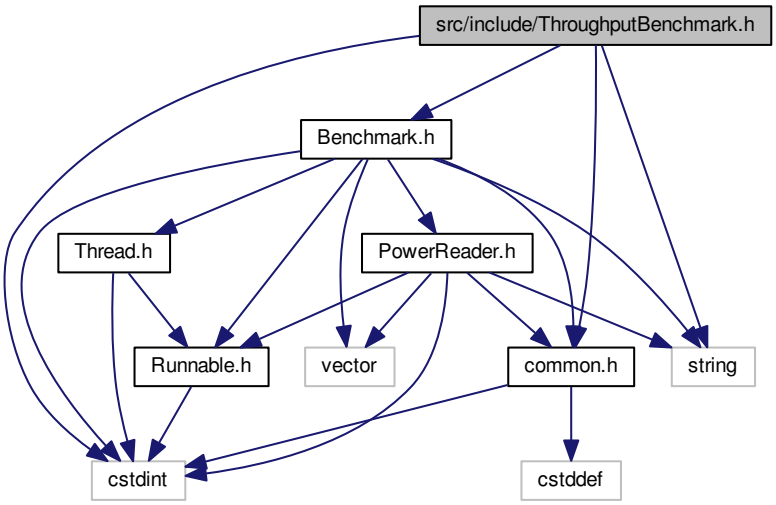


This graph shows which files directly or indirectly include this file:

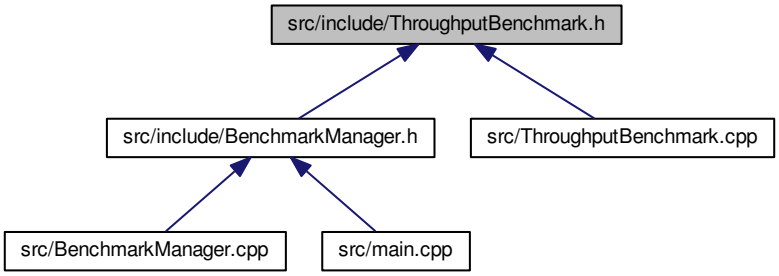


### Classes

- class xmem::ThroughputBenchmark

  *A type of benchmark that measures memory throughput.*

### 6.25.1 Detailed Description

Header file for the ThroughputBenchmark class.

## 6.26 src/include/Timer.h File Reference

Header file for the Timer class.

```
#include <common.h>
#include <cstdint>
```

Include dependency graph for Timer.h:

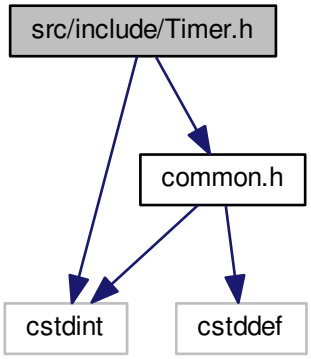


This graph shows which files directly or indirectly include this file:

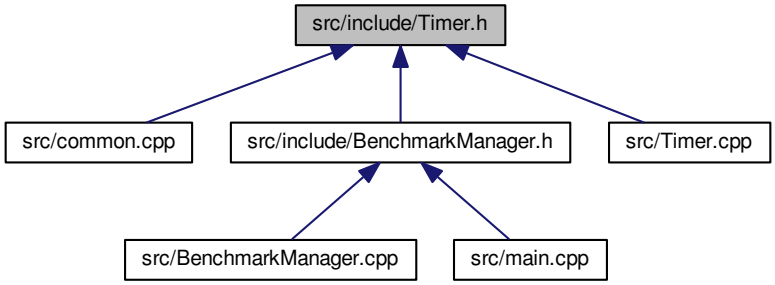


### Classes

- class xmem::Timer

  *This class abstracts some characteristics of simple high resolution stopwatch timer. However, due to the inability or complexity of abstracting shared hardware timers, this class does not actually provide start and stop functions.*

#### 6.26.1 Detailed Description

Header file for the Timer class.

## 6.27 src/include/win/win_common_third_party.h File Reference

Header file for some third-party helper code for working with Windows APIs.

### 6.27.1 Detailed Description

Header file for some third-party helper code for working with Windows APIs.

## 6.28 src/include/win/win_CPdhQuery.h File Reference

Header and implementation file for some third-party code for measuring Windows OS-exposed performance counters.

### 6.28.1 Detailed Description

Header and implementation file for some third-party code for measuring Windows OS-exposed performance counters.

## 6.29 src/include/win/WindowsDRAMPowerReader.h File Reference

Header file for the WindowsDRAMPowerReader class.

### 6.29.1 Detailed Description

Header file for the WindowsDRAMPowerReader class.

## 6.30 src/LatencyBenchmark.cpp File Reference

Implementation file for the LatencyBenchmark class.

```
#include <LatencyBenchmark.h>
#include <common.h>
#include <benchmark_kernels.h>
#include <MemoryWorker.h>
#include <LatencyWorker.h>
#include <LoadWorker.h>
#include <iostream>
#include <random>
#include <assert.h>
#include <time.h>
```

Include dependency graph for LatencyBenchmark.cpp:

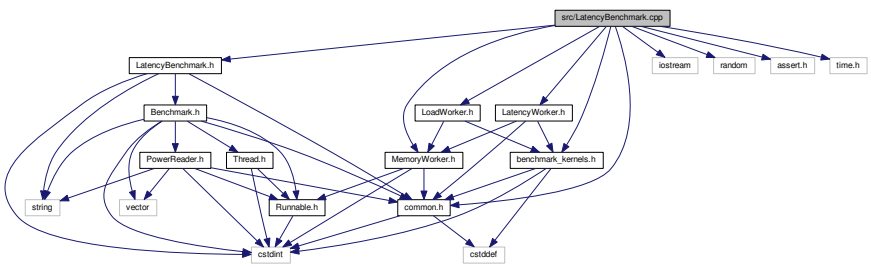

### 6.30.1 Detailed Description

Implementation file for the LatencyBenchmark class.

## 6.31 src/LatencyWorker.cpp File Reference

Implementation file for the LatencyWorker class.

```
#include <LatencyWorker.h>
#include <benchmark_kernels.h>
#include <common.h>
#include <iostream>
```

Include dependency graph for LatencyWorker.cpp:

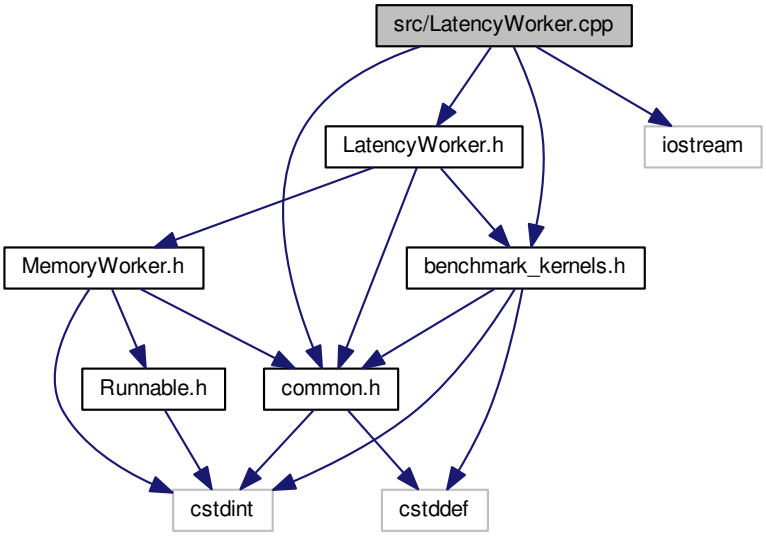


### 6.31.1 Detailed Description

Implementation file for the LatencyWorker class.

## 6.32 src/LoadWorker.cpp File Reference

Implementation file for the LoadWorker class.

```
#include <LoadWorker.h>
#include <benchmark_kernels.h>
#include <common.h>
#include <iostream>
```

Include dependency graph for LoadWorker.cpp:

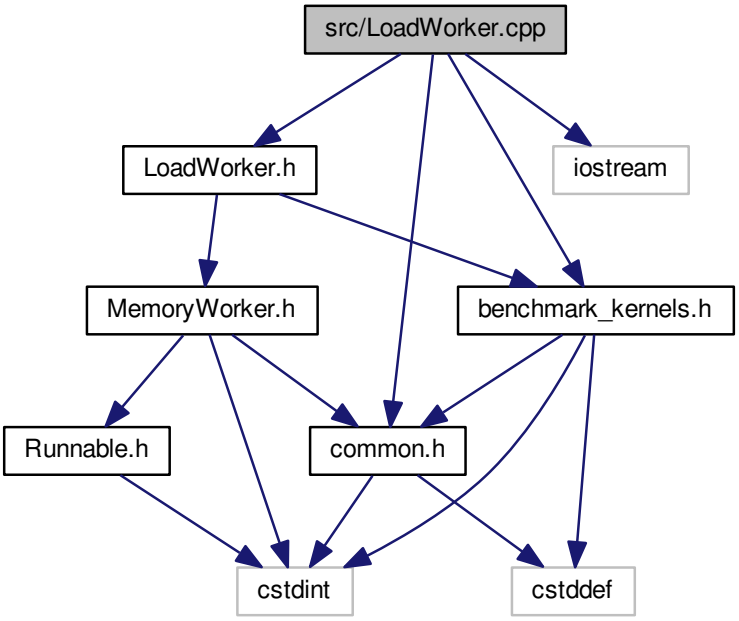


### 6.32.1 Detailed Description

Implementation file for the LoadWorker class.

## 6.33 src/main.cpp File Reference

main entry point to the tool

```
#include <common.h>
#include <build_datetime.h>
#include <Configurator.h>
#include <BenchmarkManager.h>
#include <iostream>
#include <string>
```

Include dependency graph for main.cpp:

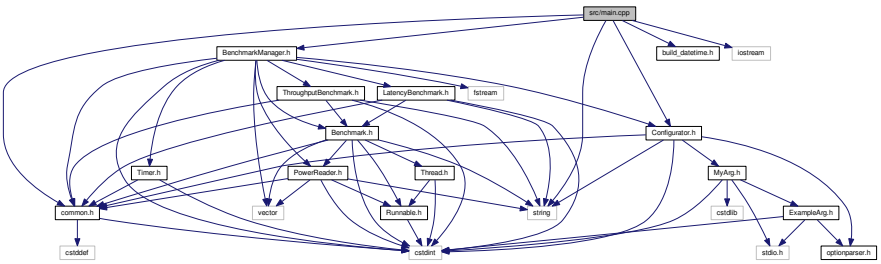

### Functions

- int main (int argc, char ∗argv[])

  *The main entry point to the program.*

### 6.33.1 Detailed Description

main entry point to the tool This tool is designed to measure bandwidth and latency of the memory system using several access patterns, strides, and working set sizes. The primary goal is to measure DRAM performance, although it can also measure cache performance depending on the configuration.

## 6.34 src/MemoryWorker.cpp File Reference

Implementation file for the MemoryWorker class.

```
#include <MemoryWorker.h>
#include <common.h>
```

Include dependency graph for MemoryWorker.cpp:

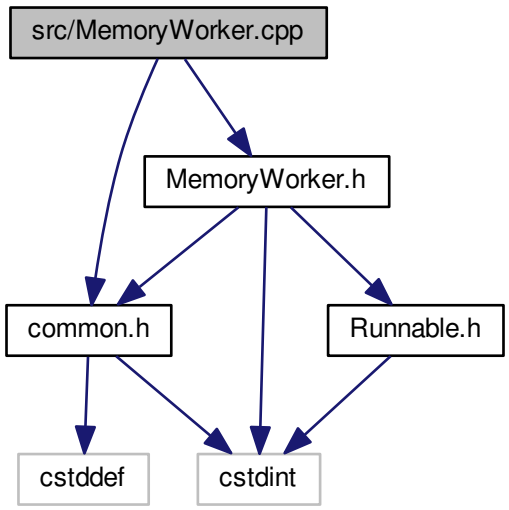


### 6.34.1 Detailed Description

Implementation file for the MemoryWorker class.

## 6.35 src/PowerReader.cpp File Reference

Implementation file for the PowerReader class.

```
#include <PowerReader.h>
#include <common.h>
#include <cstdint>
#include <vector>
#include <iostream>
```

Include dependency graph for PowerReader.cpp:

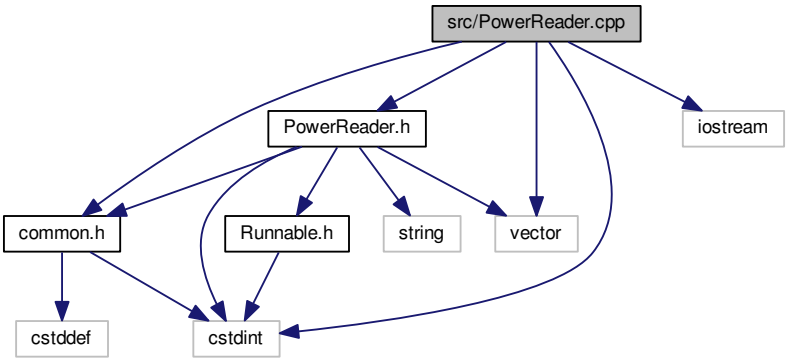


### 6.35.1 Detailed Description

Implementation file for the PowerReader class.

## 6.36 src/Runnable.cpp File Reference

Implementation file for the Runnable class.

```
#include <Runnable.h>
#include <iostream>
```

Include dependency graph for Runnable.cpp:

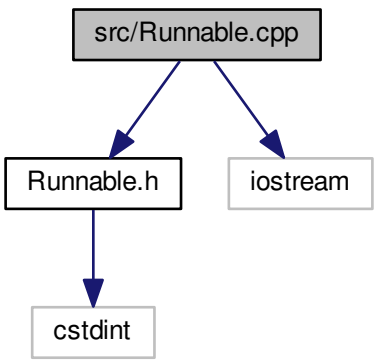

### Variables

- return **false**

### 6.36.1 Detailed Description

Implementation file for the Runnable class.

## 6.37 src/Thread.cpp File Reference

Implementation file for the Thread class.

```
#include <Thread.h>
#include <stdlib.h>
#include <iostream>
```

Include dependency graph for Thread.cpp:

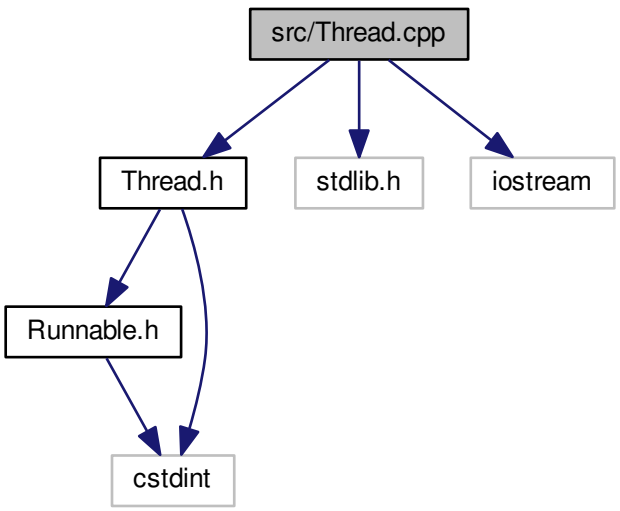


### Variables

- return **false**
- return **true**

### 6.37.1 Detailed Description

Implementation file for the Thread class.

## 6.38 src/ThroughputBenchmark.cpp File Reference

Implementation file for the ThroughputBenchmark class.

```
#include <ThroughputBenchmark.h>
#include <common.h>
#include <LoadWorker.h>
#include <Thread.h>
#include <iostream>
#include <assert.h>
#include <time.h>
```

Include dependency graph for ThroughputBenchmark.cpp:

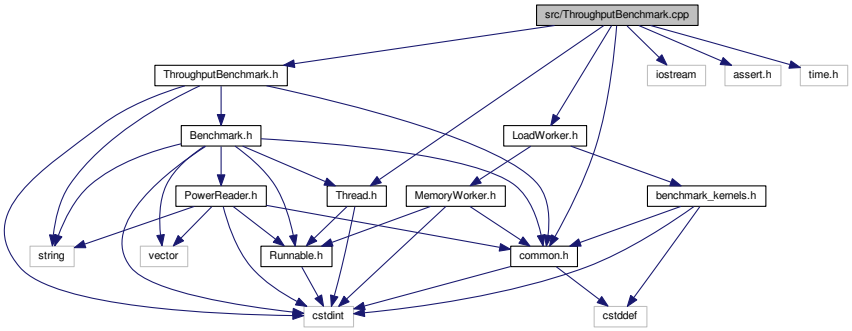


### 6.38.1 Detailed Description

Implementation file for the ThroughputBenchmark class.

## 6.39 src/Timer.cpp File Reference

Implementation file for the Timer class.

```
#include <Timer.h>
#include <common.h>
```

Include dependency graph for Timer.cpp:

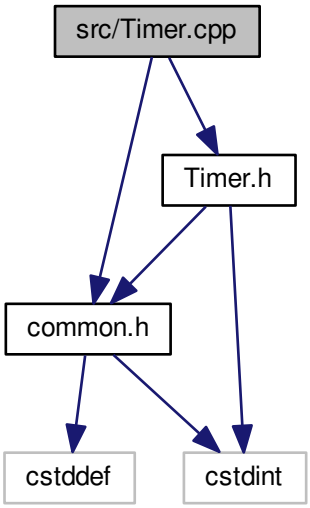


### 6.39.1 Detailed Description

Implementation file for the Timer class.

## 6.40 src/win/win_common_third_party.cpp File Reference

Implementation file for some third-party helper code for working with Windows APIs.

### 6.40.1 Detailed Description

Implementation file for some third-party helper code for working with Windows APIs.

## 6.41 src/win/WindowsDRAMPowerReader.cpp File Reference

Implementation file for the WindowsDRAMPowerReader class.

### 6.41.1 Detailed Description

Implementation file for the WindowsDRAMPowerReader class.

# Index


~Thread
    xmem::Thread, 88
_acquireLock
    xmem::Runnable, 80
_adjusted_ticks
    xmem::MemoryWorker, 54
_averageLoadMetric
    xmem::LatencyBenchmark, 42
_averageMetric
    xmem::Benchmark, 22
_average_dram_power_socket
    xmem::Benchmark, 22
_average_power
    xmem::PowerReader, 77
_bytes_per_pass
    xmem::MemoryWorker, 54
_chunk_size
    xmem::Benchmark, 22
_completed
    xmem::MemoryWorker, 54
_cpu_affinity
    xmem::MemoryWorker, 54
    xmem::PowerReader, 77
_cpu_node
    xmem::Benchmark, 22
_dram_power_readers
    xmem::Benchmark, 22
_dram_power_threads
    xmem::Benchmark, 22
_elapsed_dummy_ticks
    xmem::MemoryWorker, 54
_elapsed_ticks
    xmem::MemoryWorker, 54
_hasRun
    xmem::Benchmark, 22
_iterations
    xmem::Benchmark, 23
_len
    xmem::Benchmark, 23
    xmem::MemoryWorker, 54
_loadMetricOnIter
    xmem::LatencyBenchmark, 42
_mem_array
    xmem::Benchmark, 23
    xmem::MemoryWorker, 54
_mem_node
    xmem::Benchmark, 23
_metricOnIter
    xmem::Benchmark, 23
_metricUnits
    xmem::Benchmark, 23
_name
    xmem::Benchmark, 23
    xmem::PowerReader, 77
_ns_per_tick
    xmem::Timer, 93
_num_samples
    xmem::PowerReader, 77
_num_worker_threads
    xmem::Benchmark, 23
_obj_valid
    xmem::Benchmark, 23
_passes
    xmem::MemoryWorker, 54
_pattern_mode
    xmem::Benchmark, 23
_peak_dram_power_socket
    xmem::Benchmark, 23
_peak_power
    xmem::PowerReader, 78
_power_trace
    xmem::PowerReader, 78
_power_units
    xmem::PowerReader, 78
_releaseLock
    xmem::Runnable, 80
_run_core
    xmem::Benchmark, 18
    xmem::LatencyBenchmark, 41
    xmem::ThroughputBenchmark, 91
_rw_mode
    xmem::Benchmark, 23
_sampling_period
    xmem::PowerReader, 78
_start_power_threads
    xmem::Benchmark, 18
_stop_power_threads
    xmem::Benchmark, 18
_stop_signal
    xmem::PowerReader, 78
_stride_size
    xmem::Benchmark, 24
_ticks_per_ms
    xmem::Timer, 93
_warning
    xmem::Benchmark, 24
    xmem::MemoryWorker, 55

add

xmem::Stats, 82
append
xmem::Option, 59
arg
xmem::Option, 63

Benchmark
xmem::Benchmark, 18
BenchmarkManager
xmem::BenchmarkManager, 24
buffer_max
xmem::Stats, 82

calculateMetrics
xmem::PowerReader, 72
cancel
xmem::Thread, 89
check_arg
xmem::Descriptor, 35
clear
xmem::PowerReader, 72
clear_and_reset
xmem::PowerReader, 72
common.h
DEFAULT_NUM_NODES, 109
DEFAULT_PAGE_SIZE, 109
MIN_ELAPSED_TICKS, 109
USE_OS_TIMER, 110
completed
xmem::Thread, 89
configureFromInput
xmem::Configurator, 28
count
xmem::Option, 59
CountOptionsAction
xmem::Stats::CountOptionsAction, 34
create_and_start
xmem::Thread, 89
created
xmem::Thread, 89

DEFAULT_NUM_NODES
common.h, 109
DEFAULT_PAGE_SIZE
common.h, 109
desc
xmem::Option, 63

error
xmem::Parser, 67
extensionsEnabled
xmem::Configurator, 28

finished
xmem::Parser::Action, 13
xmem::Parser::StoreOptionAction, 84
first
xmem::Option, 59

get_ns_per_tick
xmem::Timer, 93
get_ticks_per_ms
xmem::Timer, 93
getAdjustedTicks
xmem::MemoryWorker, 50
getAverageDRAMPower
xmem::Benchmark, 19
getAverageMetric
xmem::Benchmark, 19
getAveragePower
xmem::PowerReader, 73
getAvgLoadMetric
xmem::LatencyBenchmark, 41
getBytesPerPass
xmem::MemoryWorker, 51
getCPUNode
xmem::Benchmark, 19
getChunkSize
xmem::Benchmark, 19
getElapsedDummyTicks
xmem::MemoryWorker, 51
getElapsedTicks
xmem::MemoryWorker, 52
getExitCode
xmem::Thread, 89
getIterations
xmem::Benchmark, 19
getIterationsPerTest
xmem::Configurator, 28
getLastSample
xmem::PowerReader, 73
getLen
xmem::Benchmark, 19
xmem::MemoryWorker, 52
getLoadMetricOnIter
xmem::LatencyBenchmark, 41
getMemNode
xmem::Benchmark, 20
getMetricOnIter
xmem::Benchmark, 20
getMetricUnits
xmem::Benchmark, 20
getName
xmem::Benchmark, 20
getNumSamples
xmem::PowerReader, 74
getNumThreads
xmem::Benchmark, 20
getNumWorkerThreads
xmem::Configurator, 29
getOutputFilename
xmem::Configurator, 29
getPasses
xmem::MemoryWorker, 53
getPatternMode
xmem::Benchmark, 20
getPeakDRAMPower
xmem::Benchmark, 21

getPeakPower
xmem::PowerReader, 74
getPowerTrace
xmem::PowerReader, 75
getPowerUnits
xmem::PowerReader, 75
getRWMode
xmem::Benchmark, 21
getSamplingPeriod
xmem::PowerReader, 76
getStartingTestIndex
xmem::Configurator, 29
getStrideSize
xmem::Benchmark, 21
getTarget
xmem::Thread, 89
getWorkingSetSizePerThread
xmem::Configurator, 29

hadWarning
xmem::MemoryWorker, 53
hasRun
xmem::Benchmark, 21
help
xmem::Descriptor, 35

index
xmem::Descriptor, 35
isFirst
xmem::Option, 60
isLast
xmem::Option, 60
isNUMAEnabled
xmem::Configurator, 29
isThreadRunning
xmem::Thread, 89
isThreadSuspended
xmem::Thread, 89
isValid
xmem::Benchmark, 21
isWideChar
xmem::PrintUsageImplementation, 79

join
xmem::Thread, 90

last
xmem::Option, 60
latencyTestSelected
xmem::Configurator, 29
LatencyWorker
xmem::LatencyWorker, 43
LineWrapper
xmem::PrintUsageImplementation::LineWrapper, 45
LoadWorker
xmem::LoadWorker, 48
longopt
xmem::Descriptor, 35

MIN_ELAPSED_TICKS
common.h, 109
MemoryWorker
xmem::MemoryWorker, 50

name
xmem::Option, 63
xmem::PowerReader, 76
namelen
xmem::Option, 63
next
xmem::Option, 61
xmem::PrintUsageImplementation::LinePartIterator, 44
nextRow
xmem::PrintUsageImplementation::LinePartIterator, 44
nextTable
xmem::PrintUsageImplementation::LinePartIterator, 45
nextwrap
xmem::Option, 61
nonOptions
xmem::Parser, 68
nonOptionsCount
xmem::Parser, 68

operator const Option ∗
xmem::Option, 61
operator Option ∗
xmem::Option, 62
operator=
xmem::Option, 62
Option
xmem::Option, 58
options_max
xmem::Stats, 82
optionsCount
xmem::Parser, 68

parse
xmem::Parser, 68
Parser
xmem::Parser, 66
perform
xmem::Parser::Action, 14
xmem::Parser::StoreOptionAction, 84
xmem::Stats::CountOptionsAction, 34
PowerReader
xmem::PowerReader, 71
prev
xmem::Option, 62
prevwrap
xmem::Option, 62
process
xmem::PrintUsageImplementation::LineWrapper, 46

run

xmem::Benchmark, 21
runAll
xmem::BenchmarkManager, 25
runExtDelayInjectedLoadedLatencyBenchmark
xmem::Configurator, 29
runLatencyBenchmarks
xmem::BenchmarkManager, 25
runThroughputBenchmarks
xmem::BenchmarkManager, 26

setUseOutputFile
xmem::Configurator, 30
shortopt
xmem::Descriptor, 36
src/Benchmark.cpp, 95
src/BenchmarkManager.cpp, 96
src/Configurator.cpp, 98
src/LatencyBenchmark.cpp, 130
src/LatencyWorker.cpp, 131
src/LoadWorker.cpp, 131
src/MemoryWorker.cpp, 133
src/PowerReader.cpp, 133
src/Runnable.cpp, 134
src/Thread.cpp, 135
src/ThroughputBenchmark.cpp, 135
src/Timer.cpp, 136
src/benchmark_kernels.cpp, 95
src/common.cpp, 97
src/ext/DelayInjectedLoadedLatencyBenchmark/DelayInjectedLoadedLatencyBenchmark.cpp, 99
src/ext/DelayInjectedLoadedLatencyBenchmark/delay_injected_benchmark_kernels.cpp, 99
src/include/Benchmark.h, 100
src/include/BenchmarkManager.h, 104
src/include/Configurator.h, 110
src/include/ExampleArg.h, 111
src/include/LatencyBenchmark.h, 115
src/include/LatencyWorker.h, 117
src/include/LoadWorker.h, 118
src/include/MemoryWorker.h, 119
src/include/MyArg.h, 119
src/include/PowerReader.h, 125
src/include/Runnable.h, 126
src/include/Thread.h, 126
src/include/ThroughputBenchmark.h, 127
src/include/Timer.h, 128
src/include/benchmark_kernels.h, 101
src/include/common.h, 105
src/include/ext/DelayInjectedLoadedLatencyBenchmark/DelayInjectedLoadedLatencyBenchmark.h, 115
src/include/ext/DelayInjectedLoadedLatencyBenchmark/delay_injected_benchmark_kernels.h, 112
src/include/optionparser.h, 121
src/include/win/WindowsDRAMPowerReader.h, 130
src/include/win/win_CPdhQuery.h, 130
src/include/win/win_common_third_party.h, 129
src/main.cpp, 132
src/win/WindowsDRAMPowerReader.cpp, 137
src/win/win_common_third_party.cpp, 137
started
xmem::Thread, 90
Stats
xmem::Stats, 81
stop
xmem::PowerReader, 77
StoreOptionAction
xmem::Parser::StoreOptionAction, 83

Thread
xmem::Thread, 88
throughputTestSelected
xmem::Configurator, 30
type
xmem::Descriptor, 36
xmem::Option, 62

USE_OS_TIMER
common.h, 110
useChunk32b
xmem::Configurator, 30
useLargePages
xmem::Configurator, 30
useOutputFile
xmem::Configurator, 30
useRandomAccessPattern
xmem::Configurator, 30
useReads
xmem::Configurator, 30
useSequentialAccessPattern
xmem::Configurator, 31
useStrideN1
xmem::Configurator, 31
useStrideN16
xmem::Configurator, 31
useStrideN2
xmem::Configurator, 31
useStrideN4
xmem::Configurator, 31
useStrideN8
xmem::Configurator, 31
useStrideP1
xmem::Configurator, 31
useStrideP16
xmem::Configurator, 32
useStrideP2
xmem::Configurator, 32
useStrideP4
xmem::Configurator, 32
useStrideP8
xmem::Configurator, 32
useWrites
xmem::Configurator, 32

validTarget
xmem::Thread, 90
verboseMode
xmem::Configurator, 32

xmem::Arg, 14
xmem::Benchmark, 15
_averageMetric, 22
_average_dram_power_socket, 22
_chunk_size, 22
_cpu_node, 22
_dram_power_readers, 22
_dram_power_threads, 22
_hasRun, 22
_iterations, 23
_len, 23
_mem_array, 23
_mem_node, 23
_metricOnIter, 23
_metricUnits, 23
_name, 23
_num_worker_threads, 23
_obj_valid, 23
_pattern_mode, 23
_peak_dram_power_socket, 23
_run_core, 18
_rw_mode, 23
_start_power_threads, 18
_stop_power_threads, 18
_stride_size, 24
_warning, 24
Benchmark, 18
getAverageDRAMPower, 19
getAverageMetric, 19
getCPUNode, 19
getChunkSize, 19
getIterations, 19
getLen, 19
getMemNode, 20
getMetricOnIter, 20
getMetricUnits, 20
getName, 20
getNumThreads, 20
getPatternMode, 20
getPeakDRAMPower, 21
getRWMode, 21
getStrideSize, 21
hasRun, 21
isValid, 21
run, 21
xmem::BenchmarkManager, 24
BenchmarkManager, 24
runAll, 25
runLatencyBenchmarks, 25
runThroughputBenchmarks, 26
xmem::Configurator, 26
configureFromInput, 28
extensionsEnabled, 28
getIterationsPerTest, 28
getNumWorkerThreads, 29
getOutputFilename, 29
getStartingTestIndex, 29
getWorkingSetSizePerThread, 29
isNUMAEnabled, 29
latencyTestSelected, 29
runExtDelayInjectedLoadedLatencyBenchmark, 29
setUseOutputFile, 30
throughputTestSelected, 30
useChunk32b, 30
useLargePages, 30
useOutputFile, 30
useRandomAccessPattern, 30
useReads, 30
useSequentialAccessPattern, 31
useStrideN1, 31
useStrideN16, 31
useStrideN2, 31
useStrideN4, 31
useStrideN8, 31
useStrideP1, 31
useStrideP16, 32
useStrideP2, 32
useStrideP4, 32
useStrideP8, 32
useWrites, 32
verboseMode, 32
xmem::Descriptor, 34
check_arg, 35
help, 35
index, 35
longopt, 35
shortopt, 36
type, 36
xmem::ExampleArg, 37
xmem::LatencyBenchmark, 39
_averageLoadMetric, 42
_loadMetricOnIter, 42
_run_core, 41
getAvgLoadMetric, 41
getLoadMetricOnIter, 41
xmem::LatencyWorker, 42
LatencyWorker, 43
xmem::LoadWorker, 46
LoadWorker, 48
xmem::MemoryWorker, 48
_adjusted_ticks, 54
_bytes_per_pass, 54
_completed, 54
_cpu_affinity, 54
_elapsed_dummy_ticks, 54
_elapsed_ticks, 54
_len, 54
_mem_array, 54
_passes, 54
_warning, 55
getAdjustedTicks, 50
getBytesPerPass, 51
getElapsedDummyTicks, 51
getElapsedTicks, 52
getLen, 52
getPasses, 53

hadWarning, 53
MemoryWorker, 50
xmem::MyArg, 55
xmem::Option, 56
append, 59
arg, 63
count, 59
desc, 63
first, 59
isFirst, 60
isLast, 60
last, 60
name, 63
namelen, 63
next, 61
nextwrap, 61
operator const Option *, 61
operator Option *, 62
operator=, 62
Option, 58
prev, 62
prevwrap, 62
type, 62
xmem::Parser, 65
error, 67
nonOptions, 68
nonOptionsCount, 68
optionsCount, 68
parse, 68
Parser, 66
xmem::Parser::Action, 13
finished, 13
perform, 14
xmem::Parser::StoreOptionAction, 83
finished, 84
perform, 84
StoreOptionAction, 83
xmem::PowerReader, 69
_average_power, 77
_cpu_affinity, 77
_name, 77
_num_samples, 77
_peak_power, 78
_power_trace, 78
_power_units, 78
_sampling_period, 78
_stop_signal, 78
calculateMetrics, 72
clear, 72
clear_and_reset, 72
getAveragePower, 73
getLastSample, 73
getNumSamples, 74
getPeakPower, 74
getPowerTrace, 75
getPowerUnits, 75
getSamplingPeriod, 76
name, 76
PowerReader, 71
stop, 77
xmem::PrintUsageImplementation, 78
isWideChar, 79
xmem::PrintUsageImplementation::FunctionWriter< Function >, 38
xmem::PrintUsageImplementation::IStringWriter, 39
xmem::PrintUsageImplementation::LinePartIterator, 44
next, 44
nextRow, 44
nextTable, 45
xmem::PrintUsageImplementation::LineWrapper, 45
LineWrapper, 45
process, 46
xmem::PrintUsageImplementation::OStreamWriter< OStream >, 64
xmem::PrintUsageImplementation::StreamWriter< Function, Stream >, 85
xmem::PrintUsageImplementation::SyscallWriter< Syscall >, 86
xmem::PrintUsageImplementation::TemporaryWriter< Temporary >, 87
xmem::Runnable, 79
_acquireLock, 80
_releaseLock, 80
xmem::Stats, 80
add, 82
buffer_max, 82
options_max, 82
Stats, 81
xmem::Stats::CountOptionsAction, 33
CountOptionsAction, 34
perform, 34
xmem::Thread, 88
~Thread, 88
cancel, 89
completed, 89
create_and_start, 89
created, 89
getExitCode, 89
getTarget, 89
isThreadRunning, 89
isThreadSuspended, 89
join, 90
started, 90
Thread, 88
validTarget, 90
xmem::ThroughputBenchmark, 90
_run_core, 91
xmem::Timer, 92
_ns_per_tick, 93
_ticks_per_ms, 93
get_ns_per_tick, 93
get_ticks_per_ms, 93